



TERTIARY INFOTECH ACADEMY PTE LTD

Learner Guide

Multi AI Agents Workflow for Content Creation

TGS-2023036153 · ICT-SNM-4004-1.1 Content Strategy

v1.0 · 13 September 2026

UEN 201200696W

Document Version Control Record

Version	Release date	Change	Owner
v1.0	13 September 2026	New multi-agent content strategy courseware aligned to current course page and TSC	Tertiary Infotech Academy

Contents

Learning outcomes

- LO1: Conceptualize content ideas to meet marketing objectives and map out digital storyboards as part of a content strategy.
- LO2: Identify content requirements from customer preferences and determine the frequency of marketing content.
- LO3: Determine content types and styles and decide on modes and processes for distributing content.
- LO4: Develop guidelines for content strategy execution using appropriate delivery modes and responsible AI practices.

How to use this guide

Complete the labs in order. Keep each handoff as a named artifact with evidence IDs. Lab procedures live below and in each lab folder. Use classroom mock data only.

Topic 1: Multi-AI-Agent Content Ideation and Digital Storyboarding

Campaign objective contract

Input: approved marketing brief. Control: translate objective into measurable audience action. Output: goal sheet. Accept when choose one primary conversion action. Watch for vague awareness goal. Measure goal and event IDs present.

Audience promise and proof

Input: product facts and objections. Control: pair each promise with verifiable evidence. Output: claim ledger. Accept when reject unsupported benefit claims. Watch for invented product proof. Measure supported claim rate.

Content idea scoring

Input: idea backlog and audience needs. Control: score relevance, evidence, effort and risk. Output: ranked idea board. Accept when advance ideas above agreed rubric threshold. Watch for novelty without audience fit. Measure rubric score by idea.

Agent role decomposition

Input: campaign brief and work units. Control: separate researcher, strategist, writer, designer and reviewer duties. Output: role charter. Accept when one accountable owner for each output. Watch for overlapping agent authority. Measure unowned task count.

Research-to-strategy handoff

Input: source cards with URLs and dates. Control: condense findings into claims and constraints. Output: research packet. Accept when accept only attributable findings. Watch for citation laundering. Measure traceable claim count.

Storyboard beat architecture

Input: approved message and CTA. Control: map hook, tension, proof, resolution and action. Output: five-beat storyboard. Accept when each beat advances the objective. Watch for beautiful scenes without narrative. Measure beat-to-objective coverage.

Scene-level evidence

Input: storyboard panels and product facts. Control: attach visual, voice, claim and source to each scene. Output: scene sheet. Accept when flag scenes with unsupported claims. Watch for fabricated testimonial. Measure evidence per claim.

Channel-aware story variants

Input: master storyboard and channel brief. Control: retain core story while changing length and visual grammar. Output: variant board. Accept when same claim, adapted format. Watch for message drift across channels. Measure message consistency score.

Prompt contract for ideation

Input: brand voice, audience, goal and exclusions. Control: specify role, inputs, output schema and refusal conditions. Output: ideation prompt. Accept when require structured idea records. Watch for generic idea spam. Measure schema-valid output rate.

Parallel ideation and merge

Input: three independent agent proposals. Control: deduplicate then rank without losing source provenance. Output: merged shortlist. Accept when retain evidence and dissent. Watch for majority vote on false claim. Measure duplicate and conflict count.

Human concept checkpoint

Input: shortlist and brand criteria. Control: approve, revise or reject with reasons. Output: decision record. Accept when human owns campaign direction. Watch for automatic campaign launch. Measure approval trace completeness.

Storyboard QA gate

Input: scene sheet and approval record. Control: test claim, continuity, CTA and accessibility. Output: storyboard QA report. Accept when block any critical defect. Watch for unchecked scene transition. Measure critical defect count.

CREATE prompt framework

Input: campaign role, request and examples. Control: specify Character, Request, Examples, Adjustments, Type and Extras. Output: structured content prompt. Accept when output must include a cited storyboard proposal. Watch for ambiguous prompt brief. Measure framework field coverage.

System, user and assistant roles

Input: approved policy and campaign request. Control: separate durable constraints from task input and model output. Output: role-scoped transcript. Accept when system constraints outrank untrusted source text. Watch for prompt-injection through source material. Measure role-boundary violations.

Lab 01: Campaign objective and role charter

Course: Multi AI Agents Workflow for Content Creation · TGS-2023036153 · v1.0

Outcome: LO1: Conceptualize content ideas to meet marketing objectives and map out digital storyboards as part of a content strategy.

Build: goal sheet, role charter and brief.

1. Open marketing-assets/brand-brief.md and record the campaign goal, approved claims and exclusions.
 2. Open mock-data/audience-signals.csv. Confirm every row is synthetic and do not substitute personal customer data.
 3. Open mock-data/claim-ledger.csv. Mark GL-04 as blocked because evidence is absent.
 4. Start a new workspace or chat in a tool available to you. Create four specialist roles using prompts.pdf: researcher, strategist, creator and reviewer. If the tool lacks true subagents, run the roles in separate chats and carry the structured handoff manually.
 5. Give the researcher only the local files needed. Save a source-linked research packet; include sample sizes and uncertainty.
 6. Pass the packet, campaign goal and brand constraints to the strategist. Ask for a decision record tied to the relevant learning outcome.
 7. Pass the approved brief and claim IDs to the creator. Generate the required artifact, then label it DRAFT with a version number.
 8. Give the draft and evidence packet to the reviewer. Record factual, brand, rights, privacy and accessibility findings. Revise only failed criteria, then repeat the review once.
 9. Ask a human acting as Marketing Manager to approve, revise or reject the final artifact. Record the decision, reason and time. Do not publish to a real channel during class.
- Deliverable exists: goal sheet, role charter and brief.
 - Every factual claim has a GL ID and no blocked claim appears in approved copy.
 - Agent handoffs name input, output, owner, status and source IDs.
 - At least one decision cites audience evidence rather than an invented preference.

- Reviewer records PASS, REVISE or BLOCK and human approval is distinct.

Save the brief, task ledger, prompt/output transcript, artifact, review result and human decision in an evidence/ folder named for this lab. Include one screenshot of the final approved state.

Lab 02: Research packet and idea ranking

Course: Multi AI Agents Workflow for Content Creation · TGS-2023036153 · v1.0

Outcome: LO1: Conceptualize content ideas to meet marketing objectives and map out digital storyboards as part of a content strategy.

Build: source cards, claim ledger and ranked ideas.

1. Open marketing-assets/brand-brief.md and record the campaign goal, approved claims and exclusions.
2. Open mock-data/audience-signals.csv. Confirm every row is synthetic and do not substitute personal customer data.
3. Open mock-data/claim-ledger.csv. Mark GL-04 as blocked because evidence is absent.
4. Start a new workspace or chat in a tool available to you. Create four specialist roles using prompts.pdf: researcher, strategist, creator and reviewer. If the tool lacks true subagents, run the roles in separate chats and carry the structured handoff manually.
5. Give the researcher only the local files needed. Save a source-linked research packet; include sample sizes and uncertainty.
6. Pass the packet, campaign goal and brand constraints to the strategist. Ask for a decision record tied to the relevant learning outcome.
7. Pass the approved brief and claim IDs to the creator. Generate the required artifact, then label it DRAFT with a version number.
8. Give the draft and evidence packet to the reviewer. Record factual, brand, rights, privacy and accessibility findings. Revise only failed criteria, then repeat the review once.
9. Ask a human acting as Marketing Manager to approve, revise or reject the final artifact. Record the decision, reason and time. Do not publish to a real channel during class.

- Deliverable exists: source cards, claim ledger and ranked ideas.
- Every factual claim has a GL ID and no blocked claim appears in approved copy.
- Agent handoffs name input, output, owner, status and source IDs.
- At least one decision cites audience evidence rather than an invented preference.
- Reviewer records PASS, REVISE or BLOCK and human approval is distinct.

Save the brief, task ledger, prompt/output transcript, artifact, review result and human decision in an evidence/ folder named for this lab. Include one screenshot of the final approved state.

Lab 03: Five-beat storyboard

Course: Multi AI Agents Workflow for Content Creation · TGS-2023036153 · v1.0

Outcome: LO1: Conceptualize content ideas to meet marketing objectives and map out digital storyboards as part of a content strategy.

Build: scene sheet, variants and QA decision.

1. Open marketing-assets/brand-brief.md and record the campaign goal, approved claims and exclusions.
 2. Open mock-data/audience-signals.csv. Confirm every row is synthetic and do not substitute personal customer data.
 3. Open mock-data/claim-ledger.csv. Mark GL-04 as blocked because evidence is absent.
 4. Start a new workspace or chat in a tool available to you. Create four specialist roles using prompts.pdf: researcher, strategist, creator and reviewer. If the tool lacks true subagents, run the roles in separate chats and carry the structured handoff manually.
 5. Give the researcher only the local files needed. Save a source-linked research packet; include sample sizes and uncertainty.
 6. Pass the packet, campaign goal and brand constraints to the strategist. Ask for a decision record tied to the relevant learning outcome.
 7. Pass the approved brief and claim IDs to the creator. Generate the required artifact, then label it DRAFT with a version number.
 8. Give the draft and evidence packet to the reviewer. Record factual, brand, rights, privacy and accessibility findings. Revise only failed criteria, then repeat the review once.
 9. Ask a human acting as Marketing Manager to approve, revise or reject the final artifact. Record the decision, reason and time. Do not publish to a real channel during class.
- Deliverable exists: scene sheet, variants and QA decision.
 - Every factual claim has a GL ID and no blocked claim appears in approved copy.
 - Agent handoffs name input, output, owner, status and source IDs.
 - At least one decision cites audience evidence rather than an invented preference.
 - Reviewer records PASS, REVISE or BLOCK and human approval is distinct.

Save the brief, task ledger, prompt/output transcript, artifact, review result and human decision in an evidence/ folder named for this lab. Include one screenshot of the final approved state.

Topic 2: Audience Research and Content Requirement Analysis

Preference evidence hierarchy

Input: interviews, analytics and CRM extracts. Control: rank direct behaviour above unsupported persona assumptions. Output: evidence register. Accept when label evidence age and sample size. Watch for treating guesses as facts. Measure source confidence.

Segment definition

Input: consented customer attributes. Control: cluster by need, intent and context. Output: segment cards. Accept when avoid sensitive proxy targeting. Watch for demographic stereotyping. Measure segment evidence coverage.

Persona jobs and objections

Input: segment cards and transcripts. Control: extract jobs-to-be-done, barriers and triggers. Output: persona matrix. Accept when keep verbatim evidence separate from inference. Watch for invented quotes. Measure quote traceability.

Journey stage mapping

Input: search queries and purchase path. Control: map problem, consideration, decision and retention needs. Output: journey grid. Accept when select one job per content asset. Watch for same message at every stage. Measure journey fit score.

Competitor content gap

Input: approved competitor URLs. Control: compare questions answered, proof and format. Output: gap table. Accept when seek unmet audience question. Watch for copying competitor wording. Measure originality and gap count.

Search intent classification

Input: keyword list with page evidence. Control: label informational, comparison and transactional intent. Output: intent map. Accept when match CTA to intent stage. Watch for sales CTA on learning query. Measure intent-CTA fit.

Content requirements schema

Input: persona, objective and channel constraints. Control: record length, voice, proof, format, CTA and exclusions. Output: content specification. Accept when reject missing mandatory fields. Watch for brief lost in handoff. Measure spec completeness.

Frequency from response data

Input: engagement, unsubscribe and capacity data. Control: compare marginal value with fatigue and production cost. Output: cadence proposal. Accept when reduce frequency when fatigue rises. Watch for over-posting. Measure engagement per send.

Editorial calendar capacity

Input: asset dependencies and staff capacity. Control: allocate publish slots and review buffers. Output: four-week calendar. Accept when never schedule before QA sign-off. Watch for unreviewed slot. Measure on-time approved rate.

Feedback loop design

Input: campaign performance and qualitative comments. Control: record hypotheses, tests and observed results. Output: learning log. Accept when change one variable per test. Watch for post-hoc storytelling. Measure test interpretability.

Research agent boundaries

Input: approved sources and data policy. Control: limit tools, time window and claims. Output: researcher charter. Accept when escalate inaccessible or conflicting evidence. Watch for web summary without source. Measure source coverage.

Audience analyst handoff

Input: research packet and segment matrix. Control: produce segment needs, confidence and exclusions. Output: audience brief. Accept when flag low-confidence segments. Watch for false precision. Measure confidence annotation rate.

RICCE audience brief

Input: role, intent, context, constraints and examples. Control: bind persona work to evidence rather than invented demographics. Output: RICCE brief. Accept when reject ungrounded preference claims. Watch for persona hallucination. Measure brief field completeness.

Few-shot preference extraction

Input: labelled customer comments and extraction schema. Control: show two positive and one ambiguous example. Output: structured preference records. Accept when retain ambiguous cases for human review. Watch for overfitting to examples. Measure label agreement.

Lab 04: Audience evidence and personas

Course: Multi AI Agents Workflow for Content Creation · TGS-2023036153 · v1.0

Outcome: LO2: Identify content requirements from customer preferences and determine the frequency of marketing content.

Build: evidence register, segments and persona matrix.

1. Open marketing-assets/brand-brief.md and record the campaign goal, approved claims and exclusions.
2. Open mock-data/audience-signals.csv. Confirm every row is synthetic and do not substitute personal customer data.
3. Open mock-data/claim-ledger.csv. Mark GL-04 as blocked because evidence is absent.
4. Start a new workspace or chat in a tool available to you. Create four specialist roles using prompts.pdf: researcher, strategist, creator and reviewer. If the tool lacks true subagents, run the roles in separate chats and carry the structured handoff manually.
5. Give the researcher only the local files needed. Save a source-linked research packet; include sample sizes and uncertainty.
6. Pass the packet, campaign goal and brand constraints to the strategist. Ask for a decision record tied to the relevant learning outcome.

7. Pass the approved brief and claim IDs to the creator. Generate the required artifact, then label it DRAFT with a version number.
 8. Give the draft and evidence packet to the reviewer. Record factual, brand, rights, privacy and accessibility findings. Revise only failed criteria, then repeat the review once.
 9. Ask a human acting as Marketing Manager to approve, revise or reject the final artifact. Record the decision, reason and time. Do not publish to a real channel during class.
- Deliverable exists: evidence register, segments and persona matrix.
 - Every factual claim has a GL ID and no blocked claim appears in approved copy.
 - Agent handoffs name input, output, owner, status and source IDs.
 - At least one decision cites audience evidence rather than an invented preference.
 - Reviewer records PASS, REVISE or BLOCK and human approval is distinct.

Save the brief, task ledger, prompt/output transcript, artifact, review result and human decision in an evidence/ folder named for this lab. Include one screenshot of the final approved state.

Lab 05: Requirements and cadence

Course: Multi AI Agents Workflow for Content Creation · TGS-2023036153 · v1.0

Outcome: LO2: Identify content requirements from customer preferences and determine the frequency of marketing content.

Build: content specification and four-week calendar.

1. Open marketing-assets/brand-brief.md and record the campaign goal, approved claims and exclusions.
2. Open mock-data/audience-signals.csv. Confirm every row is synthetic and do not substitute personal customer data.
3. Open mock-data/claim-ledger.csv. Mark GL-04 as blocked because evidence is absent.
4. Start a new workspace or chat in a tool available to you. Create four specialist roles using prompts.pdf: researcher, strategist, creator and reviewer. If the tool lacks true subagents, run the roles in separate chats and carry the structured handoff manually.
5. Give the researcher only the local files needed. Save a source-linked research packet; include sample sizes and uncertainty.
6. Pass the packet, campaign goal and brand constraints to the strategist. Ask for a decision record tied to the relevant learning outcome.
7. Pass the approved brief and claim IDs to the creator. Generate the required artifact, then label it DRAFT with a version number.
8. Give the draft and evidence packet to the reviewer. Record factual, brand, rights, privacy and accessibility findings. Revise only failed criteria, then repeat the review once.

9. Ask a human acting as Marketing Manager to approve, revise or reject the final artifact. Record the decision, reason and time. Do not publish to a real channel during class.

- Deliverable exists: content specification and four-week calendar.
- Every factual claim has a GL ID and no blocked claim appears in approved copy.
- Agent handoffs name input, output, owner, status and source IDs.
- At least one decision cites audience evidence rather than an invented preference.
- Reviewer records PASS, REVISE or BLOCK and human approval is distinct.

Save the brief, task ledger, prompt/output transcript, artifact, review result and human decision in an evidence/ folder named for this lab. Include one screenshot of the final approved state.

Lab 06: Research-agent handoff

Course: Multi AI Agents Workflow for Content Creation · TGS-2023036153 · v1.0

Outcome: LO2: Identify content requirements from customer preferences and determine the frequency of marketing content.

Build: validated audience brief with uncertainty labels.

1. Open marketing-assets/brand-brief.md and record the campaign goal, approved claims and exclusions.
 2. Open mock-data/audience-signals.csv. Confirm every row is synthetic and do not substitute personal customer data.
 3. Open mock-data/claim-ledger.csv. Mark GL-04 as blocked because evidence is absent.
 4. Start a new workspace or chat in a tool available to you. Create four specialist roles using prompts.pdf: researcher, strategist, creator and reviewer. If the tool lacks true subagents, run the roles in separate chats and carry the structured handoff manually.
 5. Give the researcher only the local files needed. Save a source-linked research packet; include sample sizes and uncertainty.
 6. Pass the packet, campaign goal and brand constraints to the strategist. Ask for a decision record tied to the relevant learning outcome.
 7. Pass the approved brief and claim IDs to the creator. Generate the required artifact, then label it DRAFT with a version number.
 8. Give the draft and evidence packet to the reviewer. Record factual, brand, rights, privacy and accessibility findings. Revise only failed criteria, then repeat the review once.
 9. Ask a human acting as Marketing Manager to approve, revise or reject the final artifact. Record the decision, reason and time. Do not publish to a real channel during class.
- Deliverable exists: validated audience brief with uncertainty labels.

- Every factual claim has a GL ID and no blocked claim appears in approved copy.
- Agent handoffs name input, output, owner, status and source IDs.
- At least one decision cites audience evidence rather than an invented preference.
- Reviewer records PASS, REVISE or BLOCK and human approval is distinct.

Save the brief, task ledger, prompt/output transcript, artifact, review result and human decision in an evidence/ folder named for this lab. Include one screenshot of the final approved state.

Topic 3: Multi-Channel Content Creation and Agent Workflow Coordination

Orchestration state machine

Input: approved brief and task queue. Control: move tasks through planned, drafting, review, approved and blocked. Output: task ledger. Accept when publish only from approved state. Watch for lost handoff. Measure state transition validity.

Handoff JSON contract

Input: agent output and schema version. Control: validate fields, sources, confidence and owner. Output: typed handoff packet. Accept when reject malformed payload. Watch for free-text ambiguity. Measure schema pass rate.

Researcher-writer boundary

Input: verified claim ledger. Control: pass evidence IDs and prohibited claims to writer. Output: drafting packet. Accept when writer may not invent citations. Watch for citation fabrication. Measure claim-to-source match.

Writer-editor loop

Input: draft and brand rubric. Control: revise only failed criteria with change log. Output: edited draft. Accept when cap iterations and escalate. Watch for infinite rewrite loop. Measure revision gain per cycle.

Fact-check agent

Input: draft claims and source snapshots. Control: classify supported, disputed or unverified. Output: fact-check matrix. Accept when block disputed high-impact claims. Watch for plausible falsehood. Measure verified claim share.

Brand voice gate

Input: approved voice examples and exclusions. Control: score tone, terminology and promise consistency. Output: brand review. Accept when human resolves borderline calls. Watch for robotic generic voice. Measure rubric agreement.

SEO brief to article

Input: search intent, keyword cluster and source pack. Control: structure answer, headings and internal link opportunities. Output: SEO draft. Accept when prioritize reader task over keyword repetition. Watch for keyword stuffing. Measure intent coverage.

Email adaptation

Input: master narrative and consented segment. Control: shorten to subject, preheader, body and single CTA. Output: email variant. Accept when check preference and unsubscribe handling. Watch for non-consented send. Measure CTA clarity.

Social adaptation

Input: master claim and platform constraints. Control: choose hook, visual, caption and link treatment. Output: social post pack. Accept when validate current platform limits before publishing. Watch for truncated post. Measure format validity.

Visual asset brief

Input: storyboard, brand kit and rights register. Control: specify subject, crop, alt text and licensing. Output: designer brief. Accept when reject imagery implying false evidence. Watch for misleading generated photo. Measure rights record coverage.

Multimodal consistency

Input: copy, image and video script. Control: cross-check claim, depiction and accessibility. Output: multimodal QA sheet. Accept when same offer and disclaimer across formats. Watch for visual-text contradiction. Measure cross-format consistency.

Error and retry policy

Input: agent failure log and task state. Control: retry transient faults; block validation/auth failures. Output: recovery record. Accept when bounded retries with human escalation. Watch for silent partial publish. Measure recoverable task rate.

Tool routing by role

Input: approved tool inventory and data class. Control: route research, drafting, image and review tasks to bounded tools. Output: tool assignment matrix. Accept when no private data sent to a public tool. Watch for tool sprawl and data leak. Measure routing compliance.

Self-critique editor loop

Input: draft, rubric and source ledger. Control: identify a specific defect then revise against evidence. Output: revision diff. Accept when stop after two loops and escalate. Watch for endless cosmetic rewrites. Measure defect closure rate.

Lab 07: Agent workflow state machine

Course: Multi AI Agents Workflow for Content Creation · TGS-2023036153 · v1.0

Outcome: LO3: Determine content types and styles and decide on modes and processes for distributing content.

Build: task ledger and typed handoff.

1. Open marketing-assets/brand-brief.md and record the campaign goal, approved claims and exclusions.
2. Open mock-data/audience-signals.csv. Confirm every row is synthetic and do not substitute personal customer data.

3. Open mock-data/claim-ledger.csv. Mark GL-04 as blocked because evidence is absent.
 4. Start a new workspace or chat in a tool available to you. Create four specialist roles using prompts.pdf: researcher, strategist, creator and reviewer. If the tool lacks true subagents, run the roles in separate chats and carry the structured handoff manually.
 5. Give the researcher only the local files needed. Save a source-linked research packet; include sample sizes and uncertainty.
 6. Pass the packet, campaign goal and brand constraints to the strategist. Ask for a decision record tied to the relevant learning outcome.
 7. Pass the approved brief and claim IDs to the creator. Generate the required artifact, then label it DRAFT with a version number.
 8. Give the draft and evidence packet to the reviewer. Record factual, brand, rights, privacy and accessibility findings. Revise only failed criteria, then repeat the review once.
 9. Ask a human acting as Marketing Manager to approve, revise or reject the final artifact. Record the decision, reason and time. Do not publish to a real channel during class.
- Deliverable exists: task ledger and typed handoff.
 - Every factual claim has a GL ID and no blocked claim appears in approved copy.
 - Agent handoffs name input, output, owner, status and source IDs.
 - At least one decision cites audience evidence rather than an invented preference.
 - Reviewer records PASS, REVISE or BLOCK and human approval is distinct.

Save the brief, task ledger, prompt/output transcript, artifact, review result and human decision in an evidence/ folder named for this lab. Include one screenshot of the final approved state.

Lab 08: Draft, edit and fact-check

Course: Multi AI Agents Workflow for Content Creation · TGS-2023036153 · v1.0

Outcome: LO3: Determine content types and styles and decide on modes and processes for distributing content.

Build: source-linked article and review log.

1. Open marketing-assets/brand-brief.md and record the campaign goal, approved claims and exclusions.
2. Open mock-data/audience-signals.csv. Confirm every row is synthetic and do not substitute personal customer data.
3. Open mock-data/claim-ledger.csv. Mark GL-04 as blocked because evidence is absent.
4. Start a new workspace or chat in a tool available to you. Create four specialist roles using prompts.pdf: researcher, strategist, creator and reviewer. If the tool lacks true subagents, run the roles in separate chats and carry the structured handoff manually.

5. Give the researcher only the local files needed. Save a source-linked research packet; include sample sizes and uncertainty.
 6. Pass the packet, campaign goal and brand constraints to the strategist. Ask for a decision record tied to the relevant learning outcome.
 7. Pass the approved brief and claim IDs to the creator. Generate the required artifact, then label it DRAFT with a version number.
 8. Give the draft and evidence packet to the reviewer. Record factual, brand, rights, privacy and accessibility findings. Revise only failed criteria, then repeat the review once.
 9. Ask a human acting as Marketing Manager to approve, revise or reject the final artifact. Record the decision, reason and time. Do not publish to a real channel during class.
- Deliverable exists: source-linked article and review log.
 - Every factual claim has a GL ID and no blocked claim appears in approved copy.
 - Agent handoffs name input, output, owner, status and source IDs.
 - At least one decision cites audience evidence rather than an invented preference.
 - Reviewer records PASS, REVISE or BLOCK and human approval is distinct.

Save the brief, task ledger, prompt/output transcript, artifact, review result and human decision in an evidence/ folder named for this lab. Include one screenshot of the final approved state.

Lab 09: Multi-channel asset pack

Course: Multi AI Agents Workflow for Content Creation · TGS-2023036153 · v1.0

Outcome: LO3: Determine content types and styles and decide on modes and processes for distributing content.

Build: email, social, visual brief and multimodal QA.

1. Open marketing-assets/brand-brief.md and record the campaign goal, approved claims and exclusions.
2. Open mock-data/audience-signals.csv. Confirm every row is synthetic and do not substitute personal customer data.
3. Open mock-data/claim-ledger.csv. Mark GL-04 as blocked because evidence is absent.
4. Start a new workspace or chat in a tool available to you. Create four specialist roles using prompts.pdf: researcher, strategist, creator and reviewer. If the tool lacks true subagents, run the roles in separate chats and carry the structured handoff manually.
5. Give the researcher only the local files needed. Save a source-linked research packet; include sample sizes and uncertainty.
6. Pass the packet, campaign goal and brand constraints to the strategist. Ask for a decision record tied to the relevant learning outcome.

7. Pass the approved brief and claim IDs to the creator. Generate the required artifact, then label it DRAFT with a version number.

8. Give the draft and evidence packet to the reviewer. Record factual, brand, rights, privacy and accessibility findings. Revise only failed criteria, then repeat the review once.

9. Ask a human acting as Marketing Manager to approve, revise or reject the final artifact. Record the decision, reason and time. Do not publish to a real channel during class.

- Deliverable exists: email, social, visual brief and multimodal QA.
- Every factual claim has a GL ID and no blocked claim appears in approved copy.
- Agent handoffs name input, output, owner, status and source IDs.
- At least one decision cites audience evidence rather than an invented preference.
- Reviewer records PASS, REVISE or BLOCK and human approval is distinct.

Save the brief, task ledger, prompt/output transcript, artifact, review result and human decision in an evidence/ folder named for this lab. Include one screenshot of the final approved state.

Topic 4: Content Distribution, Strategy Guidelines and Responsible AI Practices

Channel decision matrix

Input: segment preference and content specs. Control: score reach, fit, effort, privacy and measurability. Output: channel plan. Accept when pick channels with observable objective. Watch for channel sprawl. Measure weighted fit score.

Distribution dependency graph

Input: approved assets and calendar. Control: link review, legal, creative and publishing prerequisites. Output: release graph. Accept when no node runs before dependencies pass. Watch for premature release. Measure blocked dependency count.

Publish approval gate

Input: QA packet and named approver. Control: record decision, version, time and exceptions. Output: release authorization. Accept when human approval required for external publish. Watch for agent self-approval. Measure approval audit completeness.

Scheduling and idempotency

Input: approved post IDs and channel slots. Control: use unique campaign key and retry-safe publish action. Output: schedule ledger. Accept when duplicate key prevents duplicate post. Watch for double publication. Measure duplicate rate.

Localization and accessibility

Input: approved master and audience locale. Control: check language, captions, alt text and contrast. Output: accessible variant. Accept when retain claim and CTA meaning. Watch for untranslated disclaimer. Measure accessibility pass rate.

Copyright and asset provenance

Input: source files and license notes. Control: record creator, permission, allowed use and attribution. Output: rights ledger. Accept when hold asset with unknown rights. Watch for unlicensed image. Measure rights coverage.

Privacy-safe customer data

Input: CRM extract and consent register. Control: minimize, pseudonymize and restrict agent context. Output: safe audience dataset. Accept when exclude identifiers not needed for task. Watch for PII in prompt. Measure PII leakage count.

Synthetic content disclosure

Input: generated image and campaign context. Control: decide disclosure and avoid deceptive realism. Output: disclosure record. Accept when escalate misleading depictions. Watch for fake testimonial. Measure disclosure decision trace.

Claim and hallucination control

Input: draft and source ledger. Control: require source-backed factual claims and abstention. Output: claim QA report. Accept when no publication with unverified claim. Watch for confident fabricated statistic. Measure unsupported claim count.

Operational observability

Input: task ledger, token costs and errors. Control: measure latency, cost, retries and QA outcome. Output: run dashboard. Accept when pause deteriorating workflow. Watch for hidden failure. Measure cost per approved asset.

Experiment and learning plan

Input: baseline campaign and hypothesis. Control: define variant, metric, guardrail and review date. Output: test plan. Accept when stop test on safety breach. Watch for optimizing vanity metric. Measure valid experiment share.

Strategy governance playbook

Input: all approved contracts and logs. Control: set roles, cadence, escalation and revision control. Output: execution guideline. Accept when quarterly human review. Watch for orphaned workflow. Measure policy review completion.

Provider data controls

Input: approved data classification and current tool settings. Control: choose retention and training controls before upload. Output: data-control record. Accept when block confidential inputs on unapproved accounts. Watch for private data in chat history. Measure control attestation.

Synthetic-media authenticity

Input: generated photo or voice and usage context. Control: check consent, provenance, disclosure and deceptive impression. Output: media authenticity record. Accept when hold impersonation or false documentary depiction. Watch for deepfake-style deception. Measure authenticity pass rate.

Lab 10: Distribution release gate

Course: Multi AI Agents Workflow for Content Creation · TGS-2023036153 · v1.0

Outcome: LO4: Develop guidelines for content strategy execution using appropriate delivery modes and responsible AI practices.

Build: channel matrix, dependencies and approval record.

1. Open marketing-assets/brand-brief.md and record the campaign goal, approved claims and exclusions.
 2. Open mock-data/audience-signals.csv. Confirm every row is synthetic and do not substitute personal customer data.
 3. Open mock-data/claim-ledger.csv. Mark GL-04 as blocked because evidence is absent.
 4. Start a new workspace or chat in a tool available to you. Create four specialist roles using prompts.pdf: researcher, strategist, creator and reviewer. If the tool lacks true subagents, run the roles in separate chats and carry the structured handoff manually.
 5. Give the researcher only the local files needed. Save a source-linked research packet; include sample sizes and uncertainty.
 6. Pass the packet, campaign goal and brand constraints to the strategist. Ask for a decision record tied to the relevant learning outcome.
 7. Pass the approved brief and claim IDs to the creator. Generate the required artifact, then label it DRAFT with a version number.
 8. Give the draft and evidence packet to the reviewer. Record factual, brand, rights, privacy and accessibility findings. Revise only failed criteria, then repeat the review once.
 9. Ask a human acting as Marketing Manager to approve, revise or reject the final artifact. Record the decision, reason and time. Do not publish to a real channel during class.
- Deliverable exists: channel matrix, dependencies and approval record.
 - Every factual claim has a GL ID and no blocked claim appears in approved copy.
 - Agent handoffs name input, output, owner, status and source IDs.
 - At least one decision cites audience evidence rather than an invented preference.
 - Reviewer records PASS, REVISE or BLOCK and human approval is distinct.

Save the brief, task ledger, prompt/output transcript, artifact, review result and human decision in an evidence/ folder named for this lab. Include one screenshot of the final approved state.

Lab 11: Rights, privacy and disclosure

Course: Multi AI Agents Workflow for Content Creation · TGS-2023036153 · v1.0

Outcome: LO4: Develop guidelines for content strategy execution using appropriate delivery modes and responsible AI practices.

Build: rights ledger, safe data and disclosure decision.

1. Open marketing-assets/brand-brief.md and record the campaign goal, approved claims and exclusions.
 2. Open mock-data/audience-signals.csv. Confirm every row is synthetic and do not substitute personal customer data.
 3. Open mock-data/claim-ledger.csv. Mark GL-04 as blocked because evidence is absent.
 4. Start a new workspace or chat in a tool available to you. Create four specialist roles using prompts.pdf: researcher, strategist, creator and reviewer. If the tool lacks true subagents, run the roles in separate chats and carry the structured handoff manually.
 5. Give the researcher only the local files needed. Save a source-linked research packet; include sample sizes and uncertainty.
 6. Pass the packet, campaign goal and brand constraints to the strategist. Ask for a decision record tied to the relevant learning outcome.
 7. Pass the approved brief and claim IDs to the creator. Generate the required artifact, then label it DRAFT with a version number.
 8. Give the draft and evidence packet to the reviewer. Record factual, brand, rights, privacy and accessibility findings. Revise only failed criteria, then repeat the review once.
 9. Ask a human acting as Marketing Manager to approve, revise or reject the final artifact. Record the decision, reason and time. Do not publish to a real channel during class.
- Deliverable exists: rights ledger, safe data and disclosure decision.
 - Every factual claim has a GL ID and no blocked claim appears in approved copy.
 - Agent handoffs name input, output, owner, status and source IDs.
 - At least one decision cites audience evidence rather than an invented preference.
 - Reviewer records PASS, REVISE or BLOCK and human approval is distinct.

Save the brief, task ledger, prompt/output transcript, artifact, review result and human decision in an evidence/ folder named for this lab. Include one screenshot of the final approved state.

Lab 12: Run evaluation and governance

Course: Multi AI Agents Workflow for Content Creation · TGS-2023036153 · v1.0

Outcome: LO4: Develop guidelines for content strategy execution using appropriate delivery modes and responsible AI practices.

Build: dashboard, experiment plan and execution playbook.

1. Open marketing-assets/brand-brief.md and record the campaign goal, approved claims and exclusions.
2. Open mock-data/audience-signals.csv. Confirm every row is synthetic and do not substitute personal customer data.

3. Open mock-data/claim-ledger.csv. Mark GL-04 as blocked because evidence is absent.
 4. Start a new workspace or chat in a tool available to you. Create four specialist roles using prompts.pdf: researcher, strategist, creator and reviewer. If the tool lacks true subagents, run the roles in separate chats and carry the structured handoff manually.
 5. Give the researcher only the local files needed. Save a source-linked research packet; include sample sizes and uncertainty.
 6. Pass the packet, campaign goal and brand constraints to the strategist. Ask for a decision record tied to the relevant learning outcome.
 7. Pass the approved brief and claim IDs to the creator. Generate the required artifact, then label it DRAFT with a version number.
 8. Give the draft and evidence packet to the reviewer. Record factual, brand, rights, privacy and accessibility findings. Revise only failed criteria, then repeat the review once.
 9. Ask a human acting as Marketing Manager to approve, revise or reject the final artifact. Record the decision, reason and time. Do not publish to a real channel during class.
- Deliverable exists: dashboard, experiment plan and execution playbook.
 - Every factual claim has a GL ID and no blocked claim appears in approved copy.
 - Agent handoffs name input, output, owner, status and source IDs.
 - At least one decision cites audience evidence rather than an invented preference.
 - Reviewer records PASS, REVISE or BLOCK and human approval is distinct.

Save the brief, task ledger, prompt/output transcript, artifact, review result and human decision in an evidence/ folder named for this lab. Include one screenshot of the final approved state.

Sources and further reading

- <https://www.tertiarycourses.com.sg/multi-ai-agents-workflow-for-content-creation.html>
- <https://techcommunity.microsoft.com/blog/educatordeveloperblog/creating-a-fun-multi-agent-content-strategy-system-with-microsoft-agent-framework/4495105>
- <https://www.mindstudio.ai/blog/ai-content-creation-sub-agents-research-to-post>
- <https://n8n.io/workflows/10293-multi-agent-ai-content-creator-for-seo-blogs-and-newsletters-with-openrouter-dall-e-gemini/>
- https://dev.to/pavel_polivka/building-a-multi-agent-content-management-system-with-ai-29i7
- Multi-Agent Development with Claude Code: Subagents, Team Orchestration, and Long-Running Systems for Autonomous AI Development (reference ebook)