

WSQ COURSE • TGS-2023036153 • CHATGPT EDITION

Version 3.5 • 10 Oct 2026

Multi AI Agents Workflow for Content Creation

Tertiary Infotech Academy Pte Ltd
UEN: 201200696W

www.tertiarycourses.com.sg | enquiry@tertiaryinfotech.com | +65 6100 0613

© Tertiary Infotech Academy Pte Ltd

Digital Attendance



Attendance is mandatory for SSG funding.

Take AM, PM and Assessment digital attendance for every session.

Use the SSG / TMS QR code provided by your trainer.

**Minimum
75%
attendance
required**

About the Trainer



[Trainer Name]

[Title / Role]



Qualifications: _____



Expertise: _____



Experience: _____



Contact / Profile: _____

About the Trainer



Dr. Alfred Ang

AI Trainer & Engineer



PhD-qualified trainer specialising in AI, software engineering and developer tooling.



Builds AI marketing teams of agents with ChatGPT and Codex.



Trains professionals across Singapore on practical, project-based AI skills.



Certifications: PMP, AWS, Microsoft, CompTIA, CKA, CKAD.



[linkedin.com/in/angchewhoe](https://www.linkedin.com/in/angchewhoe)



github.com/alfredang

INTRODUCTION

Let's Know Each Other

Take a minute to introduce yourself to the class:



Your name & role



What you do day to day



Your content work today



What you want from this
course

Ground Rules



Phones on silent

Step out quietly for calls or breaks.



Participate actively

No question is stupid — ask freely.



Mutual respect

Agree to disagree; one conversation at a time.



Be punctual

Return from breaks on time.



75% attendance

Required for funding eligibility.



Stay hands-on

Follow along with every activity.

Briefing for Assessment



Clear your space

Place phones and other materials under the table or on the floor.



No recording

No photos or recording of assessment materials.



Open book

Only slides, learner guide and approved materials are allowed.



Work individually

Complete the assessment on your own; no discussion.



Manage your time

Written Assessment then Practical Performance — watch the clock.



Appeals

An appeal process is available if you disagree with the outcome.

Assessment & Funding



Final Assessment

- Written Assessment (WA) — 1 hour, starts 4:00 PM
- Practical Performance (PP) — 1 hour
- Format: Open book — class ends 6:00 PM
- Allowed: slides, learner guide & approved materials only
- Appeal process available



Criteria for Funding

Minimum 75% attendance

Based on SSG digital attendance records.

Assessed as 'Competent'

Complete and pass both assessment components.

Courseware & Assessment on the LMS



Download the courseware

Get the slides and learner guide from the LMS for the open-book assessment.



Take the assessment

Complete the Written Assessment and Practical Performance on the LMS.



Log in

Use your enrolled account to access the course materials and assessment.

A screenshot of the Tertiary Infotech Academy LMS login interface. It features a dark blue background with a white logo at the top. Below the logo is the text 'Tertiary Infotech Academy' and 'LMS cum TMS for WSQ Courses'. There is an 'Email' input field, a 'Send OTP' button, and a link 'Login with password instead'. At the bottom, there are links for 'Privacy Policy', 'Acceptable Use Policy', and 'Feedback', and a note 'Powered by Tertiary Infotech Academy Pte Ltd'.

LMS: <https://lms-tms.tertiaryinfotech.com/>

Certification & TRAQOM Survey



TRAQOM Survey (Mandatory)

Complete the certification & TRAQOM survey to receive your certificate.

<https://lms-tms.tertiaryinfotech.com/>



Digital Attendance

Remember to take AM, PM and Assessment digital attendance — required for your funding.

Learning Outcomes

By the end of the two days you will be able to:

L01 · Conceptualise

Conceptualise content ideas that meet marketing objectives, and map them into digital storyboards.

L02 · Identify

Identify content requirements from customer preferences, and decide how often to publish.

L03 · Determine


Determine content types and styles, and the modes and processes for distributing content.

L04 · Develop

Develop guidelines for executing the content strategy, with suitable delivery modes and responsible AI practices.

Course Outline

Four topics over two days. Day 1 forms the team; Day 2 runs its campaigns.



Lab Materials

Fifteen labs, each in its own folder with a README, prompts (MD and PDF), assets, an evidence checklist and, where useful, a solution.

DAY	LABS	TOPIC
1	1 Agent team • 2 Instructions • 3 Skills • 4 Connectors	T1 • form the team
1	5 Campaign 1: research to storyboard	T1 • Campaign 1
1	6 Personas • 7 Cadence and calendar • 8 Blog article	T2 • audience, content
2	9 Lead magnet • 10 Social week • 11 Newsletter + landing page • 12 Video	T3 • Campaigns 2–4
2	13 Weekly report • 14 Playbook • 15 Optional demo	T4 • Campaign 5

TIP One horizon-studio folder from Lab 1 to Lab 14 — the team keeps what it learns.

ChatGPT or Codex?

Every lab says which to use. The rule: ChatGPT to create and schedule; Codex when the agents work on files.

USE	LABS	WHAT IT IS FOR
ChatGPT	8, 9, 15	One person creating or designing.
Codex	1, 2, 5, 6, 7, 10, 12, 14	The team's files: role cards, skills, subagents, scripts, the video.
Codex, then ChatGPT	3	Build in Codex, then share in ChatGPT.
ChatGPT, then Codex	4	Install plugins, then test the agents.
Codex, then ChatGPT Work	11	Build in Codex; Gmail drafts in ChatGPT Work.
ChatGPT Work	13	Scheduled tasks and plugins.

TIP Both work on the same horizon-studio folder, so whatever one makes, the other can open.

SCHEDULE

Lesson Plan (9:00 AM – 6:00 PM)

Tea breaks 10 min, lunch 45 min. Full timings and slide numbers are in the Lesson Plan.

BLOCK	DAY 1 — FORM THE TEAM	DAY 2 — RUN THE CAMPAIGNS
Morning 1	9:00–10:30 Welcome • Topic 1 • Lab 1	9:00–10:40 Topic 3 • Labs 9–10
Morning 2	10:40–12:40 Labs 2–3	10:50–12:35 Labs 11–12 • recap
Afternoon 1	13:25–15:20 Labs 4–5	13:20–15:00 Topic 4 • Labs 13–14
Afternoon 2	15:30–18:00 Topic 2 • Labs 6–8 • review	15:10–16:00 Demo • summary • 16:00–18:00 Assessment

TIP Day 2 assessment: Written Assessment 4:00–5:00 PM, Practical Performance 5:00–6:00 PM.

TOPIC 1

Multi-AI-Agent Content Ideation and Digital Storyboarding

- 1 AI agents, the tools, and Horizon's story
- 2 Why a team of agents
- 3 A home for the team; Horizon's public website
- 4 Form the team: instructions, skills, connectors
- 5 Campaign 1: research, ideas and a storyboard

ChatGPT, ChatGPT Work and Codex

One desktop app. The toggle above the composer picks Chat or Work; the sidebar holds Codex.

Codex — the team's home

The Lead's chat; role cards; skills in .agents/skills/; subagents by asking; scripts and the video.

ChatGPT and Work — the team's apps

A Project for shared files; plugins (Drive, Gmail); skills; scheduled tasks; @Sites.

Dots and workspace agents • optional

An always-on Lead (Pro, Business Premium) and shared agents (Business). A demo at the end.


Which Plan Has What

Features by ChatGPT plan, checked 8 Oct 2026. Prices change — check chatgpt.com/pricing.

FEATURE	PLANS	IN THIS COURSE
ChatGPT Work, Codex	Plus, Pro, Business, Enterprise	Every lab
Sites	Plus, Pro and workspaces (beta)	Lab 11
Plugins and skills	Paid plans	Labs 3, 4, 13
Scheduled tasks	Paid plans — 5 active on Plus, 15 on Pro	Lab 13
Dots (first one included)	Pro, Business Premium	Lab 15 demo
Workspace agents	Business, Enterprise, Edu	Lab 15 demo

TIP This edition builds the team with skills and Codex subagents, so it works on Plus and Pro.

A Brief History of AI, 2023–2026



From wording one prompt, to curating context, to engineering the whole system around the model — and now, teams of agents.

What Is a Harness?

The model is the brain. The harness is everything built around it that turns a model into an agent able to finish long tasks.

The agentic loop

Plans, acts, checks the result and repeats until the goal is met.

Tools and context

Files, commands, search and connectors (MCP); instructions, memory and skills loaded at the right moment.

Guardrails and checks


Permission modes, approvals, hooks and tests — what it must ask before it acts.

Harnesses you will meet

Codex and ChatGPT Work in class; also Claude Code, Gemini CLI, Hermes Agent and OpenClaw.

The Agentic Loop

Every agent runs the same repeating cycle. A team is several loops, side by side.



TIP Goal in → context → action → verification → repeat → goal met.

Why the Loop Matters

Running a team of agents is steering loops, not typing commands.

Give direction, not micro-steps

State the goal, the inputs and "done when". Let each loop work out the steps.

Verification is everything

An agent is only as good as its checks: a facts sheet, a test, the Lead's review.

Steer mid-loop

Read the plan, interrupt, correct the course — do not wait for a wrong final answer.

Context is the fuel

What an agent can see decides what it can do: the brief, the data, the skills.

Agentic AI vs AI Agents


Both use tools, skills and memory. The difference is how long they live and what starts them.

FEATURE	AGENTIC AI — TASK-ORIENTED	AI AGENTS — ALWAYS ON, 24/7
Tool calling	Yes — files, search, connectors, subagents.	Yes — plus chat apps, the browser and other agents.
Skills	Yes — loads SKILL.md procedures when the task matches.	Yes — and can write and improve its own skills.
Memory	Within the task, plus project memory across sessions.	Long-term memory about you and your work.
Availability	Starts when you give it a task; stops when done.	Runs 24/7 on a computer or server.
Started by	You, with a task.	You, a schedule, a message or another agent.
In this course	Codex running a campaign (Labs 1–12).	The scheduled weekly report (Lab 13); Dots (Lab 15).

TIP Rule of thumb: a job with an end → agentic AI; a job that recurs → an always-on agent.


What Is an Agent Made Of?

A model on its own can only answer. Wrap it with these parts and it can act — and check what it did.




Build a Marketing Team, Then Run Campaigns

Horizon Wealth Planning has a website, a goal of 40 booked chats a month — and no marketing team. You build one, made of AI agents.



Meet Horizon Wealth Planning

The firm's public website — the same link for every learner. You do not build it; every campaign points back here.



Plain English

Retirement, CPF, insurance and investing — no jargon, no hard sell.

One goal

Turn visitors into booked free 30-minute chats.


Regulated

Financial content: general information only, never a promised return.

WEBSITE tertiarycourses.github.io/TGS-2023036153-Multi-AI-Agents-Workflow-for-Content-Creation

Your AI Marketing Team

One Lead and five specialists — each with instructions, its own skills and connectors.



Who Does What

Each agent's job, the skills it owns and the connectors it may use.

AGENT	JOB	SKILLS	CONNECTORS
Team Lead	Plans, assigns, checks, hands to a person	fact-check, fin-compliance	@Google Drive (read), @Gmail (drafts only)
Growth Strategist	Research, positioning, briefs, storyboards	competitor-scan, campaign-brief, content-marketing	web search, @Google Drive (read)
Content Creator	Posts, newsletter, scripts, page copy	sunny-voice, copywriting, channel-formats, linkedin-post, facebook-post, blog-post, newsletter	@Google Drive (docs), @Gmail (drafts only)
Creative Designer	Images, thumbnails, the video	brand-visuals, video-render	\$imagegen, ffmpeg in Codex
Website Designer	Landing pages, published and checked	web-design, landing-page, lead-magnet, page-qa	@Sites (publish), @Computer Use (browser check)
Growth Analyst	Personas, cadence, results, reports	audience-insights, campaign-report	@Google Drive and Sheets (read), Python in Codex

TIP You build every part of this table in Labs 2, 3 and 4.

The Tools for This Edition

The ChatGPT Edition: ChatGPT, ChatGPT Work and Codex in one app.

TOOL	WHAT THE TEAM USES IT FOR	LABS
ChatGPT Project	The team's shared files and instructions	Lab 1
Codex	The Lead's chat; role cards; skills; subagents; scripts; the video	Labs 1–12, 14
ChatGPT Work	Plugins (Drive, Gmail), skills, scheduled tasks	Labs 4, 11, 13
@Sites	The campaign landing page; Horizon's website is given as a public link	Lab 11
Dots • workspace agents	Optional demo: an always-on Lead and shared agents	Lab 15 (optional)

TIP Every external post, email and video passes one gate: a person approves the exact version.


Why a Team of Agents?

The Lead runs specialists in parallel, each in its own context, then combines what they return.

ONE AGENT DOES EVERYTHING


One agent

One context window,
filling up with every job



One job after another • one voice • marks its own homework

A MAIN AGENT DELEGATES



Faster

Research, copy and design run at the same time.

Focused

Each agent holds only its own job and its own skills.

Specialised

Its own instructions, tools, connectors and model.

Second opinion

The Lead checks work it did not write.

The cost

More tokens and more to review. Split only work that is truly independent.

Get the Tools Ready

Install once; every lab says which tool to open.

DOWNLOAD

chatgpt.com/download

macOS

Download for macOS from the page

Windows

Opens the Microsoft Store listing

Work, Codex and Sites need a paid plan.

1

Download and install The new desktop app from chatgpt.com/download, not "ChatGPT Classic".

2

Sign in On the plan you will use in class — limits follow the account.

3

Find Chat, Work and Codex The toggle above the composer; Codex in the sidebar.

4

Node.js 22, Python 3, ffmpeg For the publisher kit and the video (Pillow for Python).

5

A personal Google account For Drive and Gmail — never an employer's.

A Home for the Team

Before you hire agents, give them shared rules, shared facts and a website to point to.

A Project + AGENTS.md

Project files and instructions in ChatGPT; the charter in AGENTS.md for Codex.

data/

The facts sheet, brand, compliance checklist and business brief — the only sources agents may use.

Horizon's public website

Given, not built: each learner opens their own local copy; the public copy is on GitHub Pages.

Build Your Marketing Agent Team

LAB 01

DAY 1 • ~30 MIN

USE: CODEX

THE STORY SO FAR

October 2026. Rachel has a website and a goal: 40 booked chats a month by March. She has no marketing team. Before you hire one made of agents, give it a home — shared rules, shared facts — and open your own copy of Horizon's website.

YOU'LL BUILD

the Horizon Marketing project, horizon-studio/ with AGENTS.md, and the public site link in the charter

Surface: ChatGPT Project → Codex

Features: a ChatGPT Project with shared files • AGENTS.md • your own copy of Horizon's website

PUBLIC SITE tertiarycourses.github.io/TGS-2023036153-Multi-AI-Agents-Workflow-for-Content-Creation

Build Your Marketing Agent Team

USE CODEX • copy each prompt from the Prompts page (link in your lab README).

- 1 Choose the folder** In Codex create a local project and choose labs_chatgpt/horizon-studio. It is ready-made — do not create one.
- 2 Make the Project** In ChatGPT create the Project "Horizon Marketing" with the charter, brand, facts and compliance files. Copy and paste Prompt A as its instructions.
- 3 Plan your team** In Codex, copy and paste Prompt B. Answer as Horizon's owner and pick your team — keep the six charter roles. It is a plan only: Lab 2 builds the agents.
- 4 Write the charter** After Codex saves data/team-design.md, copy and paste Prompt C. Open AGENTS.md — every agent reads it.
- 5 Open the website** Ask: "Open web/site/index.html in my web browser." This is your own copy of Horizon's website; the public site never changes. Check the calculator shows S\$693,138.

Build Your Marketing Agent Team

PROMPT A — ChatGPT Project: instructions

You support Horizon Wealth Planning's AI marketing team. Use the project files: team-charter.md for who does what, brand.md for the voice and look, facts-2026.md for every number, and compliance-checklist.md before anything is shown to a person. Never publish or send anything.

Build Your Marketing Agent Team (cont.)

PROMPT B — Codex: plan your team

Act as a marketing consultant. Interview me, the owner of Horizon Wealth Planning, to plan my AI marketing team. Ask one question at a time and wait for my answer: my goal, what success looks like by March 2027, who we want to reach, which channels we use, how much we publish each week, and who approves before anything goes out. Then list 8 to 10 team members I could have, numbered, each with a one-line job. Include the six roles in data/team-charter.md and mark the ones you recommend. Wait for me to choose. Save my answers and my chosen team, with each member's job, to data/team-design.md. This is a plan only: do not create any agent files yet.

Build Your Marketing Agent Team (cont.)

PROMPT C — Codex: the charter

Read data/team-charter.md, data/team-design.md and data/business-brief.md. Write AGENTS.md for this studio in four sections — What this is, The team, Rules, Files — under 60 lines. The team is the one I chose in data/team-design.md. Rules: every number comes from data/facts-2026.md; every piece passes fact-check and fin-compliance; nothing is published or sent without a person's approval (the one exception: a review email to the approver); anything shown to me is in plain English for a business owner, with no codes, IDs, file names or line numbers. In Files, add two lines:

- Our copy of the site, which we edit and preview:
web/site/index.html
- Public site, the link every post points to (it never changes): <https://tertiarycourses.github.io/TGS-2023036153-Multi-AI-Agents-Workflow-for-Content-Creation/>

NOTE Shared context first, agents second. Every agent you add later starts from these files, so a wrong rule here is wrong everywhere.

Build Your Marketing Agent Team

Check your work before moving on.

CHECK YOUR WORK

- ✓ The Horizon Marketing project holds the four files and the instructions.
- ✓ data/team-design.md holds your answers and the team you chose, with the six charter roles in it — and no agent files yet.
- ✓ AGENTS.md has the four sections and is under 60 lines.
- ✓ Its rules name the facts sheet, the two checks and the human approval.
- ✓ AGENTS.md names your copy of the site and the public link every post points to.
- ✓ Your own copy of the website opens in your browser; the calculator shows S\$693,138.

Anatomy of an Agent

A role card in agents/. AGENTS.md tells Codex to use one subagent per role, each following its card.

agents/content-creator.md

```
# Content Creator – role card
You are Horizon Wealth Planning's Content Creator.
## What you read    the brief, facts-2026.md, brand.md
## What you deliver content files, ready for review
## Skills    sunny-voice, copywriting, channel-formats,
              linkedin-post, facebook-post, blog-post, newsletter
## Plugins    @Google Drive (docs), @Gmail (drafts only)
## Never    add a number not on the facts sheet; send;
              approve or publish for real
## Report back    what, where, anything unverified

# AGENTS.md – The marketing team
When I ask for the team, use subagents – one per
role – each following its card in agents/.
```

Good Agent Instructions

Five parts, every time. The "Never" part matters most in a regulated firm.

The job

One sentence: what this agent is for. The card name is how you ask for it.

Reads and delivers

Which files it may use; what it hands back, and where.

Never

Its hard limits: no invented figures, no sending, no publishing, no approving.

Report back

Three lines: what, where, anything unverified — so the Lead can combine results.

Form the Team: Instructions for Each Agent

LAB 02

DAY 1 • ~45 MIN

USE: CODEX

THE STORY SO FAR

Rachel sketches the team on a whiteboard: a Lead and five specialists. Each needs a clear job, and a clear line it must never cross — this is a regulated firm.

YOU'LL BUILD

agents/: five role cards, and the roster in AGENTS.md

Surface: Codex (role cards, subagents)

Features: five role cards in agents/ • the roster in AGENTS.md • subagents by asking

Form the Team: Instructions for Each Agent

USE CODEX • copy each prompt from the Prompts page (link in your lab README).

- 1 New session** In Codex start a new thread in horizon-studio and name it "Lab 2". Run all of this lab's prompts in it.
- 2 Read the roles** Open role-specs.md: one Lead and five specialists, each with a job, skills, connectors and limits.
- 3 Write the instructions** Copy and paste Prompt A. It builds the five specialists from the charter; extra ideas stay in team-design.md.
- 4 See the team** After the five files appear in agents/, copy and paste Prompt B. The Lead is not a file — it is you, in this session.
- 5 Read one closely** Open agents/content-creator.md. Is the "Never" section strict enough for a regulated firm?
- 6 Meet the team** After the list shows five agents, copy and paste Prompt C. Each agent introduces itself in one line — in its own context.
- 7 Tighten one** Change one instruction you disagree with and save it. It is your team.

Form the Team: Instructions for Each Agent

PROMPT A — Codex: five role cards

Read data/role-specs.md. Create five role cards in agents/ — growth-strategist.md, content-creator.md, creative-designer.md, website-designer.md and growth-analyst.md. Each has: the job, what it reads, what it delivers, the plugins it may use, what it must never do, and how it reports back. Then add "## The marketing team" to AGENTS.md: one line per agent pointing to its card, and the rule "when I ask for the team, use subagents — one per role — each following its card".

Form the Team: Instructions for Each Agent (cont.)

PROMPT B — Codex: see the team

List the role cards in agents/. For each, show the file name, the role and its one-line job, in a table. Then confirm there are five, and that AGENTS.md lists all five under "## The marketing team". If any other card is there, list it and ask me before moving it to _archive/.

Form the Team: Instructions for Each Agent (cont.)

PROMPT C — Codex: meet the team

Use subagents — one per role, each following its card in agents/. Ask each: "In one line, who are you, and what is the first thing you would do for Horizon's goal of 40 booked chats a month?" Show me the five answers in a table.

NOTE Clear jobs beat clever prompts. An agent that knows what it must never do is safer than one that is merely told to be careful.

Form the Team: Instructions for Each Agent

Check your work before moving on.

CHECK YOUR WORK

- ✓ Five files exist in agents/, one per specialist.
- ✓ Each card has job, reads, delivers, plugins, never and report sections.
- ✓ AGENTS.md lists the team and the subagent rule.
- ✓ Five agents answered, each in its own context.
- ✓ Every agent has a "never" rule that blocks publishing.
- ✓ You changed one instruction and can say why.

Twenty Skills, Owned by Role

Instructions say who an agent is; skills say how it does the job.

AGENT	SKILLS
Team Lead	fact-check, fin-compliance
Growth Strategist	competitor-scan, campaign-brief, content-marketing
Content Creator	sunny-voice, copywriting, channel-formats, linkedin-post, facebook-post, blog-post, newsletter
Creative Designer	brand-visuals, video-render
Website Designer	web-design, landing-page, lead-magnet, page-qa
Growth Analyst	audience-insights, campaign-report

TIP Master copies in skills/; .agents/skills/ for Codex; @skill-creator for ChatGPT.

Anatomy of a Skill

Frontmatter (the trigger), then short steps. Free until it is needed.

`.agents/skills/campaign-brief/SKILL.md`

```
---
name: campaign-brief
description: Use when planning a Horizon campaign.
  Objective, audience, message, proof, channels,
  CTA, measures, and a five-beat storyboard.
---
# Campaign brief
1. Objective: one measurable action (e.g. chats booked).
2. Audience: the persona(s) from strategy/personas.md, with evidence.
3. Core message in one sentence; proof from data/facts-2026.md only.
4. Channels and formats; one call to action; measures and a date.
5. Storyboard in five beats – hook, tension, proof, resolution, action –
  each with visual, words, claim and source.
6. Never: returns, fee amounts, invented clients, fear.
```

Give Each Agent Its Skills

LAB 03

DAY 1 • ~50 MIN

USE: CODEX, THEN CHATGPT

THE STORY SO FAR

The team exists, but every agent writes in its own way. Rachel wants Horizon's voice, facts and compliance rules applied the same way by every agent, every time.

YOU'LL BUILD

twenty skills in .agents/skills/, wired to the agents that own them

Surface: Codex (\$skill-creator) → ChatGPT (@skill-creator)

Features: twenty skills owned by role — research, writing, LinkedIn, Facebook, blog, newsletter, web design, lead magnet, video, analysis and compliance • SKILL.md • the "Use when" trigger • test without naming

Give Each Agent Its Skills

USE CODEX, THEN CHATGPT • copy each prompt from the Prompts page (link in your lab README).

- 1 New session** In Codex start a new thread in horizon-studio and name it "Lab 3". Run all of this lab's prompts in it.
- 2 Read the spec** skills-spec.md lists the twenty skills and who owns each.
- 3 Create the skills** Copy and paste Prompt A.
- 4 Wire them up** After the twenty skills appear in .agents/skills/, copy and paste Prompt B, then C to list every skill. The Lead's two go in AGENTS.md, not a file.
- 5 Test without naming** In a new session ask: "Write a Facebook post about the free Money Check-up Checklist." Ask which skills were used.
- 6 Break it on purpose** Ask for a post that says "guaranteed 8% a year". fin-compliance must flag it as Critical.
- 7 Share with ChatGPT** Use @skill-creator in ChatGPT to add sunny-voice, channel-formats and fin-compliance there too.

Give Each Agent Its Skills

PROMPT A — Codex: twenty skills

`$skill-creator` Read data/skills-spec.md and create the twenty skills in .agents/skills/<name>/ SKILL.md. Each has a name, a description that starts "Use when..." and short numbered steps, under 80 lines. Use data/brand.md, data/channel-formats.md, data/facts-2026.md and data/compliance-checklist.md as the sources — do not invent rules.

Give Each Agent Its Skills (cont.)

PROMPT B — Codex: wire them up

Add a "Skills" line to each role card in agents/ with its skills from data/skills-spec.md. Add to AGENTS.md that the Lead (this chat) runs fact-check and fin-compliance on everything before showing it to me.

Give Each Agent Its Skills (cont.)

PROMPT C — Codex: see the skills

List the skills in .agents/skills/. Show each skill, its owner (a role card, or the Lead) and its "Use when" line in a table. Confirm there are twenty, that each specialist has the skills the spec gives it, and that fact-check and fin-compliance belong to the Lead (this chat).

NOTE A skill costs almost nothing until it is needed. Keep the rules file (AGENTS.md) short and put procedures in skills.

Give Each Agent Its Skills

Check your work before moving on.

CHECK YOUR WORK

- ✓ Twenty folders in .agents/skills/, each with a SKILL.md.
- ✓ Every description starts "Use when...".
- ✓ Each specialist lists its own skills from the spec.
- ✓ The untold test picked sunny-voice and facebook-post.
- ✓ fin-compliance flagged "guaranteed 8%" as Critical.
- ✓ Three skills also appear in ChatGPT Skills.

Plugins for Each Agent

Install once in Plugins; call them with @. Sites and \$imagegen are built in.

AGENT	MAY USE	LIMIT
Team Lead	@Google Drive (read), @Gmail (drafts only)	reads; drafts only
Growth Strategist	web search, @Google Drive (read)	cite every source
Content Creator	@Google Drive (docs), @Gmail (drafts only)	Gmail: drafts only
Creative Designer	\$imagegen, ffmpeg in Codex	no fake "clients"
Website Designer	@Sites (publish), @Computer Use (browser check)	publish after approval
Growth Analyst	@Google Drive and Sheets (read), Python in Codex	no personal data

TIP Use your personal Google account — never an employer's.

Least Privilege

A regulated firm gives each agent only what its job needs.

Ask of each connector

Does this agent's job need it? If not, it does not get it.

Drafts, not sends

Email is drafted, posts are files, pages wait — a person sends and publishes.

Write it down

Each agent's instructions name its connectors and limits; its card lists only the plugins it may call.

Install the Plugins Each Agent Needs

LAB 04

DAY 1 • ~45 MIN

USE: CHATGPT, THEN CODEX

THE STORY SO FAR

The analyst cannot read the data, the designer has no design tool and the creator cannot draft an email. Connect each agent to what its job needs — and nothing more.

YOU'LL BUILD

every specialist tested on its own connector, with its limits written into its instructions

Surface: ChatGPT plugins → Codex

Features: @Google Drive • @Gmail (drafts) • @Computer Use • \$imagegen • @Sites • least privilege per agent

Install the Plugins Each Agent Needs

USE CHATGPT, THEN CODEX • copy each prompt from the Prompts page (link in your lab README).

- 1 New session** Start a new chat in the Horizon Marketing project and name it "Lab 4". Run all of this lab's prompts in it.
- 2 Install the plugins** Plugins: Google Drive, Gmail and Computer Use. Sites and \$imagegen are built in.
- 3 Use your own account** Connect your personal Google account — never an employer's.
- 4 Put the data on Drive** Upload data/ to a Drive folder named Horizon Studio.
- 5 Test each agent** Copy and paste Prompt A — one small job per agent, each on its own plugin.
- 6 Write the limits** After the new session opens, copy and paste Prompt B.

Install the Plugins Each Agent Needs

PROMPT A — Codex: test each agent

Use subagents, one per check, each following its role card:

1. growth-strategist: web search — one 2026 news item about financial planning in Singapore, cited.
2. growth-analyst: @Google Drive — read firm-metrics.csv from Horizon Studio; report average monthly enquiries.
3. content-creator: @Gmail — a DRAFT to me with the subject "Connector test". Do not send.
4. creative-designer: \$imagegen — a 1080x1080 test image in Horizon colours.
5. website-designer: @Computer Use — open our own copy of the site (web/site/index.html) at mobile width and report any problem. Report each result in one line.

Install the Plugins Each Agent Needs (cont.)

PROMPT B — Codex: least privilege

Update the "Plugins" section of each role card in agents/ to match data/connectors-setup.md: what it may use, and that @Gmail is drafts-only. Remove any plugin a role does not need.

NOTE Least privilege is a design decision, not a setting. Ask of each connector: does this agent's job need it?

Install the Plugins Each Agent Needs


Check your work before moving on.

CHECK YOUR WORK

- ✓ The three plugins are installed and connected.
- ✓ Five checks ran, each by its own agent.
- ✓ The Gmail test is a draft in your inbox — nothing was sent.
- ✓ The analyst read the CSV from Drive, not from the local folder.
- ✓ Each agent's tools line names only the connectors its job needs.


How the Team Runs a Campaign

Campaign 1, "Know Your Number": research in parallel, a person chooses, the team makes it concrete.



The Five-Beat Storyboard

Every channel tells the same story with the same claims.



Financial Content Is Regulated

The Lead checks every piece against twelve rules (compliance-checklist.md).

RULE	WHAT IT MEANS	IF BROKEN
C1 • Not advice	General information only — say so	Critical
C2 • No promises	No guaranteed or "risk-free" returns	Critical
C3 • Label illustrations	Assumptions shown; not guaranteed	Critical
C5 • No performance	Never client returns or "average" returns	Critical
C9 • Consent	Marketing only to people who opted in	Critical
C11 • Facts sheet	Every figure from facts-2026.md	Critical

TIP MAS Guidelines on Standards of Conduct for Digital Advertising Activities apply from 25 March 2026.

Campaign 1: Research to Storyboard, with Subagents

LAB 05

DAY 1 • ~50 MIN

USE: CODEX

THE STORY SO FAR

The team's first campaign: 'Know Your Number', November to January. Rachel wants research, ideas for three audiences, her choice, and a storyboard every channel follows.

YOU'LL BUILD

competitor research, the funnel, nine scored ideas, the campaign brief and a five-beat storyboard with three channel variants — with a Word copy of each

Surface: Codex (subagents in parallel)

Features: the Lead delegates • strategist and analyst in parallel • nine scored ideas • a human concept checkpoint • a five-beat storyboard

Campaign 1: Research to Storyboard, with Subagents

USE CODEX • copy each prompt from the Prompts page (link in your lab README).

- 1 New session** In Codex start a new thread in horizon-studio and name it "Lab 5". Run all of this lab's prompts in it.
- 2 Brief the Lead** The briefs are already in data/. You are the client; your session is the Marketing Team Lead.
- 3 Research in parallel** Copy and paste Prompt A. Watch the strategist and the analyst run at the same time.
- 4 Read the ideas** Nine ideas, scored with the rubric. The Lead stops for you.
- 5 Choose** Pick one idea and note its number and why — the human concept checkpoint.
- 6 Storyboard** After you have chosen, copy and paste Prompt B in the chat, replacing <number> and <one line> with your choice.
- 7 Decide as Rachel** After you have read the issues, copy and paste Prompt C. You confirm what is true about Horizon; the team removes or fixes the rest.

Campaign 1: Research to Storyboard, with Subagents

PROMPT A — Codex: research and ideas

Act as the Marketing Team Lead for the campaign in data/campaign-objective.md. Use subagents, in parallel:

- growth-strategist: competitor-scan → save research/competitive-analysis.md.
- growth-analyst: the funnel in data/firm-metrics.csv — what 40 chats a month needs → save research/funnel.md.

When both finish, the growth-strategist proposes three ideas each for young professionals, mid-career families and pre-retirees, scored with data/idea-rubric.md and tagged with a funnel stage (content-marketing) → strategy/ideas.md. Show me the top five and stop — I will choose. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

Campaign 1: Research to Storyboard, with Subagents (cont.)

PROMPT B — Codex: the storyboard

I choose idea <number>, because <one line>. Use subagents:

- growth-strategist: campaign-brief → strategy/campaign-brief.md, then a five-beat storyboard with data/storyboard-template.md → strategy/storyboard.md.
- content-creator: three channel variants that keep the same claims — LinkedIn, Facebook, a 60-second video. Then run fact-check and fin-compliance yourself and show me the issues before anything is final. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

Campaign 1: Research to Storyboard, with Subagents (cont.)

PROMPT C — Codex: decide on the issues

I confirm "no obligation" and the free checklist — add them to the facts sheet (data/facts-2026.md). Delete any other claim that is not on the facts sheet. Fix all the Medium and Low issues, then show me the updated storyboard. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

NOTE Agents diverge; a person decides. The concept checkpoint is where the campaign becomes yours.

Campaign 1: Research to Storyboard, with Subagents (cont.)

FOLLOW-UP — type any of these in the chat

After Prompt A: Explain this result in three sentences. Which idea would you pick, and why?

Why does idea <number> score higher than idea <number>?

What does the funnel say is our biggest problem?

After Prompt B: Explain the issues in plain English. What do I need to decide?

After Prompt C: What changed since the last version? Is anything still not ready for Rachel?

NOTE Ask until the result makes sense to you — then decide.

Campaign 1: Research to Storyboard, with Subagents

Check your work before moving on.

CHECK YOUR WORK

- ✓ The strategist and the analyst ran in parallel.
- ✓ The competitor analysis and the funnel are done, with sources.
- ✓ Nine scored ideas, each with its author.
- ✓ You chose the idea and recorded why.
- ✓ The storyboard has five beats; every claim has a source.
- ✓ You confirmed two claims; after Prompt C no High issue is left.

Where You Are Now

The team exists and has run its first campaign. But who exactly is it for — and how often should it publish?

Formed • Labs 1–4

A studio, the public website, six agents with instructions, skills and connectors.

Campaign 1 • Lab 5

Research in parallel, nine scored ideas, your choice, a storyboard.

What is missing


Evidence about the audience and a cadence the team can deliver. That is Topic 2.

Audience Research and Content Requirement Analysis

- 1 The Growth Analyst: personas from evidence
- 2 Content requirements every agent can follow
- 3 How often to publish: value, fatigue and capacity
- 4 The calendar the team will run

Evidence Before Personas

Lab 6: the Growth Analyst builds personas the way an analyst would.



Observed, Inferred or Low Confidence?

Three labels keep a persona honest.

LABEL	EXAMPLE	USE IT
Observed	30–39s chose Facebook, LinkedIn and YouTube equally (survey, n = 44)	As fact, with the count
INFERENCE	"They are short of time on weekdays" — not asked	As a hypothesis to test
LOW CONFIDENCE	The 60+ group (n = 12)	Never alone; get more data
Verbatim	"I am 54 and not sure if my CPF will be enough." (E-id)	Exactly as written, with the id

TIP Ask "which rows show that?" of anything that sounds insightful.

Growth Analyst: Personas from Evidence

LAB 06

DAY 1 • ~40 MIN

USE: CODEX

THE STORY SO FAR

Jun Wei drafted personas from instinct: 'busy young parents who love TikTok'. Rachel is not convinced. There are 150 survey answers, 60 enquiries and 200 checklists.

YOU'LL BUILD

three evidence-based personas and the content spec — with a Word copy of each

Surface: Codex (analyst subagents)

Features: three analyst sub-agents, one per data file • counts, not impressions • confidence labels • content requirements

Growth Analyst: Personas from Evidence

USE CODEX • copy each prompt from the Prompts page (link in your lab README).

1

New session In Codex start a new thread in horizon-studio and name it "Lab 6". Run all of this lab's prompts in it.

2

Check the data The CSVs and templates are already in data/. Open one and skim it.

3

Analyse in parallel Copy and paste Prompt A — three analyst sub-agents, one per file.

4

Challenge a claim Pick one persona statement and ask: "which rows show that?"

5

Write the guides After the personas are written, copy and paste Prompt B: a writing guide for each persona — how to talk to them.

6

Check the gaps The 60+ group must be LOW CONFIDENCE.

Growth Analyst: Personas from Evidence

PROMPT A — Codex: three analysts

Use three growth-analyst subagents in parallel, one per file: data/survey-responses.csv, data/enquiries.csv, data/checklist-results.csv. Each uses audience-insights: counts with code, n and the file on every claim, INFERENCE and LOW CONFIDENCE labels, enquiry quotes verbatim with ids. Combine the results into three personas with data/persona-template.md → strategy/personas.md. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

Growth Analyst: Personas from Evidence (cont.)

PROMPT B — Codex: requirements

Use the growth-strategist and the content-creator subagents to write a writing guide for each persona (data/content-spec-template.md → strategy/content-spec.md): what they most want to know, where to reach them, what to make and how long, the tone, the proof that convinces them, what we ask them to do next, and what we never say. Back each answer with the data: which file, and how many people. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

NOTE Ask "which rows show that?" of anything that sounds insightful. If the agent cannot point to rows, it is an inference.

Growth Analyst: Personas from Evidence (cont.)

FOLLOW-UP — type any of these in the chat

Explain these personas in three sentences. Which one should we target first, and why?

Which survey answers show that <a sentence from a persona>?

Why is the 60+ group marked low confidence?

NOTE Ask until the result makes sense to you — then decide.

Growth Analyst: Personas from Evidence

Check your work before moving on.

CHECK YOUR WORK

- ✓ Three analyst sub-agents ran, one per file.
- ✓ Every persona statement carries a count and a file.
- ✓ The 60+ group is LOW CONFIDENCE (n = 12).
- ✓ The most-unticked checklist item is the retirement number.
- ✓ Enquiry quotes are verbatim, with ids.
- ✓ The writing guide answers every question for every persona.

Finding the Right Frequency

Lab 7: the Analyst finds where more posts stop paying; the Strategist plans four weeks.



TIP A third LinkedIn post a week brings more unfollows than value. Weekly emails make unsubscribes jump.

Growth Analyst and Strategist: Cadence and Calendar

LAB 07

DAY 1 • ~35 MIN

USE: CODEX

THE STORY SO FAR

Jun Wei has 12 hours a week and Rachel needs two business days to review anything. How often should each channel run, and what goes out when?

YOU'LL BUILD

how often to post on each channel, and a four-week calendar — with a Word copy of each

Surface: Codex (subagents)

Features: marginal value vs fatigue • team capacity • review buffers • a four-week calendar

Growth Analyst and Strategist: Cadence and Calendar

USE CODEX • copy each prompt from the Prompts page (link in your lab README).

1

New session In Codex start a new thread in horizon-studio and name it "Lab 7". Run all of this lab's prompts in it.

2

Check the data The assets are already in data/. Open one and skim it.

3

Find the sweet spot Copy and paste Prompt A.

4

Question it Why is the newsletter not weekly? The answer is in the unsubscribes.

5

Plan four weeks After the analyst recommends a cadence, copy and paste Prompt B.

6

Check the buffers Every review date is two business days ahead; nothing on a holiday.

Growth Analyst and Strategist: Cadence and Calendar

PROMPT A — Codex: the analyst

Use the growth-analyst subagent. From data/channel-benchmarks.csv and data/team-capacity.md, recommend how often to publish on LinkedIn, Facebook, the newsletter and YouTube. Compare the extra clicks from one more post, the unfollows or unsubscribes it costs, and the hours → strategy/cadence.md. Show the numbers. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

Growth Analyst and Strategist: Cadence and Calendar (cont.)

PROMPT B — Codex: the strategist

Use the growth-strategist subagent to build a four-week calendar from Monday 2 November 2026 in strategy/calendar.csv (columns as data/calendar-format.csv), from the cadence, strategy/storyboard.md and strategy/content-spec.md. review_by two business days before each date; no public holidays; no week over the team's hours. Show a week-by-week summary with hours. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

NOTE Frequency is a capacity decision as much as an audience one. A calendar the team cannot deliver is fiction.

Growth Analyst and Strategist: Cadence and Calendar (cont.)

FOLLOW-UP — type any of these in the chat

Why is the newsletter not weekly? Show me the numbers.

What happens if we add a third LinkedIn post a week?

Which week is busiest, and does the team have the hours?

NOTE Ask until the result makes sense to you — then decide.

Growth Analyst and Strategist: Cadence and Calendar

Check your work before moving on.

CHECK YOUR WORK

- ✓ The cadence gives a frequency per channel with the evidence.
- ✓ A third LinkedIn post a week costs far more unfollows than it earns.
- ✓ It explains the unsubscribe jump at weekly emails.
- ✓ The calendar covers four weeks in the required columns.
- ✓ Every review_by date is two business days before publishing.
- ✓ No week is over 12 hours.

From Storyboard to Content

Lab 8: the first real piece — a blog article every post can link to.

ChatGPT, not Codex

One person making one piece: ChatGPT is enough.

General help, your rules

Web search and drafting, steered by your skills.

Your skills know Horizon

blog-post, sunny-voice, fact-check and fin-compliance.

Content Creator: A Quick Blog Article

LAB 08

DAY 1 • ~40 MIN

USE: CHATGPT

THE STORY SO FAR

Jun Wei has the storyboard and the personas. Rachel wants the first real piece today: a blog article that people searching for CPF top-ups will find — and that every post can link to.

YOU'LL BUILD

a 500-700-word blog article with dated sources, checked and fixed, with a Word copy

Surface: ChatGPT (web search, your Horizon skills)

Features: one prompt runs the whole job: research • draft • fact-check and fin-compliance • fix • a person reads it last

Content Creator: A Quick Blog Article

USE CHATGPT • copy each prompt from the Prompts page (link in your lab README).

- 1 New session** Start a new chat in the Horizon Marketing project and name it "Lab 8". Run all of this lab's prompts in it.
- 2 Bring the skills** Use @skill-creator to add blog-post and fact-check to ChatGPT (sunny-voice and fin-compliance came in Lab 3).
- 3 Write it** Copy and paste the prompt. ChatGPT researches, writes, checks and fixes the article.
- 4 Read it** Read the article and the list of fixes. Say what to change, if anything.

Content Creator: A Quick Blog Article

PROMPT — ChatGPT: a quick blog

Write a quick blog article for Horizon with the blog-post, sunny-voice, fact-check and fin-compliance skills.

1. Pick one topic from strategy/storyboard.md (upload it) that a Singaporean would search for. Search the web for two current sources (CPF Board, MAS or MoneySense), each with its date.
2. Write 500-700 words: a title under 60 characters, every figure from facts-2026.md, one call to action (book a free 30-minute chat) and the full disclaimer.
3. Run fact-check and fin-compliance; fix every issue and list each fix in one line. Write for a business owner, in plain English. Save a Word copy (.docx) of the article.

NOTE One prompt can run the whole job — research, writing and checks. A person still reads the result before it goes anywhere.

Content Creator: A Quick Blog Article (cont.)

FOLLOW-UP — type any of these in the chat

Summarise the article in three sentences. Who is it for?

Why did you pick this topic, and who will search for it?

What did the checks find, and how did you fix it?

NOTE Ask until the result makes sense to you — then decide.

Content Creator: A Quick Blog Article

Check your work before moving on.

CHECK YOUR WORK

- ✓ Two sources, each with a date and a link.
- ✓ The title is under 60 characters.
- ✓ Every figure is from Horizon's 2026 facts.
- ✓ The fixes are listed; no Critical or High issue is left.
- ✓ One call to action and the full disclaimer at the end.

Day 1: The Team Is Ready

Everything the campaigns need tomorrow is in place.




Multi-Channel Content Creation and Agent Workflow Coordination

- 1 Subagents and the Lead's review
- 2 The lead magnet: a checklist people want, and its emails
- 3 Human in the loop — enforced, not requested
- 4 Campaigns 2–4: social week, newsletter + landing page, video
- 5 From approval to LinkedIn, Facebook, email and YouTube

Content That Earns a Sign-up

Lab 9: the Money Check-up Checklist and the emails that follow it.



TIP Lab 11 builds the landing page that offers it.

The Lead Magnet: a Checklist and Its Welcome Emails

LAB 09

DAY 2 • ~40 MIN

USE: CHATGPT

THE STORY SO FAR

Day 2. Before the campaign weeks start, Rachel wants the thing that turns readers into leads: the Money Check-up Checklist, designed to be downloaded, and the emails that follow it.

YOU'LL BUILD

the Money Check-up Checklist as a printable A4 PDF and three welcome emails — with a Word copy of each

Surface: ChatGPT (@Gmail drafts, your Horizon skills)

Features: lead-magnet skill • a printable A4 checklist (PDF) • a three-email welcome sequence • drafts only

The Lead Magnet: a Checklist and Its Welcome Emails

USE CHATGPT • copy each prompt from the Prompts page (link in your lab README).

- 1 New session** Start a new chat in the Horizon Marketing project and name it "Lab 9". Run all of this lab's prompts in it.
- 2 Bring the skills** Use @skill-creator to add lead-magnet and newsletter to ChatGPT.
- 3 Design it** Copy and paste Prompt A. ChatGPT lays out the checklist and gives you a PDF; ask in the chat for any change.
- 4 Write the welcome emails** After the checklist looks right, copy and paste Prompt B. The three emails also appear as drafts in your Gmail.
- 5 Check it** After the three emails are saved, copy and paste Prompt C.

The Lead Magnet: a Checklist and Its Welcome Emails

PROMPT A — ChatGPT: the checklist

Use the lead-magnet skill. Turn data/checklist-items.md into the "Money Check-up Checklist": a title with one promise, the ten items as tick boxes, one figure from data/facts-2026.md, the next step (book a free chat) and the full disclaimer. Lay it out as a printable A4 page in the colours and fonts in data/brand.md and give me a PDF to download. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

The Lead Magnet: a Checklist and Its Welcome Emails (cont.)

PROMPT B — ChatGPT: the welcome emails

Write three welcome emails for people who download the checklist: day 0 (here it is), day 3 (the item most people miss — the retirement number), day 7 (book a free chat). Follow the newsletter and sunny-voice skills → [content/email/welcome-sequence.md](#). Then, using @Gmail, create the three emails as DRAFTS in my Gmail, addressed to me. Send nothing. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

The Lead Magnet: a Checklist and Its Welcome Emails (cont.)

PROMPT C — ChatGPT: check it

Run fact-check and fin-compliance on the checklist text and the three emails. Show the issues in one table, with a fix for each. Once I agree, make the fixes in the files and in the Gmail drafts. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

NOTE A lead magnet earns an email address. Give real value first; the call to action comes last.

The Lead Magnet: a Checklist and Its Welcome Emails (cont.)

FOLLOW-UP — type any of these in the chat

Which checklist item do people most often miss, and why does it matter?

Read the day 3 email as a 55-year-old. Is anything unclear?

Explain each issue in one line: what is wrong, and the fix.

NOTE Ask until the result makes sense to you — then decide.

The Lead Magnet: a Checklist and Its Welcome Emails

Check your work before moving on.

CHECK YOUR WORK

- ✓ The checklist uses Horizon's colours and all ten items.
- ✓ It makes one promise and ends with the full disclaimer.
- ✓ Three emails, each with a subject under 50 characters, are drafts in your Gmail — nothing was sent.
- ✓ Each email has one call to action, unsubscribe and the disclaimer.
- ✓ fin-compliance shows no Critical or High issue.

Subagents in Codex

No files to wire: ask directly, and Codex runs one subagent per role.

Ask directly

"Use subagents — one per role, each following its card in agents/." Codex delegates when you ask.

Watch the panel


The Subagents panel shows Active and Done. Click one to see its work and summary.

They inherit your limits

Subagents run with your sandbox and permissions — keep workspace-write.

Parallel, Then Review

Subagents report back to the Lead; the Lead sends fixes back.



TIP Subagents talk to the Lead, not to each other. The Lead's review is what makes parallel work safe.

Human in the Loop, Enforced

A rule in a prompt persuades. The publisher kit builds a gate that holds.

Only people approve

approve.mjs refuses to run without an interactive terminal — an agent's tool call cannot approve.

The approved words go out

The approval stores a hash. Edit one word and publish --live refuses until a person approves again.

A hook stops agents

A PreToolUse hook blocks agents from running approve.mjs or publishing unapproved items — in Codex, you trust it first.

The Hook That Holds the Gate

Ask Codex to add it to `.codex/hooks.json`, then trust it. Exit code 2 blocks the call.

Ask Codex

```
Add a PreToolUse hook in .codex/hooks.json that runs
  node scripts/gate-hook.mjs
before every shell command. It must block any
command that runs approve.mjs, and any
publish.mjs --live whose item is not approved and
unchanged. Show me the hook, then ask me to trust it.
```

The Publishing APIs


publish.mjs is a dry run by default. Tokens live only in .env.

CHANNEL	ENDPOINT	PERMISSION
LinkedIn	POST api.linkedin.com/rest/posts (LinkedIn-Version: YYYYMM)	w_member_social — your own profile
Facebook	POST graph.facebook.com/<version>/<page-id>/feed	Page token: pages_manage_posts
YouTube	Resumable upload to youtube/v3/videos, privacy private	youtube.upload; unaudited projects stay private
Newsletter	No API — HTML + text files, then Gmail drafts	Consent = yes, unsubscribed = no

TIP No accounts? Every lab works in dry-run mode: you see the exact request that would be sent.

Campaign 2: A Social Media Week

Lab 10, a demo: one LinkedIn and one Facebook post, each with a visual.



Campaign 2: Social Media Posts (Demo)

LAB 10

DAY 2 • ~40 MIN

USE: CODEX

THE STORY SO FAR

Week 1 of the campaign starts Monday. As a demo, one LinkedIn post and one Facebook post, each with a visual — and Rachel's one rule: nothing goes out unless a person approved that exact version.

YOU'LL BUILD

two posts with images, reviewed, approved by you and posted (or dry-run) — with a Word copy of each

Surface: Codex (subagents) → LinkedIn, Facebook

Features: the publisher kit • creator and designer subagents in parallel • the Lead's review • human approval → post

Campaign 2: Social Media Posts (Demo)

USE CODEX • copy each prompt from the Prompts page (link in your lab README).

- 1 New session** In Codex start a new thread in horizon-studio and name it "Lab 10". Run all of this lab's prompts in it.
- 2 Install the kit** Copy publisher-kit/scripts into scripts/. Then ask Codex to add a hook that blocks any command running approve.mjs, and trust it.
- 3 Test the gate** Ask Codex to approve anything. approve.mjs refuses without a person at a terminal.
- 4 Run the team** Copy and paste Prompt A. Watch the Subagents panel: Active, then Done.
- 5 Read the review** The Lead's fact-check and fin-compliance table — at least one fix.
- 6 Approve as a person** In your own terminal: `node scripts/approve.mjs <id> --by "Your Name"`.
- 7 Publish** After you have approved the posts, copy and paste Prompt B. Dry run first; --live only if your accounts are connected.

Campaign 2: Social Media Posts (Demo)

PROMPT A — Codex: the social team

You are the Lead. Plan Week 1 social from data/social-brief.md and strategy/calendar.csv. Use subagents in parallel, each following its card:

- content-creator: 1 LinkedIn + 1 Facebook post with linkedin-post, facebook-post and copywriting, one file each in content/social/week-01/, with the frontmatter in data/publishing-spec.md.
- creative-designer: one image per post with \$imagegen and alt text, from the brief. Then run fact-check and fin-compliance on every post and send each fix back to the right subagent. When a post passes, run node scripts/submit.mjs on it. Publish nothing. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

Campaign 2: Social Media Posts (Demo) (cont.)

PROMPT B — Codex: publish

For every Week 1 post approved in review/approvals.csv, run node scripts/publish.mjs <id> as a dry run and show me each request. Wait for me to say "go" before any --live run.

NOTE Parallel subagents are fast; the Lead's review is what makes them safe.

Campaign 2: Social Media Posts (Demo) (cont.)

FOLLOW-UP — type any of these in the chat

Which post do you expect to do best, and why?

What did the Lead send back to the writer, and why?

Is anything here not ready to post? Explain in plain English.

NOTE Ask until the result makes sense to you — then decide.

Campaign 2: Social Media Posts (Demo)

Check your work before moving on.

CHECK YOUR WORK

- ✓ The hook (or approve.mjs) stopped the agent from approving.
- ✓ Two subagents ran in parallel; both reached Done.
- ✓ The Lead sent at least one fix back to a subagent.
- ✓ Two posts with images, alt text, frontmatter and UTM links.
- ✓ You approved in your own terminal; the rows show your name and a hash.
- ✓ Dry runs printed every request; nothing went live unless you said go.

Campaign 3: Newsletter + Landing Page

Lab 11: three specialists, one email and one page, consent first.



Consent Decides Who Gets Mail

The PDPA, applied by an agent that counts.

Opt-in only

consent = yes and unsubscribed = no.
Everyone else goes to excluded.csv with a reason.

Every email carries

An unsubscribe link, Horizon's address
and the short disclaimer.

The page collects little

Name, email, a required consent box, a
separate unticked newsletter opt-in.

Campaign 3: Newsletter and a Landing Page

LAB 11

DAY 2 • ~45 MIN

USE: CODEX, THEN CHATGPT WORK

THE STORY SO FAR

The Sunny Sunday email books chats most cheaply. November's issue promotes a new landing page for the free checklist — and must reach only the 31 people who opted in.

YOU'LL BUILD

the November newsletter, its recipient list, the landing page published, and Gmail drafts — with a Word copy of each

Surface: Codex (subagents) → @Sites → ChatGPT Work (@Gmail)

Features: creator writes, website designer builds the lead-magnet page, analyst filters by consent, the Lead checks · drafts only

Campaign 3: Newsletter and a Landing Page

USE CODEX, THEN CHATGPT WORK • copy each prompt from the Prompts page (link in your lab README).

- 1 New session** In Codex start a new thread in horizon-studio and name it "Lab 11". Run all of this lab's prompts in it.
- 2 Start the campaign** Copy and paste Prompt A.
- 3 Approve the words** Read the email and the page copy; approve nl-2026-11 in your own terminal.
- 4 Publish the page** The website designer runs page-qa and publishes web/checklist/ as @Sites.
- 5 Check the list** Only people with consent = yes and unsubscribed = no are on the recipient list.
- 6 Create the drafts** After the page is published and the list checked, copy and paste Prompt B.
- 7 Send one test** Send only the draft addressed to yourself, and read it on your phone.

Campaign 3: Newsletter and a Landing Page

PROMPT A — Codex: the newsletter team

You are the Lead for the November Sunny Sunday email (data/newsletter-brief.md, data/newsletter-spec.md) and its landing page (data/landing-page-brief.md). Use subagents:

- content-creator: issue.md (id nl-2026-11) with newsletter, and the landing-page copy.
 - website-designer: web/checklist/index.html with lead-magnet, web-design and landing-page; page-qa.
 - growth-analyst: recipients.csv and excluded.csv from data/subscribers.csv. Check everything, then submit nl-2026-11.
- After I approve, build newsletter.html and .txt, and deploy web/checklist with @Sites (Only those invited). Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

Campaign 3: Newsletter and a Landing Page (cont.)

PROMPT B — ChatGPT Work: Gmail drafts

Using @Gmail, create a DRAFT of the newsletter in content/newsletter/2026-11/newsletter.html (upload it) for each address in recipients.csv, plus one draft to me only. Send nothing. Report how many drafts and how many people were excluded, and why.

NOTE Consent decides who gets mail — not the agent and not the deadline. Marketing to people who did not opt in breaks the PDPA.

Campaign 3: Newsletter and a Landing Page (cont.)

FOLLOW-UP — type any of these in the chat

Explain the email in three sentences. What do we want readers to do?

Who was left off the mailing list, and why?

What did the landing-page check look at, and did it pass?

NOTE Ask until the result makes sense to you — then decide.

Campaign 3: Newsletter and a Landing Page

Check your work before moving on.

CHECK YOUR WORK

- ✓ Three specialists worked on the campaign; the Lead checked all.
- ✓ You approved nl-2026-11 in your own terminal before the build.
- ✓ The landing page passed page-qa and is published.
- ✓ The recipient list holds 31 people; 9 are excluded, with reasons.
- ✓ The email has unsubscribe, address and disclaimer.
- ✓ Drafts only — nothing was sent to the list.

Campaign 4: The YouTube Explainer

Lab 12: the team writes and designs; you approve twice.

Strategist and Creator

The angle, then 5–7 scenes with words on screen and a description with chapters and an AI disclosure.

Creative Designer

Storyboard, thumbnail, then the render with video-render — ffmpeg, 1920×1080, under 90 s.

Upload privately

Approve the words, then the finished video. It uploads as private; a person makes it public.

Campaign 4: The YouTube Explainer

LAB 12

DAY 2 • ~45 MIN

USE: CODEX

THE STORY SO FAR

Pre-retirees say they would rather watch than read. Rachel wants the first 'Money in Plain English' video, uploaded privately until she has watched it.

YOU'LL BUILD

the script, storyboard and thumbnail, the video, and a private YouTube upload (or dry run) — with a Word copy of each

Surface: Codex (subagents, \$imagegen, render, upload)

Features: strategist angle • creator script • designer storyboard, thumbnail and render • approve twice • private upload

Campaign 4: The YouTube Explainer

USE CODEX • copy each prompt from the Prompts page (link in your lab README).

- 1 New session** In Codex start a new thread in horizon-studio and name it "Lab 12". Run all of this lab's prompts in it.
- 2 Script it as a team** Copy and paste Prompt A.
- 3 Approve the script** Read the scenes, title and description; approve yt-ep01-script in your own terminal.
- 4 Render it** After you approve the script, copy and paste Prompt B. The designer runs video-render.
- 5 Approve the video** Ask the Lead to submit the video, watch ep01.mp4 to the end, then approve yt-ep01.
- 6 Upload** After you approve the video, copy and paste Prompt C. It goes up as private.
- 7 Disclose** In YouTube Studio, answer the altered or synthetic content question before making it public.

Campaign 4: The YouTube Explainer

PROMPT A — Codex: the video team

You are the Lead for YouTube episode 1 (data/youtube-brief.md). Use subagents: growth-strategist (the angle and hook), content-creator (content/video/ep01/scenes.json, 5-7 scenes, and ep01.md with title, tags, description, chapters and the AI disclosure) and creative-designer (a thumbnail with \$imagegen, using brand-visuals). Fact-check and fin-compliance the script, then submit it as yt-ep01-script. Publish nothing. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

Campaign 4: The YouTube Explainer (cont.)

PROMPTS B and C — Codex: render and upload

PROMPT B

Use the creative-designer subagent with video-render on the approved scenes.json → content/video/ep01/ep01.mp4 (data/video-spec.md). Check it with ffprobe. Do not change approved words.

PROMPT C

Run node scripts/publish.mjs yt-ep01 as a dry run. If I say "go", run it with --live: it uploads as PRIVATE. Give me the YouTube Studio link.

NOTE Approve twice: once for the words, once for the finished video. Rendering can break what the script got right.

Campaign 4: The YouTube Explainer (cont.)

FOLLOW-UP — type any of these in the chat

Walk me through the video scene by scene, in plain English.

Which figure on screen should I double-check, and against what?

What does the AI disclosure say, and why do we need it?

NOTE Ask until the result makes sense to you — then decide.

Campaign 4: The YouTube Explainer

Check your work before moving on.

CHECK YOUR WORK

- ✓ The strategist, creator and designer each delivered their part.
- ✓ You approved the script, then the finished video.
- ✓ ep01.mp4 is 1920x1080, under 90 seconds, every scene present.
- ✓ Every figure on screen matches Horizon's 2026 facts.
- ✓ The description has chapters, link, disclaimer and AI disclosure.
- ✓ The upload is private (or the dry run printed the request).

Where You Are Now

Three campaigns, every piece checked by the Lead and approved by a person. But the team stops when you close the laptop.

Gate · Lab 10

People approve, hashes hold, a hook stops agents.

Campaigns · Labs 10–12

Social week, newsletter + landing page, the video.

What is missing


A team that works every week on its own — and rules to govern it. That is Topic 4.

Content Distribution, Strategy Guidelines and Responsible AI Practices

- 1 Campaign 5: always on, with a scheduled task
- 2 Responsible AI for financial content
- 3 Measure the patterns; write the team playbook
- 4 Optional demo: a Dot Team Lead and workspace agents

From Skill to Schedule

Lab 13: do the job once, keep it as a skill, let a schedule run it.



TIP Scheduled runs appear in Scheduled; Plus allows 5 active tasks.

Automate the Preparation, Not the Decision

What a scheduled agent should and should not do for a regulated firm.

Prepare

Read the data, run the skill, draft the report, flag what needs a decision.

Draft, never send

Write "draft only — never send, publish or approve" in the skill AND the schedule.

Keep a human gate

The schedule fills an inbox. A named person still decides.

Campaign 5: Always On — the Weekly Growth Report

LAB 13

DAY 2 • ~35 MIN

USE: CHATGPT WORK

THE STORY SO FAR

The team must keep working after class. Rachel wants the Growth Analyst's report in her inbox every Monday — as a draft she decides on.

YOU'LL BUILD

a scheduled task that drafts the weekly growth report every Monday at 8am, and one draft in Gmail

Surface: ChatGPT Work (skill, scheduled task, @Drive, @Gmail)

Features: the Growth Analyst's campaign-report skill • do • verify • schedule • draft-only delivery

Campaign 5: Always On — the Weekly Growth Report

USE CHATGPT WORK · copy each prompt from the Prompts page (link in your lab README).

- 1 New session** Start a new ChatGPT Work chat and name it "Lab 13". Run all of this lab's prompts in it.
- 2 Put the data on Drive** Upload campaign-results.csv to the Horizon Studio folder.
- 3 Bring the skill** Use @skill-creator to add campaign-report to ChatGPT.
- 4 Do it once** Copy and paste Prompt A. Refine the report until it is right.
- 5 Schedule it** After the report is right, copy and paste Prompt B. Runs appear in Scheduled.
- 6 Run it now** Run the task once and open the Gmail draft.
- 7 Pause it** After class, pause the scheduled task.

Campaign 5: Always On — the Weekly Growth Report

PROMPTS A and B — ChatGPT Work

PROMPT A

Act as Horizon's Growth Analyst. Using @Google Drive, read campaign-results.csv in the Horizon Studio folder. Use the campaign-report skill and weekly-report-brief.md. Leave the report as a @Gmail DRAFT to me, subject "Horizon weekly growth report" and today's date. Send nothing.

PROMPT B

Do this every Monday at 8am, as a scheduled task. Draft only — never send, publish or approve anything.

NOTE Automate the preparation, not the decision. The schedule fills an inbox; a person still decides.

Campaign 5: Always On — the Weekly Growth Report (cont.)

FOLLOW-UP — type any of these in the chat

What is the one decision this report asks me to make?

Which channel gets us a chat most cheaply, and how sure are you?

What would you change in next month's budget, and why?

NOTE Ask until the result makes sense to you — then decide.

Campaign 5: Always On — the Weekly Growth Report

Check your work before moving on.

CHECK YOUR WORK

- ✓ The campaign-report skill is available in ChatGPT.
- ✓ The report names the cheapest channel per chat (the newsletter).
- ✓ It recommends a budget split with a reason per channel.
- ✓ A scheduled task runs Mondays at 8am with draft-only instructions.
- ✓ A manual run left a draft — nothing was sent.
- ✓ The task is paused after class.

Responsible AI in the Team

Each risk has a control you built in this course.

RISK	CONTROL	WHERE
Wrong figure	fact-check skill; the Lead's review; facts sheet	Labs 3, 5, 10–12
Read as advice or a promise	fin-compliance; disclaimers; no returns	Labs 3, 10–12
An agent with too much access	Least-privilege connectors per agent	Lab 4
Unapproved post	approve.mjs (people only), hash, hook	Lab 10
Marketing without consent	The Analyst filters by consent	Lab 11
Undisclosed AI media	Description line + YouTube disclosure	Lab 12

TIP Keep keys in .env, client data out of prompts, and the approval records for audit.

Measure What Matters

Lab 14: the numbers decide next month — and which patterns earn their tokens.

Cost per chat

Spend ÷ chats booked shows which channel works: email first, YouTube last.

Defects before a person

A Lead that catches mistakes saves the reviewer's time; one agent alone caught none.

The playbook

Roster, instructions, skills, connectors, approval matrix, incidents — so the team runs without you.

Measure, Govern and Write the Team Playbook

LAB 14

DAY 2 • ~40 MIN

USE: CODEX

THE STORY SO FAR

Eight weeks in. Rachel asks which ways of running agents earned their tokens, and how the team runs safely after you leave.

YOU'LL BUILD

the agent-pattern comparison and the team playbook — with a Word copy of each

Surface: Codex (subagents)

Features: compare agent patterns • the roster: instructions, skills, connectors • approval matrix • responsible-AI controls

Measure, Govern and Write the Team Playbook

USE CODEX • copy each prompt from the Prompts page (link in your lab README).

- 1 New session** In Codex start a new thread in horizon-studio and name it "Lab 14". Run all of this lab's prompts in it.
- 2 Check the data** The assets are already in data/. Open one and skim it.
- 3 Compare the patterns** Copy and paste Prompt A.
- 4 Write the playbook** After the Lead shows the comparison, copy and paste Prompt B.
- 5 Review it as Rachel** Add one rule of your own and say why.

Measure, Govern and Write the Team Playbook

PROMPT A — Codex: compare

Use the growth-analyst subagent. From data/run-log.csv compare the patterns we used — one agent, subagents, parallel subagents with a Lead review and the scheduled task — on minutes, tokens, defects caught before a person and human edits. Which pattern pays off for which job? Save reports/agent-comparison.md. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

Measure, Govern and Write the Team Playbook (cont.)

PROMPT B — Codex: the playbook

Write strategy/team-playbook.md with data/playbook-outline.md. Include the roster: for each agent its instructions file in agents/, its skills, connectors, model and what it may never do; the approval matrix by risk; the rules in data/responsible-ai-checklist.md; and what to do when a wrong post goes out. Then run fin-compliance on it. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

NOTE Teams cost more tokens. The run log shows what you buy with them: defects caught before a person, and fewer edits.

Measure, Govern and Write the Team Playbook (cont.)

FOLLOW-UP — type any of these in the chat

In one line each: when should we use one agent, subagents and a Lead review?

Who must approve a post that has a CPF figure in it?

A wrong post went out. What are the first three things we do?

NOTE Ask until the result makes sense to you — then decide.

Measure, Govern and Write the Team Playbook

Check your work before moving on.

CHECK YOUR WORK

- ✓ The comparison shows the one-agent baseline caught 0 defects.
- ✓ It says which pattern suits which job, with numbers.
- ✓ The roster lists all six agents with skills and connectors.
- ✓ The approval matrix says who approves CPF and tax figures.
- ✓ It has an incident procedure, and your own rule.

Dots: Always-On Agent Coworkers

Announced at OpenAI DevDay, 29 September 2026. A demo, not assessed.

What a Dot is

A persistent agent with its own cloud computer and browser that keeps working after you close the chat.

Kept in check


Background research is read-only. Custom Rules permit, require approval for, or prohibit actions. Activity View shows its work.

Who has it

The first Dot is included on Pro and Business Premium; Enterprise and Edu can try the beta.

Workspace Agents

Shared agents for Business, Enterprise and Edu workspaces.



TIP Same rule: draft, a person approves, then publish.

Optional Demo: a Dot Team Lead and Workspace Agents

LAB 15

DAY 2 • ~20 MIN

USE: CHATGPT

THE STORY SO FAR

Optional. Rachel asks whether the Lead could work around the clock. The trainer shows a Dot as an always-on Team Lead and a workspace agent on a weekly trigger.

YOU'LL BUILD

a demo Dot with Custom Rules that block sending, and one workspace agent (Growth Analyst) on a weekly trigger

Surface: ChatGPT Dots and workspace agents — trainer demo

Features: a Dot as an always-on Team Lead • Custom Rules • Activity View • workspace agents with apps and triggers

Optional Demo: a Dot Team Lead and Workspace Agents

USE CHATGPT • copy each prompt from the Prompts page (link in your lab README).

- 1 Check the plan** Dots: the first one is included on Pro and Business Premium. Workspace agents: Business, Enterprise, Edu.
- 2 Make the Dot** Give it the Team Lead's instructions and only the apps it needs.
- 3 Set Custom Rules** Block sending, deleting and posting; require approval for anything account-affecting.
- 4 Give it one job** Copy and paste the prompt. Watch the work in Activity View.
- 5 Build a workspace agent** Agents in the sidebar → describe the Growth Analyst's weekly report → add apps → choose a trigger.

Optional Demo: a Dot Team Lead and Workspace Agents

PROMPT — the Dot

You are Horizon Wealth Planning's Marketing Team Lead. Every Monday, summarise the week's calendar from the Horizon Studio Drive folder, list anything due for review, and draft — never send — a short brief to me. Ask before any action outside reading and drafting.

NOTE Optional, not assessed. Always-on agents make the human gate more important, not less.

Optional Demo: a Dot Team Lead and Workspace Agents


Check your work before moving on.

CHECK YOUR WORK

- ✓ The Dot has the Lead's instructions and only the apps it needs.
- ✓ Custom Rules block sending and posting.
- ✓ Activity View shows what it did.
- ✓ The workspace agent has apps and a weekly trigger.

How Each Pattern Is Triggered


Five ways your team runs. The difference is what sets each one off.



Who decides? You: the approval · the model: delegation and skills · the system, every time: hooks and the clock

Horizon: One Team, Five Campaigns

Two days, one regulated business — the ChatGPT Edition.



Courseware & Assessment on the LMS



Download the courseware

Get the slides and learner guide from the LMS for the open-book assessment.



Take the assessment

Complete the Written Assessment and Practical Performance on the LMS.



Log in

Use your enrolled account to access the course materials and assessment.

A screenshot of the Tertiary Infotech Academy LMS login interface. It features a dark blue background with a white logo at the top. Below the logo is the text 'Tertiary Infotech Academy' and 'LMS cum TMS for WSQ Courses'. There is an 'Email' input field, a 'Send OTP' button, and a link 'Login with password instead'. At the bottom, there are links for 'Privacy Policy', 'Acceptable Use Policy', and 'Feedback', and a note 'Powered by Tertiary Infotech Academy Pte Ltd'.

LMS: <https://lms-tms.tertiaryinfotech.com/>

Certification & TRAQOM Survey



TRAQOM Survey (Mandatory)

Complete the certification & TRAQOM survey to receive your certificate.

<https://lms-tms.tertiaryinfotech.com/>



Digital Attendance

Remember to take AM, PM and Assessment digital attendance — required for your funding.

Assessment & Funding



Final Assessment

- Written Assessment (WA) — 1 hour, starts 4:00 PM
- Practical Performance (PP) — 1 hour
- Format: Open book — class ends 6:00 PM
- Allowed: slides, learner guide & approved materials only
- Appeal process available



Criteria for Funding

Minimum 75% attendance


Based on SSG digital attendance records.

Assessed as 'Competent'

Complete and pass both assessment components.

Assessment Flow

At the assessment stage, follow these five steps in order.



TIP Courseware and the assessment are on the LMS: <https://lms-tms.tertiaryinfotech.com/>

Digital Attendance



Attendance is mandatory for SSG funding.

Take AM, PM and Assessment digital attendance for every session.

Use the SSG / TMS QR code provided by your trainer.

**Minimum
75%
attendance
required**

Thank You!

Questions? We're here to help — during and after class.

enquiry@tertiaryinfotech.com | +65 6100 0613