



Tertiary Infotech Academy Pte Ltd

UEN: 201200696W

LESSON PLAN

For

Multi AI Agents Workflow for Content Creation (Claude Edition)

TGS Ref No: TGS-2023036153

Conducted by

Tertiary Infotech Academy Pte Ltd

UEN: 201200696W

Version 3.5

DOCUMENT VERSION CONTROL RECORD

Version Number	Effective Date of Release	Summary of Included Changes	Author
1.0	13 Sep 2026	Multi-agent content strategy courseware aligned to the current course page and TSC; GreenLeaf Café case; 12 labs.	Tertiary Infotech Academy
2.0	8 Oct 2026	Revamped in the house design of the Agentic AI Applications with Codex course. One central case — Horizon Wealth Planning, a Singapore financial-planning firm — across 14 labs, each in its own folder: Claude Cowork sub-agents (research, ideation, personas, cadence, skills, measurement), Codex (website, compliance subagent, approval gate, video), Claude Code subagents and agent teams (social, newsletter and video teams with LinkedIn, Facebook and YouTube publishing behind a human approval gate) and Hermes Agent bots on a Kanban board. Practice Exam slide (Claude Certified Associate — Foundations). WA (6 SAQ, K1-K6) and PP (4 tasks, A1-A7) rewritten for the v2.0 labs; assessment 4:00-6:00 PM Day 2 (WA 1 h, PP 1 h).	Dr. Alfred Ang
2.1	8 Oct 2026	Two Claude tools only: Claude (Chat and Cowork are one product) and Claude Code, which replaces Codex for the website, compliance subagent, approval gate and video. The website is published and shared as a Claude artifact instead of GitHub Pages. New slides: Claude plans and pricing; the agentic loop, why it matters and the loop in action; permission modes; publishing a Claude artifact. Topic 4 runs a scheduled weekly brief in Claude (Lab 12) and the playbook (Lab 13); Hermes Agent becomes an optional demo at the end (Lab 14). PP Task 4 updated to the Topic 4 labs.	Dr. Alfred Ang
3.0	8 Oct 2026	Two editions from one source — Claude and ChatGPT — each with its own deck, Lesson Plan, Learner Guide, labs and assessment. Rebuilt around an AI marketing team: a Marketing Team Lead and five specialists (Growth Strategist, Content Creator, Creative Designer, Website Designer, Growth Analyst), each with instructions, two skills and its own connectors. The team runs five campaigns: research to storyboard, a social week, newsletter and landing page, a YouTube explainer and an always-on weekly report. Claude edition: Claude Code subagents and agent teams; the website as a Claude artifact; optional Hermes Agent demo. 13 labs.	Dr. Alfred Ang
3.2	9 Oct 2026	Horizon's website is given, not built: Lab 1 now opens one public link (a shared Claude artifact) with a GitHub Pages backup for learners who cannot open artifacts or Sites, and adds the link to the team charter. The website links are printed and clickable on the scenario slide; the Practice Exam slide links to the exam. Lab 4 checks the public site; broken site URLs fixed.	Dr. Alfred Ang
3.3	9 Oct 2026	No-code Claude Edition: every lab says which app to use (Claude, or Claude Code) on its slides, README	Dr. Alfred Ang

		and Learner Guide page; Labs 6–9 and 13–15 run in Claude, Labs 3, 4 and 11 in both. Two new content labs: Lab 8 (a blog article from the storyboard, Day 1) and Lab 9 (the lead magnet and its welcome emails, Day 2); later labs renumbered 10–15. A person approves by typing "approved <id> by <name>" — the gate comes pre-installed. The optional demo is now code-free. Skills widened for content creation: twenty skills, adding content-marketing, copywriting, linkedin-post, facebook-post, blog-post, newsletter, web-design and lead-magnet, used in Labs 5 and 8–11. Lab 4 connects every tool the labs use (Google Drive, Gmail, Google Calendar, Firecrawl for web research); the Website Designer uses the built-in browser — Playwright removed.	
3.4	10 Oct 2026	Day 2 ends with agent teams at work: the optional demo is replaced by Lab 15, a hands-on agent-team lab in two parts — an advocate and a critic argue the pros and cons of a "Turning 55" CPF campaign while the Lead judges, then a QA loop in which a QA Auditor sends a failed draft straight back to the writer until it passes. New slides: subagents vs agent teams (diagram), a team that argues, and the QA loop. Hermes Agent references removed. Lab 15 runs 30 minutes; the course summary 15:45-16:00.	Dr. Alfred Ang
3.5	10 Oct 2026	Labs made for learners who are not IT people. Every lab runs in Code mode, in its own session named after the lab. From Lab 5 every result opens as a plain-English artifact with a Word copy, and each lab lists follow-up questions to ask. Prompts are copied from the Prompts page. Horizon's website gains a blog; each learner works on their own local copy of the site: Lab 8 (now one quick-blog prompt) posts an article to it and Lab 9 adds the checklist, with the welcome emails as Gmail drafts. Canva removed. Lab 10 is a two-post demo (40 min; Lab 9 40 min). Labs 10-11: review in the artifact or by a review email, then approve in the chat.	Dr. Alfred Ang

TABLE OF CONTENTS

Course Overview	5
Learning Outcomes	5
Daily Schedule	5
Day 1 — Form the Team.....	5
Day 2 — Run the Campaigns.....	6
Topic-by-topic Breakdown	7
Topic 1 — Multi-AI-Agent Content Ideation and Digital Storyboarding · Slides 16–74.....	7
Lab 1 — Build Your Marketing Agent Team · Slides 36–40.....	8
Lab 2 — Form the Team: Instructions for Each Agent · Slides 43–48.....	8
Lab 3 — Give Each Agent Its Skills · Slides 51–56.....	8
Lab 4 — Connect the Tools Each Agent Needs · Slides 59–63.....	8
Lab 5 — Campaign 1: Research to Storyboard, with Subagents · Slides 67–73.....	9
Topic 2 — Audience Research and Content Requirement Analysis · Slides 75–97.....	9
Lab 6 — Growth Analyst: Personas from Evidence · Slides 78–83.....	9
Lab 7 — Growth Analyst and Strategist: Cadence and Calendar · Slides 85–90.....	9
Lab 8 — Content Creator: A Quick Blog Article · Slides 92–96.....	10
Topic 3 — Multi-Channel Content Creation and Agent Workflow Coordination · Slides 98–138.....	10
Lab 9 — The Lead Magnet: a Checklist and Its Welcome Emails · Slides 100–107.....	10
Lab 10 — Campaign 2: Social Media Posts (Demo) · Slides 115–122.....	11
Lab 11 — Campaign 3: Newsletter and a Landing Page · Slides 125–130.....	11
Lab 12 — Campaign 4: The YouTube Explainer · Slides 132–137.....	11
Topic 4 — Content Distribution, Strategy Guidelines and Responsible AI Practices · Slides 139–174.....	11
Lab 13 — Campaign 5: Always On — the Weekly Growth Report · Slides 142–147.....	12
Lab 14 — Measure, Govern and Write the Team Playbook · Slides 150–155.....	12
Lab 15 — Agent Teams: A Strategy Debate and a QA Loop · Slides 159–165.....	12
Resources Required	12
Assessment	13

Course Overview

A two-day, hands-on WSQ course in content strategy with multi-agent AI workflows. Learners take one business — Horizon Wealth Planning, an independent financial-planning firm in Singapore with a website and no marketing team — and build it an AI marketing team: a Marketing Team Lead and five specialists (Growth Strategist, Content Creator, Creative Designer, Website Designer, Growth Analyst), each with instructions, skills and connectors. Day 1 forms the team and runs its first campaign; Day 2 runs a social week, a newsletter with a landing page, a YouTube explainer and an always-on weekly report, with a person approving every external piece. This Claude Edition uses Claude Code subagents and agent teams, Claude connectors and scheduled tasks, and publishes pages as Claude artifacts.

Course code TGS-2023036153 · 2 days · 16 hours including a 2-hour assessment · 174-slide deck v3.5 · 15 labs, each in its own folder. TSC: ICT-SNM-4004-1.1 Content Strategy.

Learning Outcomes

By the end of this course, participants will be able to:

- **LO1.** Conceptualize content ideas to meet marketing objectives and map out digital storyboards as part of a content strategy.
- **LO2.** Identify content requirements from customer preferences and determine the frequency of marketing content.
- **LO3.** Determine content types and styles and decide on modes and processes for distributing content.
- **LO4.** Develop guidelines for content strategy execution using appropriate delivery modes and responsible AI practices.

Daily Schedule

Both days run 9:00 AM – 6:00 PM with two 10-minute tea breaks and a 45-minute lunch. The Slides column gives the deck pages for each teaching block.

Day 1 — Form the Team

Time	Duration	Topic / Activity	Method	Slides
9:00 AM–9:45 AM	45 min	Welcome, attendance, trainer, ground rules, learning outcomes; overview — the Claude suite, Claude, Claude Code, plans and pricing, AI 2023–2026, the harness, the agentic loop, agentic AI vs AI agents, what an agent is made of; the Horizon scenario and your AI marketing team, lesson plan, assessment briefing	Presentation, discussion	1–15
9:45 AM–10:00 AM	15 min	Topic 1 — why a team of agents; getting the tools ready; a home for the team	Presentation, live demo	16–35
10:00 AM–10:30 AM	30 min	Lab 1 — Build Your Marketing Agent Team	Hands-on lab	36–40

10:30 AM–10:40 AM	10 min	Tea break	—	—
10:40 AM–10:55 AM	15 min	Agent instructions: job, inputs, outputs, limits	Presentation, live demo	41–42
10:55 AM–11:40 AM	45 min	Lab 2 — Form the team: instructions for each agent	Hands-on lab	43–48
11:40 AM–11:50 AM	10 min	Skills: twenty procedures, owned by role	Presentation	49–50
11:50 AM–12:40 PM	50 min	Lab 3 — Give each agent its skills	Hands-on lab	51–56
12:40 PM–1:25 PM	45 min	Lunch	—	—
1:25 PM–1:35 PM	10 min	Connectors and least privilege	Presentation	57–58
1:35 PM–2:20 PM	45 min	Lab 4 — Connect the Tools Each Agent Needs	Hands-on lab	59–63
2:20 PM–2:30 PM	10 min	Campaign 1: how the team works; the five-beat storyboard; regulated content	Presentation	64–66
2:30 PM–3:20 PM	50 min	Lab 5 — Campaign 1: research to storyboard, with subagents	Hands-on lab	67–73
3:20 PM–3:30 PM	10 min	Tea break	—	—
3:30 PM–3:35 PM	5 min	Topic 1 recap	Discussion	74
3:35 PM–3:45 PM	10 min	Topic 2 — evidence before personas; data agents count	Presentation	75–77
3:45 PM–4:25 PM	40 min	Lab 6 — Growth Analyst: personas from evidence	Hands-on lab	78–83
4:25 PM–4:35 PM	10 min	Finding the right frequency	Presentation	84
4:35 PM–5:10 PM	35 min	Lab 7 — Growth Analyst and Strategist: cadence and calendar	Hands-on lab	85–90
5:10 PM–5:15 PM	5 min	From storyboard to content: A quick blog in one prompt	Presentation	91
5:15 PM–5:55 PM	40 min	Lab 8 — Content Creator: A Quick Blog Article	Hands-on lab	92–96
5:55 PM–6:00 PM	5 min	Day 1 review and Q&A	Discussion	97

Day 2 — Run the Campaigns

Time	Duration	Topic / Activity	Method	Slides
9:00 AM–9:05 AM	5 min	Topic 3 — content types and styles; the lead magnet	Presentation	98–99
9:05 AM–9:45 AM	40 min	Lab 9 — The Lead Magnet: a Checklist and Its Welcome Emails	Hands-on lab	100–107
9:45 AM–10:00 AM	15 min	Subagents vs agent teams; human in the loop, enforced; the publishing APIs	Presentation, live demo	108–114
10:00 AM–10:40 AM	40 min	Lab 10 — Campaign 2: Social Media Posts (Demo)	Hands-on lab	115–122

10:40 AM–10:50 AM	10 min	Tea break	—	—
10:50 AM–10:55 AM	5 min	Campaign 3: newsletter, landing page and consent	Presentation	123–124
10:55 AM–11:40 AM	45 min	Lab 11 — Campaign 3: newsletter and a landing page	Hands-on lab	125–130
11:40 AM–11:45 AM	5 min	Campaign 4: the video team	Presentation	131
11:45 AM–12:30 PM	45 min	Lab 12 — Campaign 4: the YouTube explainer	Hands-on lab	132–137
12:30 PM–12:35 PM	5 min	Topic 3 recap	Discussion	138
12:35 PM–1:20 PM	45 min	Lunch	—	—
1:20 PM–1:35 PM	15 min	Topic 4 — always on: from skill to schedule	Presentation, live demo	139–141
1:35 PM–2:10 PM	35 min	Lab 13 — Campaign 5: the weekly growth report	Hands-on lab	142–147
2:10 PM–2:20 PM	10 min	Responsible AI for financial content; measuring what matters	Presentation	148–149
2:20 PM–3:00 PM	40 min	Lab 14 — Measure, govern and write the team playbook	Hands-on lab	150–155
3:00 PM–3:10 PM	10 min	Tea break	—	—
3:10 PM–3:15 PM	5 min	Agent teams at work: subagents vs agent teams; a team that argues; the QA loop	Presentation	156–158
3:15 PM–3:45 PM	30 min	Lab 15 — Agent Teams: A Strategy Debate and a QA Loop	Hands-on lab	159–165
3:45 PM–4:00 PM	15 min	Course summary: the team, the campaigns, which app for which job; Q&A and assessment briefing	Discussion	166–167
4:00 PM–5:00 PM	60 min	Written Assessment (WA) — open book, 1 hour	Assessment	—
5:00 PM–6:00 PM	60 min	Practical Performance (PP) — 1 hour; TRAQOM, attendance, submit, sign the record	Assessment	168–174

Topic-by-topic Breakdown

Topic 1 — Multi-AI-Agent Content Ideation and Digital Storyboarding • Slides 16–74

- AI agents, the tools, and Horizon's story
- Why a team of agents
- A home for the team; Horizon's public website
- Form the team: instructions, skills, connectors
- Campaign 1: research, ideas and a storyboard

Lab 1 — Build Your Marketing Agent Team • Slides 36–40

- **Use:** Claude Code — Code mode (the `</>` button) in the Claude desktop app, in your horizon-studio folder
- **Surface:** Claude Code
- **Time:** 30 minutes (Day 1)
- **Features practised:** CLAUDE.md as the team charter · your own copy of Horizon's website
- **Output:** horizon-studio/ with CLAUDE.md, the shared data, and your own copy of the site running at `http://localhost:8080/`
- **Evidence:** Claude shows horizon-studio as the chosen folder. Your website opens at `http://localhost:8080/`; the calculator shows S\$693,138.
- **Lab folder:** labs_claude/lab-01-studio-and-website/

Lab 2 — Form the Team: Instructions for Each Agent • Slides 43–48

- **Use:** Claude Code — Code mode (the `</>` button) in the Claude desktop app, in your horizon-studio folder
- **Surface:** Claude Code (subagent files)
- **Time:** 45 minutes (Day 1)
- **Features practised:** five subagent files in `.claude/agents/` · name, description, tools, model · the Lead is your session
- **Output:** `.claude/agents/`: growth-strategist, content-creator, creative-designer, website-designer, growth-analyst
- **Evidence:** Five files exist in `.claude/agents/`, one per specialist. You changed one instruction and can say why.
- **Lab folder:** labs_claude/lab-02-agent-instructions/

Lab 3 — Give Each Agent Its Skills • Slides 51–56

- **Use:** Claude Code — Code mode (the `</>` button) in the Claude desktop app, in your horizon-studio folder
- **Surface:** Claude Code (`.claude/skills/`)
- **Time:** 50 minutes (Day 1)
- **Features practised:** twenty skills owned by role — research, writing, LinkedIn, Facebook, blog, newsletter, web design, lead magnet, video, analysis and compliance · SKILL.md · the "Use when" trigger · test without naming
- **Output:** twenty skills in `.claude/skills/`, wired to the agents that own them
- **Evidence:** Twenty folders in `.claude/skills/`, each with a SKILL.md. fin-compliance flagged "guaranteed 8%" as Critical.
- **Lab folder:** labs_claude/lab-03-agent-skills/

Lab 4 — Connect the Tools Each Agent Needs • Slides 59–63

- **Use:** Claude Code — Code mode (the `</>` button) in the Claude desktop app, in your horizon-studio folder
- **Surface:** Claude desktop app (Customize → Connectors), Claude Code
- **Time:** 45 minutes (Day 1)

- **Features practised:** Customize → Connectors: Google Drive · Gmail (drafts) · Google Calendar · Firecrawl · /plugin: Marketing (by Anthropic) · built-in browser · least privilege per agent
- **Output:** every specialist tested on its own connector, with its limits written into its instructions
- **Evidence:** /mcp lists Google Drive, Gmail, Google Calendar and Firecrawl; /plugin shows Marketing. Each agent's tools line names only the connectors its job needs.
- **Lab folder:** labs_claude/lab-04-agent-connectors/

Lab 5 — Campaign 1: Research to Storyboard, with Subagents · Slides 67–73

- **Use:** Claude Code — Code mode (the `</>` button) in the Claude desktop app, in your horizon-studio folder
- **Surface:** Claude Code (subagents in parallel)
- **Time:** 50 minutes (Day 1)
- **Features practised:** the Lead delegates · strategist and analyst in parallel · nine scored ideas · a human concept checkpoint · a five-beat storyboard
- **Output:** competitor research, the funnel, nine scored ideas, the campaign brief and a five-beat storyboard with three channel variants — shown as artifacts, with a Word copy of each
- **Evidence:** The strategist and the analyst ran in parallel. You confirmed two claims; after Prompt C no High issue is left.
- **Lab folder:** labs_claude/lab-05-campaign-1-research-to-storyboard/

Topic 2 — Audience Research and Content Requirement Analysis · Slides 75–97

- The Growth Analyst: personas from evidence
- Content requirements every agent can follow
- How often to publish: value, fatigue and capacity
- The calendar the team will run

Lab 6 — Growth Analyst: Personas from Evidence · Slides 78–83

- **Use:** Claude Code — Code mode (the `</>` button) in the Claude desktop app, in your horizon-studio folder
- **Surface:** Claude Code (three analyst subagents in parallel)
- **Time:** 40 minutes (Day 1)
- **Features practised:** three analyst sub-agents, one per data file · counts, not impressions · confidence labels · content requirements
- **Output:** three evidence-based personas and the content spec — shown as artifacts, with a Word copy of each
- **Evidence:** Three analyst sub-agents ran, one per file. The writing guide answers every question for every persona.
- **Lab folder:** labs_claude/lab-06-personas-from-evidence/

Lab 7 — Growth Analyst and Strategist: Cadence and Calendar · Slides 85–90

- **Use:** Claude Code — Code mode (the `</>` button) in the Claude desktop app, in your horizon-studio folder
- **Surface:** Claude Code (subagents)
- **Time:** 35 minutes (Day 1)

- **Features practised:** marginal value vs fatigue · team capacity · review buffers · a four-week calendar
- **Output:** how often to post on each channel, and a four-week calendar — shown as artifacts, with a Word copy of each
- **Evidence:** The cadence gives a frequency per channel with the evidence. No week is over 12 hours.
- **Lab folder:** labs_claude/lab-07-cadence-and-calendar/

Lab 8 — Content Creator: A Quick Blog Article · Slides 92–96

- **Use:** Claude Code — Code mode (the `</>` button) in the Claude desktop app, in your horizon-studio folder
- **Surface:** Claude Code (Firecrawl, your Horizon skills) → your copy of Horizon's website
- **Time:** 40 minutes (Day 1)
- **Features practised:** one prompt runs the whole job: research · draft · fact-check and fin-compliance · fix · post to the blog on your copy of the website · a person reads it last
- **Output:** a 500-700-word blog article with dated sources, checked and fixed and posted on your own copy of Horizon's website, opened in your browser, with a Word copy
- **Evidence:** Two sources, each with a date and a link. Your article is first under "From our blog" on your copy of the website.
- **Lab folder:** labs_claude/lab-08-blog-article/

Topic 3 — Multi-Channel Content Creation and Agent Workflow Coordination · Slides 98–138

- Subagents or an agent team?
- The lead magnet: a checklist people want, and its emails
- Human in the loop — enforced, not requested
- Campaigns 2–4: social week, newsletter + landing page, video
- From approval to LinkedIn, Facebook, email and YouTube

Lab 9 — The Lead Magnet: a Checklist and Its Welcome Emails · Slides 100–107

- **Use:** Claude Code — Code mode (the `</>` button) in the Claude desktop app, in your horizon-studio folder
- **Surface:** Claude Code (artifact, Gmail drafts, Marketing plugin, your Horizon skills)
- **Time:** 40 minutes (Day 2)
- **Features practised:** lead-magnet skill · a printable A4 checklist as an artifact · the Marketing plugin's email sequence · newsletter and sunny-voice · drafts only
- **Output:** the Money Check-up Checklist as a printable A4 page (artifact and PDF) and on your copy of Horizon's website, and three welcome emails — shown as artifacts, with a Word copy of each
- **Evidence:** The checklist uses Horizon's colours and all ten items. The checklist is on your copy of the website: email and consent first, then the ten tick boxes.
- **Lab folder:** labs_claude/lab-09-lead-magnet/

Lab 10 — Campaign 2: Social Media Posts (Demo) • Slides 115–122

- **Use:** Claude Code — Code mode (the `</>` button) in the Claude desktop app, in your horizon-studio folder
- **Surface:** Claude Code (agent team) → LinkedIn, Facebook
- **Time:** 40 minutes (Day 2)
- **Features practised:** the publisher kit and its hook · an agent team: creator, designer, Lead · teammates message each other · human approval → you post
- **Output:** two posts with images, reviewed, approved by you and posted by you (or, optionally, through the API) — shown as artifacts, with a Word copy of each
- **Evidence:** The gate stopped Claude from approving. You posted the approved text and image yourself (or, optionally, through the API after you said go).
- **Lab folder:** labs_claude/lab-10-campaign-2-social-week/

Lab 11 — Campaign 3: Newsletter and a Landing Page • Slides 125–130

- **Use:** Claude Code — Code mode (the `</>` button) in the Claude desktop app, in your horizon-studio folder
- **Surface:** Claude Code (agent team) → Claude artifact, Gmail
- **Time:** 45 minutes (Day 2)
- **Features practised:** creator writes, website designer builds the lead-magnet page, analyst filters by consent, the Lead checks · drafts only
- **Output:** the November newsletter, its recipient list, the landing page published, and Gmail drafts — shown as artifacts, with a Word copy of each
- **Evidence:** Three specialists worked on the campaign; the Lead checked all. Drafts only — nothing was sent to the list.
- **Lab folder:** labs_claude/lab-11-campaign-3-newsletter-and-landing-page/

Lab 12 — Campaign 4: The YouTube Explainer • Slides 132–137

- **Use:** Claude Code — Code mode (the `</>` button) in the Claude desktop app, in your horizon-studio folder
- **Surface:** Claude Code (agent team, render, upload)
- **Time:** 45 minutes (Day 2)
- **Features practised:** strategist angle · creator script · designer storyboard, thumbnail and render · approve twice · private upload
- **Output:** the script, storyboard and thumbnail, the video, and a private YouTube upload (or dry run) — shown as artifacts, with a Word copy of each
- **Evidence:** The strategist, creator and designer each delivered their part. The upload is private (or the dry run printed the request).
- **Lab folder:** labs_claude/lab-12-campaign-4-youtube-explainer/

Topic 4 — Content Distribution, Strategy Guidelines and Responsible AI Practices • Slides 139–174

- Campaign 5: always on, with a scheduled task
- Responsible AI for financial content
- Measure the patterns; write the team playbook
- Agent teams at work: argue a strategy, then a QA loop

Lab 13 — Campaign 5: Always On — the Weekly Growth Report • Slides 142–147

- **Use:** Claude Code — Code mode (the `</>` button) in the Claude desktop app, in your horizon-studio folder
- **Surface:** Claude Code (skill, routine, Drive, Gmail)
- **Time:** 35 minutes (Day 2)
- **Features practised:** the Growth Analyst's campaign-report skill · do · verify · schedule · draft-only delivery
- **Output:** a scheduled task that drafts the weekly growth report every Monday at 8am, and one draft in Gmail
- **Evidence:** The report used the campaign-report skill from your folder. The task is paused after class.
- **Lab folder:** labs_claude/lab-13-campaign-5-weekly-growth-report/

Lab 14 — Measure, Govern and Write the Team Playbook • Slides 150–155

- **Use:** Claude Code — Code mode (the `</>` button) in the Claude desktop app, in your horizon-studio folder
- **Surface:** Claude Code (subagents)
- **Time:** 40 minutes (Day 2)
- **Features practised:** compare agent patterns · the roster: instructions, skills, connectors · approval matrix · responsible-AI controls
- **Output:** the agent-pattern comparison and the team playbook — shown as artifacts, with a Word copy of each
- **Evidence:** The comparison shows the one-agent baseline caught 0 defects. It has an incident procedure, and your own rule.
- **Lab folder:** labs_claude/lab-14-team-playbook/

Lab 15 — Agent Teams: A Strategy Debate and a QA Loop • Slides 159–165

- **Use:** Claude Code — Code mode (the `</>` button) in the Claude desktop app, in your horizon-studio folder
- **Surface:** Claude Code (agent teams: Lead, teammates, shared task list, direct messages)
- **Time:** 30 minutes (Day 2)
- **Features practised:** an advocate and a critic argue the pros and cons · the Lead judges · a writer and a QA Auditor loop until PASS · FAIL goes straight back to the writer · human approval
- **Output:** the decision (pros, cons, verdict), then an email, a video script and a thumbnail that passed QA, submitted for your approval — shown as artifacts, with a Word copy of each
- **Evidence:** The advocate and the critic messaged each other; every point cites a file. Both teams were shut down and cleaned up.
- **Lab folder:** labs_claude/lab-15-agent-teams-debate-and-qa-loop/

Resources Required

- Trainer: the v3.5 Claude Edition slide deck, this Lesson Plan, the Learner Guide and the lab pack (labs_claude/, 15 folders, with reference solutions for Labs 1, 2, 3, 10 and 12).

- Each learner: a laptop with the Claude desktop app on a Pro plan or above (Claude Code and artifacts are not on the Free plan), signed in with their claude.ai account. Learners type no commands: each lab says whether to use Claude or Claude Code.
- Installed by the trainer before class: Claude Code, Node.js 22, Python 3 with Pillow and ffmpeg (they run the approval gate and the video). A personal Google account (Drive, Gmail, Calendar). A free Firecrawl account (1,000 credits a month, no card) for web research; Anthropic's Marketing plugin (free, added in Lab 4).
- Lab 15 (agent teams) runs in the same studio folder; agent teams come switched on in `.claude/settings.json`.
- Optional, for live publishing: a LinkedIn developer app, a test Facebook Page and a Google Cloud project with the YouTube Data API. Every lab also works in dry-run mode without them.
- All fixtures are synthetic. Every learner email address resolves to the learner's own Gmail via plus-addressing.

Assessment

- Day 2, 4:00 – 6:00 PM: Written Assessment (WA, short-answer, 1 hour) from 4:00 PM, then Practical Performance (PP, 1 hour) from 5:00 PM.
- Open book: slides, Learner Guide and approved materials only.
- Funding criteria: minimum 75% attendance, and assessed as Competent in both the WA and the PP.
- Assessment flow: TRAQOM survey → assessment digital attendance → WA and PP → submit answers on the LMS → sign the Assessment Summary Record.
- Courseware and assessment are on the LMS: <https://lms-tms.tertiaryinfotech.com/>