

WSQ COURSE • TGS-2023036153 • CLAUDE EDITION

Version 3.1 • 9 Oct 2026

Multi AI Agents Workflow for Content Creation

Tertiary Infotech Academy Pte Ltd
UEN: 201200696W

www.tertiarycourses.com.sg | enquiry@tertiaryinfotech.com | +65 6100 0613

© Tertiary Infotech Academy Pte Ltd

Digital Attendance



Attendance is mandatory for SSG funding.

Take AM, PM and Assessment digital attendance for every session.

Use the SSG / TMS QR code provided by your trainer.

**Minimum
75%
attendance
required**

About the Trainer



[Trainer Name]

[Title / Role]



Qualifications: _____



Expertise: _____



Experience: _____



Contact / Profile: _____

About the Trainer



Dr. Alfred Ang

AI Trainer & Engineer



PhD-qualified trainer specialising in AI, software engineering and developer tooling.



Builds AI marketing teams of agents with Claude and Claude Code.



Trains professionals across Singapore on practical, project-based AI skills.



Certifications: PMP, AWS, Microsoft, CompTIA, CKA, CKAD.



[linkedin.com/in/angchewhoe](https://www.linkedin.com/in/angchewhoe)



github.com/alfredang

INTRODUCTION

Let's Know Each Other

Take a minute to introduce yourself to the class:



Your name & role



What you do day to day



Your content work today



What you want from this
course

Ground Rules



Phones on silent

Step out quietly for calls or breaks.



Participate actively

No question is stupid — ask freely.



Mutual respect

Agree to disagree; one conversation at a time.



Be punctual

Return from breaks on time.



75% attendance

Required for funding eligibility.



Stay hands-on

Follow along with every activity.

Learning Outcomes

By the end of the two days you will be able to:

L01 · Conceptualise

Conceptualise content ideas that meet marketing objectives, and map them into digital storyboards.

L02 · Identify

Identify content requirements from customer preferences, and decide how often to publish.

L03 · Determine


Determine content types and styles, and the modes and processes for distributing content.

L04 · Develop

Develop guidelines for executing the content strategy, with suitable delivery modes and responsible AI practices.


Build a Marketing Team, Then Run Campaigns

Horizon Wealth Planning has a website, a goal of 40 booked chats a month — and no marketing team. You build one, made of AI agents.



Meet Horizon Wealth Planning

The firm's public website — the same link for every learner. You do not build it; every campaign points back here.



WEBSITE claude.ai/artifact/J82VCbzFogYps2vep8UYhL

BACKUP tertiarycourses.github.io/TGS-2023036153-Multi-AI-Agents-Workflow-for-Content-Creation

Plain English

Retirement, CPF, insurance and investing — no jargon, no hard sell.

One goal


Turn visitors into booked free 30-minute chats.

Regulated

Financial content: general information only, never a promised return.

Your AI Marketing Team

One Lead and five specialists — each with instructions, two skills and its own connectors.



Who Does What

Each agent's job, the skills it owns and the connectors it may use.

AGENT	JOB	SKILLS	CONNECTORS
Team Lead	Plans, assigns, checks, hands to a person	fact-check, fin-compliance	Google Drive (read), Gmail (drafts only)
Growth Strategist	Research, positioning, briefs, storyboards	competitor-scan, campaign-brief	web search and fetch, Google Drive (read)
Content Creator	Posts, newsletter, scripts, page copy	sunny-voice, channel-formats	Google Drive (docs), Gmail (drafts only)
Creative Designer	Images, thumbnails, the video	brand-visuals, video-render	Canva, Bash for Pillow and ffmpeg
Website Designer	Landing pages, published and checked	landing-page, page-qa	artifacts (publish), Playwright MCP (browser check)
Growth Analyst	Personas, cadence, results, reports	audience-insights, campaign-report	Google Drive and Sheets (read), Bash for Python

TIP You build every part of this table in Labs 2, 3 and 4.

The Tools for This Edition


The Claude Edition: two Claude products, one studio.

TOOL	WHAT THE TEAM USES IT FOR	LABS
Claude Code	The Lead's session; subagent files; agent teams; skills; hooks; publishing artifacts	Labs 1–10, 12
Claude	Connectors (Drive, Gmail, Canva), skills, scheduled tasks	Labs 4, 11
Claude artifacts	Horizon's public website (given) and the campaign landing page, shared by link	Labs 1, 9
Hermes Agent	Optional demo: the team on an open-source agent	Lab 13 (optional)

TIP Every external post, email and video passes one gate: a person approves the exact version.

Course Outline

Four topics over two days. Day 1 forms the team; Day 2 runs its campaigns.



Lab Materials

Thirteen labs, each in its own folder with a README, prompts (MD and PDF), assets, an evidence checklist and, where useful, a solution.

DAY	LABS	TOPIC
1	1 Studio + website • 2 Instructions • 3 Skills • 4 Connectors	T1 • form the team
1	5 Campaign 1: research to storyboard	T1 • Campaign 1
1	6 Personas • 7 Cadence and calendar	T2 • audience
2	8 Social week • 9 Newsletter + landing page • 10 Video	T3 • Campaigns 2–4
2	11 Weekly report • 12 Playbook • 13 Optional demo	T4 • Campaign 5

TIP One horizon-studio folder from Lab 1 to Lab 12 — the team keeps what it learns.

SCHEDULE

Lesson Plan (9:00 AM – 6:00 PM)

Tea breaks 10 min, lunch 45 min. Full timings and slide numbers are in the Lesson Plan.

BLOCK	DAY 1 — FORM THE TEAM	DAY 2 — RUN THE CAMPAIGNS
Morning 1	9:00–10:30 Welcome • Topic 1 • Lab 1	9:00–10:10 Topic 3 • Lab 8
Morning 2	10:40–12:40 Labs 2–3	10:20–12:05 Labs 9–10 • recap
Afternoon 1	13:25–15:20 Labs 4–5	12:50–14:30 Topic 4 • Labs 11–12
Afternoon 2	15:30–18:00 Topic 2 • Labs 6–7 • review	14:40–16:00 Demo • summary • 16:00–18:00 Assessment

TIP Day 2 assessment: Written Assessment 4:00–5:00 PM, Practical Performance 5:00–6:00 PM.

Briefing for Assessment



Clear your space

Place phones and other materials under the table or on the floor.



No recording

No photos or recording of assessment materials.



Open book

Only slides, learner guide and approved materials are allowed.



Work individually

Complete the assessment on your own; no discussion.



Manage your time

Written Assessment then Practical Performance — watch the clock.



Appeals

An appeal process is available if you disagree with the outcome.

Assessment & Funding



Final Assessment

- Written Assessment (WA) — 1 hour, starts 4:00 PM
- Practical Performance (PP) — 1 hour
- Format: Open book — class ends 6:00 PM
- Allowed: slides, learner guide & approved materials only
- Appeal process available



Criteria for Funding

Minimum 75% attendance

Based on SSG digital attendance records.

Assessed as 'Competent'

Complete and pass both assessment components.

Courseware & Assessment on the LMS



Download the courseware

Get the slides and learner guide from the LMS for the open-book assessment.



Take the assessment

Complete the Written Assessment and Practical Performance on the LMS.



Log in

Use your enrolled account to access the course materials and assessment.

A screenshot of the Tertiary Infotech Academy LMS login interface. It features a dark blue background with a white logo at the top. Below the logo is the text 'Tertiary Infotech Academy' and 'LMS cum TMS for WSQ Courses'. There is an 'Email' input field, a 'Send OTP' button, and a link 'Login with password instead'. At the bottom, there are links for 'Privacy Policy', 'Acceptable Use Policy', and 'Feedback', and a note 'Powered by Tertiary Infotech Academy Pte Ltd'.

LMS: <https://lms-tms.tertiaryinfotech.com/>

Certification & TRAQOM Survey



TRAQOM Survey (Mandatory)

Complete the certification & TRAQOM survey to receive your certificate.

<https://lms-tms.tertiaryinfotech.com/>



Digital Attendance

Remember to take AM, PM and Assessment digital attendance — required for your funding.

Multi-AI-Agent Content Ideation and Digital Storyboarding

- 1 From one chatbot to a team of agents
- 2 A home for the team; Horizon's public website
- 3 Form the team: instructions, skills, connectors
- 4 Campaign 1: research, ideas and a storyboard

From Prompts to Agent Teams

Prompt Engineering

2023

ChatGPT goes mainstream. Craft the wording of one prompt.

Tools and MCP

2024

Models call tools. Anthropic open-sources the Model Context Protocol (Nov 2024).

Context Engineering

2024

Fill the context with exactly what the next step needs: files, memory, retrieval, skills.

Harness Engineering

2025

Agents ship inside a harness — Claude Code, Codex: the loop, tools and checks around the model.

Agent Teams

2026

Claude Code subagents and agent teams; open-source Hermes Agent.

One agent is a helper. A team with roles, a reviewer and a human gate is a marketing department.


Why a Team of Agents?

The Lead runs specialists in parallel, each in its own context, then combines what they return.

ONE AGENT DOES EVERYTHING


One agent

One context window,
filling up with every job



One job after another • one voice • marks its own homework

A MAIN AGENT DELEGATES



Faster

Research, copy and design run at the same time.

Focused

Each agent holds only its own job and its own skills.

Specialised

Its own instructions, tools, connectors and model.

Second opinion

The Lead checks work it did not write.

The cost

More tokens and more to review. Split only work that is truly independent.

The Agentic Loop

Every agent runs the same repeating cycle. A team is several loops, side by side.



TIP Goal in → context → action → verification → repeat → goal met.

Why the Loop Matters

Running a team of agents is steering loops, not typing commands.

Give direction, not micro-steps

State the goal, the inputs and "done when". Let each loop work out the steps.

Verification is everything

An agent is only as good as its checks: a facts sheet, a test, the Lead's review.

Steer mid-loop

Read the plan, interrupt, correct the course — do not wait for a wrong final answer.

Context is the fuel

What an agent can see decides what it can do: the brief, the data, the skills.

Claude and Claude Code

One family, two ways in. Chat and Cowork are now one product: Claude.

Claude Code — the team's home

Terminal, VS Code or the desktop Code tab. The Lead's session; subagent files; agent teams; skills; hooks; artifacts.

Claude — the team's apps

Desktop and web. Connectors (Drive, Gmail, Canva), skills, scheduled tasks, artifacts.

Hermes Agent • optional

Open source (MIT). The same roles on any model and a Kanban board. A demo at the end.

Claude Plans and Pricing

Claude Code is included in every paid plan. Prices in US dollars, checked 8 Oct 2026.

PLAN	PRICE	FOR THIS COURSE
Free	\$0	Chat only — no Claude Code, no artifact publishing
Pro	\$17/month billed annually · \$20 monthly	Enough for every lab
Max	From \$100/month — 5× or 20× Pro usage	Heavy daily use of agent teams
Team	Standard \$20/seat/month annual (\$25 monthly) · Premium \$100 (\$125)	A marketing team sharing skills and artifacts
Enterprise	\$20/seat/month + usage at API rates	Admin controls, retention, sharing policy
API	Per million tokens in/out: Opus 5.5 \$4/\$20 · Sonnet 5.5 \$2/\$10 · Fable 5.1 \$10/\$50	Pay as you go — e.g. Hermes

TIP Agent teams use many more tokens than one session. Check current prices at claude.com/pricing.

Get the Tools Ready

Install once; every lab says which tool to open.

DOWNLOAD

claude.com/download

Claude desktop

macOS and Windows — Claude and Claude Code

Claude Code CLI

For agent teams in a terminal

Claude Code and artifacts need a Pro, Max, Team or Enterprise plan.

1

Claude desktop Install from claude.com/download and sign in on a paid plan.

2

Claude Code CLI `curl -fsSL https://claude.ai/install.sh | bash` (Windows: `irm https://claude.ai/install.ps1 | iex`).

3

Sign in with /login Artifacts need your claude.ai account, not an API key.

4

Node.js 22, Python 3, ffmpeg For the publisher kit and the video (Pillow for Python).

5

A personal Google account For Drive and Gmail — never an employer's.

A Home for the Team

Before you hire agents, give them shared rules, shared facts and a website to point to.

CLAUDE.md

The team charter every agent reads at the start of every session. Under 60 lines.

data/

The facts sheet, brand, compliance checklist and business brief — the only sources agents may use.

Horizon's public website

Given, not built: one public link every learner opens, with a GitHub Pages backup if artifacts or Sites are blocked.

Set Up the Studio and Explore Horizon's Website

LAB 01

DAY 1 • ~30 MIN

THE STORY SO FAR

October 2026. Rachel has a website and a goal: 40 booked chats a month by March. She has no marketing team. Before you hire one made of agents, give it a home — shared rules, shared facts — and put the website online.

YOU'LL BUILD

horizon-studio/ with CLAUDE.md, the shared data, a local copy of the site, and the public site link in the charter

Surface: Claude Code

Features: CLAUDE.md as the team charter • Horizon's public website

WEBSITE claude.ai/artifact/J82VCbzFogYps2vep8UYhL **BACKUP** tertiarycourses.github.io/TGS-2023036153-Multi-AI-Agents-Workflow-for-Content-Creation

Set Up the Studio and Explore Horizon's Website

Claude Code • full step-by-step with reasoning is in your Learner Guide.

- 1 Make the studio** Create horizon-studio/ with data/ and web/site/. Copy the assets into data/ and horizon-site.html to web/site/index.html.
- 2 Start Claude Code** In a terminal in horizon-studio run `claude`, then `/login` with your `claude.ai` account. Trust the folder.
- 3 Write the charter** Paste Prompt A. Read `CLAUDE.md` — every agent will read it at the start of every session.
- 4 Open the website** Open <https://claude.ai/artifact/J82VCbzFqYps2vep8UYhL> — Horizon's public website, the same for every learner. You do not build or publish it. If it will not open, use the backup: <https://tertiarycourses.github.io/TGS-2023036153-Multi-AI-Agents-Workflow-for-Content-Creation/>
- 5 Know the site** Paste Prompt B. Then open the link on your phone and check the calculator shows S\$693,138.

Set Up the Studio and Explore Horizon's Website

PROMPT A — Claude Code: the charter

Read data/team-charter.md and data/business-brief.md. Write CLAUDE.md for this studio in four sections — What this is, The team, Rules, Files — under 60 lines. Rules must include:
every number comes from data/facts-2026.md; every piece passes fact-check and fin-compliance; nothing is published or sent without a person's approval.

Set Up the Studio and Explore Horizon's Website (cont.)

PROMPT B — Claude Code: know the site

Read web/site/index.html, the local copy of Horizon's public website. In five bullets: who it is for, the offer, the main call to action, what the calculator does, and the disclaimer. Then add this line to the Files section of CLAUDE.md:
Public site (link in every post): <https://claude.ai/artifact/J82VCbzFoqYps2vep8UYhL>

NOTE Shared context first, agents second. Every agent you add later starts from these files, so a wrong rule here is wrong everywhere.

Set Up the Studio and Explore Horizon's Website

Check your work before moving on.

CHECK YOUR WORK

- ✓ horizon-studio/ holds data/ (six files) and web/site/index.html.
- ✓ CLAUDE.md has the four sections and is under 60 lines.
- ✓ Its rules name the facts sheet, the two checks and the human approval.
- ✓ CLAUDE.md names the public site link every post points to.
- ✓ The public site opens on a phone; the calculator shows S\$693,138.

Anatomy of an Agent

A Markdown file in `.claude/agents/`. The same file works as a subagent or as an agent-team teammate.

`.claude/agents/content-creator.md`

```
---
name: content-creator
description: Use for Horizon copy – posts, newsletter,
  video scripts and page copy, in the Sunny voice.
tools: Read, Write, Edit, Grep, Glob
model: sonnet
skills:
  - sunny-voice
  - channel-formats
---
You are Horizon Wealth Planning's Content Creator.
## What you read    the brief, facts-2026.md, brand.md
## What you deliver content files, ready for review
## Connectors      Google Drive (docs), Gmail (drafts only)
## Never           add a number not on the facts sheet; send;
                  run approve.mjs or publish --live
## Report back     what, where, anything unverified
```

Good Agent Instructions

Five parts, every time. The "Never" part matters most in a regulated firm.

The job

One sentence: what this agent is for. The description decides when Claude delegates to it.

Reads and delivers

Which files it may use; what it hands back, and where.

Never

Its hard limits: no invented figures, no sending, no publishing, no approving.

Report back

Three lines: what, where, anything unverified — so the Lead can combine results.

Form the Team: Instructions for Each Agent

LAB 02

DAY 1 • ~45 MIN

THE STORY SO FAR

Rachel sketches the team on a whiteboard: a Lead and five specialists. Each needs a clear job, and a clear line it must never cross — this is a regulated firm.

YOU'LL BUILD

.claude/agents/: growth-strategist, content-creator, creative-designer, website-designer, growth-analyst

Surface: Claude Code (subagent files)

Features: five subagent files in .claude/agents/ • name, description, tools, model • the Lead is your session

Form the Team: Instructions for Each Agent

Claude Code (subagent files) • full step-by-step with reasoning is in your Learner Guide.

- 1 Read the roles** Open role-specs.md: one Lead and five specialists, each with a job, skills, connectors and limits.
- 2 Write the instructions** Paste Prompt A.
- 3 Read one closely** Open .claude/agents/content-creator.md. Is the "Never" section strict enough for a regulated firm?
- 4 Meet the team** Paste Prompt B. Each agent introduces itself in one line — in its own context.
- 5 Tighten one** Change one instruction you disagree with and save it. It is your team.

Form the Team: Instructions for Each Agent

PROMPT A — Claude Code: five subagents

Read data/role-specs.md. Create five project subagents in .claude/agents/ — growth-strategist, content-creator, creative-designer, website-designer and growth-analyst. Each file has:

- frontmatter: name; a description that says when to use it; tools; model (opus for the strategist, sonnet for the others)
- a body with: the job, what it reads, what it delivers, the connectors it may use, what it must never do, and how it reports back. No skills line yet — we add skills in Lab 3.

Form the Team: Instructions for Each Agent (cont.)

PROMPT B — Claude Code: meet the team

Use each of the five subagents once, in parallel. Ask each: "In one line, who are you, and what is the first thing you would do for Horizon's goal of 40 booked chats a month?" Show me the five answers in a table.

NOTE Clear jobs beat clever prompts. An agent that knows what it must never do is safer than one that is merely told to be careful.

Form the Team: Instructions for Each Agent

Check your work before moving on.

CHECK YOUR WORK

- ✓ Five files exist in `.claude/agents/`, one per specialist.
- ✓ Each has a name, a "use when" description, tools and a model.
- ✓ Your session acts as the Marketing Team Lead.
- ✓ Five agents answered, each in its own context.
- ✓ Every agent has a "never" rule that blocks publishing.
- ✓ You changed one instruction and can say why.

Twelve Skills, Owned by Role

Instructions say who an agent is; skills say how it does the job.

AGENT	SKILL 1	SKILL 2
Team Lead	fact-check	fin-compliance
Growth Strategist	competitor-scan	campaign-brief
Content Creator	sunny-voice	channel-formats
Creative Designer	brand-visuals	video-render
Website Designer	landing-page	page-qa
Growth Analyst	audience-insights	campaign-report

TIP Master copies in skills/; .claude/skills/ for Claude Code; Customize → Skills for Claude.

Anatomy of a Skill

Frontmatter (the trigger), then short steps. Free until it is needed.

```
.claude/skills/campaign-brief/SKILL.md
```

```
---
name: campaign-brief
description: Use when planning a Horizon campaign.
  Objective, audience, message, proof, channels,
  CTA, measures, and a five-beat storyboard.
---
# Campaign brief
1. Objective: one measurable action (e.g. chats booked).
2. Audience: the persona(s) from strategy/personas.md, with evidence.
3. Core message in one sentence; proof from data/facts-2026.md only.
4. Channels and formats; one call to action; measures and a date.
5. Storyboard in five beats – hook, tension, proof, resolution, action –
  each with visual, words, claim and source.
6. Never: returns, fee amounts, invented clients, fear.
```

Give Each Agent Its Skills

LAB 03

DAY 1 • ~50 MIN

THE STORY SO FAR

The team exists, but every agent writes in its own way. Rachel wants Horizon's voice, facts and compliance rules applied the same way by every agent, every time.

YOU'LL BUILD

twelve skills in `.claude/skills/`, wired to the agents that own them

Surface: Claude Code (`.claude/skills/`) → Claude

Features: twelve skills: two per specialist, two for the Lead • `SKILL.md` • the "Use when" trigger • test without naming

Give Each Agent Its Skills

Claude Code (.claude/skills/) → Claude · full step-by-step with reasoning is in your Learner Guide.

- 1 Read the spec** skills-spec.md lists the twelve skills and who owns each.
- 2 Create the skills** Paste Prompt A.
- 3 Wire them up** Paste Prompt B — each agent now knows its two skills; the Lead knows fact-check and fin-compliance.
- 4 Test without naming** In a new session ask: "Write a Facebook post about the free Money Check-up Checklist." Ask which skills were used.
- 5 Break it on purpose** Ask for a post that says "guaranteed 8% a year". fin-compliance must flag it as Critical.
- 6 Share with Claude** Upload sunny-voice, channel-formats and fin-compliance under Customize → Skills in Claude too.

Give Each Agent Its Skills

PROMPT A — Claude Code: twelve skills

Read data/skills-spec.md. Create the twelve skills in .claude/skills/<name>/SKILL.md. Each has frontmatter (name, and a description that starts "Use when...") and short numbered steps, under 80 lines. Use data/brand.md, data/channel-formats.md, data/facts-2026.md and data/compliance-checklist.md as the sources — do not invent rules.

Give Each Agent Its Skills (cont.)

PROMPT B — Claude Code: wire them up

Add a skills: list to each subagent in `.claude/agents/` with its two skills from `data/skills-spec.md`. Add to `CLAUDE.md` that the Lead (this session) runs fact-check and fin-compliance on everything before showing it to me.

NOTE A skill costs almost nothing until it is needed. Keep the rules file (`CLAUDE.md`) short and put procedures in skills.

Give Each Agent Its Skills

Check your work before moving on.

CHECK YOUR WORK

- ✓ Twelve folders in `.claude/skills/`, each with a `SKILL.md`.
- ✓ Every description starts "Use when...".
- ✓ Each specialist lists its own two skills.
- ✓ The untold test picked sunny-voice and channel-formats.
- ✓ fin-compliance flagged "guaranteed 8%" as Critical.
- ✓ Three skills also appear in Claude under Customize → Skills.

Connectors for Each Agent

Connect once in claude.ai; Claude Code picks them up (/mcp).

AGENT	MAY USE	LIMIT
Team Lead	Google Drive (read), Gmail (drafts only)	reads; drafts only
Growth Strategist	web search and fetch, Google Drive (read)	cite every source
Content Creator	Google Drive (docs), Gmail (drafts only)	Gmail: drafts only
Creative Designer	Canva, Bash for Pillow and ffmpeg	no fake "clients"
Website Designer	artifacts (publish), Playwright MCP (browser check)	publish after approval
Growth Analyst	Google Drive and Sheets (read), Bash for Python	no personal data

TIP Use your personal Google account — never an employer's.

Least Privilege

A regulated firm gives each agent only what its job needs.

Ask of each connector

Does this agent's job need it? If not, it does not get it.

Drafts, not sends

Email is drafted, posts are files, pages wait — a person sends and publishes.

Write it down

Each agent's instructions name its connectors and limits; its tools line removes what it does not need.

Connect the Tools Each Agent Needs

LAB 04

DAY 1 • ~45 MIN

THE STORY SO FAR

The analyst cannot read the data, the designer has no design tool and the creator cannot draft an email. Connect each agent to what its job needs — and nothing more.

YOU'LL BUILD

every specialist tested on its own connector, with its limits written into its instructions

Surface: claude.ai connectors → Claude Code (/mcp)

Features: Google Drive • Gmail (drafts) • Canva • Playwright MCP • least privilege per agent

Connect the Tools Each Agent Needs

claude.ai connectors → Claude Code (/mcp) • full step-by-step with reasoning is in your Learner Guide.

- 1 Connect in claude.ai** Settings → Connectors: Google Drive, Gmail and Canva. Use your personal account.
- 2 Check in Claude Code** Run /mcp. The claude.ai connectors appear in the list.
- 3 Add a browser** Run: `claude mcp add playwright -s project -- npx -y @playwright/mcp@latest`
- 4 Put the data on Drive** Upload data/ to a Drive folder named Horizon Studio.
- 5 Test each agent** Paste Prompt A — one small job per agent, each on its own connector.
- 6 Write the limits** Paste Prompt B.

Connect the Tools Each Agent Needs

PROMPT A — Claude Code: test each agent

Run these five checks, each by its own subagent:

1. growth-strategist: web search — one 2026 news item about financial planning in Singapore, cited.
2. growth-analyst: Google Drive — read firm-metrics.csv from Horizon Studio; report average monthly enquiries.
3. content-creator: Gmail — a DRAFT to me with the subject "Connector test". Do not send.
4. creative-designer: Canva — list my designs, or create a 1080x1080 test design in Horizon colours.
5. website-designer: Playwright — open Horizon's public site at 375px wide and report any problem. Report each result in one line.

Connect the Tools Each Agent Needs (cont.)

PROMPT B — Claude Code: least privilege

Update the "Connectors" section of each subagent in `.claude/agents/` to match `data/connectors-setup.md`: what it may use, and that Gmail is drafts-only. Tighten each tools line so no agent has a tool its job does not need.

NOTE Least privilege is a design decision, not a setting. Ask of each connector: does this agent's job need it?

Connect the Tools Each Agent Needs


Check your work before moving on.

CHECK YOUR WORK

- ✓ /mcp lists Google Drive, Gmail, Canva and Playwright.
- ✓ Five checks ran, each by its own agent.
- ✓ The Gmail test is a draft in your inbox — nothing was sent.
- ✓ The analyst read the CSV from Drive, not from the local folder.
- ✓ Each agent's instructions list only the tools its job needs.


How the Team Runs a Campaign

Campaign 1, "Know Your Number": research in parallel, a person chooses, the team makes it concrete.



The Five-Beat Storyboard

Every channel tells the same story with the same claims.



Financial Content Is Regulated

The Lead checks every piece against twelve rules (compliance-checklist.md).

RULE	WHAT IT MEANS	IF BROKEN
C1 • Not advice	General information only — say so	Critical
C2 • No promises	No guaranteed or "risk-free" returns	Critical
C3 • Label illustrations	Assumptions shown; not guaranteed	Critical
C5 • No performance	Never client returns or "average" returns	Critical
C9 • Consent	Marketing only to people who opted in	Critical
C11 • Facts sheet	Every figure from facts-2026.md	Critical

TIP MAS Guidelines on Standards of Conduct for Digital Advertising Activities apply from 25 March 2026.

Campaign 1: Research to Storyboard, with Subagents

LAB 05

DAY 1 • ~50 MIN

THE STORY SO FAR

The team's first campaign: 'Know Your Number', November to January. Rachel wants research, ideas for three audiences, her choice, and a storyboard every channel follows.

YOU'LL BUILD

research/ (competitors, funnel), strategy/ideas.md, strategy/storyboard.md with three channel variants

Surface: Claude Code (subagents in parallel)

Features: the Lead delegates • strategist and analyst in parallel • nine scored ideas • a human concept checkpoint • a five-beat storyboard

Campaign 1: Research to Storyboard, with Subagents

Claude Code (subagents in parallel) • full step-by-step with reasoning is in your Learner Guide.

- 1 Brief the Lead** Copy the assets into data/. You are the client; your session is the Marketing Team Lead.
- 2 Research in parallel** Paste Prompt A. Watch the strategist and the analyst run at the same time.
- 3 Read the ideas** Nine ideas, scored with the rubric. The Lead stops for you.
- 4 Choose** Pick one idea and say why — the human concept checkpoint.
- 5 Storyboard** Paste Prompt B with your choice.
- 6 Check the claims** Every beat's claim must have a source in facts-2026.md.

Campaign 1: Research to Storyboard, with Subagents

PROMPT A — Claude Code: research and ideas

Act as the Marketing Team Lead for the campaign in data/campaign-objective.md. Use subagents, in parallel:

- growth-strategist: competitor-scan → save research/competitive-analysis.md.
- growth-analyst: the funnel in data/firm-metrics.csv — what 40 chats a month needs → save research/funnel.md.

When both finish, the growth-strategist proposes three ideas each for young professionals, mid-career families and pre-retirees, scored with data/idea-rubric.md → strategy/ideas.md. Show me the top five and stop — I will choose.

Campaign 1: Research to Storyboard, with Subagents (cont.)

PROMPT B — Claude Code: the storyboard

I choose idea <number>, because <one line>. Use subagents:

- growth-strategist: campaign-brief and a five-beat storyboard with data/storyboard-template.md → strategy/storyboard.md.
- content-creator: three channel variants that keep the same claims — LinkedIn, Facebook, a 60-second video. Then run fact-check and fin-compliance yourself and show me the issues before anything is final.

NOTE Agents diverge; a person decides. The concept checkpoint is where the campaign becomes yours.

Campaign 1: Research to Storyboard, with Subagents

Check your work before moving on.

CHECK YOUR WORK

- ✓ The strategist and the analyst ran in parallel.
- ✓ research/ has the competitor analysis and the funnel, cited.
- ✓ strategy/ideas.md holds nine scored ideas with their authors.
- ✓ You chose the idea and recorded why.
- ✓ The storyboard has five beats; every claim has a source.
- ✓ The Lead's fact-check and fin-compliance report came back.

Where You Are Now

The team exists and has run its first campaign. But who exactly is it for — and how often should it publish?

Formed • Labs 1–4

A studio, the public website, six agents with instructions, skills and connectors.

Campaign 1 • Lab 5

Research in parallel, nine scored ideas, your choice, a storyboard.

What is missing


Evidence about the audience and a cadence the team can deliver. That is Topic 2.

Audience Research and Content Requirement Analysis

- 1 The Growth Analyst: personas from evidence
- 2 Content requirements every agent can follow
- 3 How often to publish: value, fatigue and capacity
- 4 The calendar the team will run

Evidence Before Personas

Lab 6: the Growth Analyst builds personas the way an analyst would.



Observed, Inferred or Low Confidence?

Three labels keep a persona honest.

LABEL	EXAMPLE	USE IT
Observed	30–39s chose Facebook, LinkedIn and YouTube equally (survey, n = 44)	As fact, with the count
INFERENCE	"They are short of time on weekdays" — not asked	As a hypothesis to test
LOW CONFIDENCE	The 60+ group (n = 12)	Never alone; get more data
Verbatim	"I am 54 and not sure if my CPF will be enough." (E-id)	Exactly as written, with the id

TIP Ask "which rows show that?" of anything that sounds insightful.

Growth Analyst: Personas from Evidence

LAB 06

DAY 1 • ~40 MIN

THE STORY SO FAR

Jun Wei drafted personas from instinct: 'busy young parents who love TikTok'. Rachel is not convinced. There are 150 survey answers, 60 enquiries and 200 checklists.

YOU'LL BUILD

strategy/personas.md and strategy/content-spec.md

Surface: Claude Code (analyst subagents)

Features: three analyst subagents, one per data file • counts, not impressions • confidence labels • content requirements

Growth Analyst: Personas from Evidence

Claude Code (analyst subagents) • full step-by-step with reasoning is in your Learner Guide.

- 1 Add the data** Copy the CSVs and templates into data/.
- 2 Analyse in parallel** Paste Prompt A — three analyst subagents, one per file.
- 3 Challenge a claim** Pick one persona statement and ask: "which rows show that?"
- 4 Write the requirements** Paste Prompt B.
- 5 Check the gaps** The 60+ group must be LOW CONFIDENCE.

Growth Analyst: Personas from Evidence

PROMPT A — Claude Code: three analysts

Use three growth-analyst subagents in parallel, one per file: data/survey-responses.csv, data/enquiries.csv, data/checklist-results.csv. Each uses audience-insights: counts with code, n and the file on every claim, INFERENCE and LOW CONFIDENCE labels, enquiry quotes verbatim with ids. Combine the results into three personas with data/persona-template.md → strategy/personas.md.

Growth Analyst: Personas from Evidence (cont.)

PROMPT B — Claude Code: requirements

Use the growth-strategist and the content-creator subagents to write strategy/content-spec.md with data/content-spec-template.md: for each persona the top 3 questions, channels, formats, length, tone, proof, call to action and never-say — each with the evidence (file and count) behind it.

NOTE Ask "which rows show that?" of anything that sounds insightful. If the agent cannot point to rows, it is an inference.

Growth Analyst: Personas from Evidence

Check your work before moving on.

CHECK YOUR WORK

- ✓ Three analyst subagents ran, one per file.
- ✓ Every persona statement carries a count and a file.
- ✓ The 60+ group is LOW CONFIDENCE (n = 12).
- ✓ The most-unticked checklist item is the retirement number.
- ✓ Enquiry quotes are verbatim, with ids.
- ✓ content-spec.md has every field for every persona.

Finding the Right Frequency

Lab 7: the Analyst finds where more posts stop paying; the Strategist plans four weeks.



TIP A third LinkedIn post a week brings more unfollows than value. Weekly emails make unsubscribes jump.

Growth Analyst and Strategist: Cadence and Calendar

LAB 07

DAY 1 • ~35 MIN

THE STORY SO FAR

Jun Wei has 12 hours a week and Rachel needs two business days to review anything. How often should each channel run, and what goes out when?

YOU'LL BUILD

strategy/cadence.md and strategy/calendar.csv

Surface: Claude Code (subagents)

Features: marginal value vs fatigue • team capacity • review buffers • a four-week calendar

Growth Analyst and Strategist: Cadence and Calendar

Claude Code (subagents) • full step-by-step with reasoning is in your Learner Guide.

- 1 Add the data** Copy the assets into data/.
- 2 Find the sweet spot** Paste Prompt A.
- 3 Question it** Why is the newsletter not weekly? The answer is in the unsubscribes.
- 4 Plan four weeks** Paste Prompt B.
- 5 Check the buffers** Every review date is two business days ahead; nothing on a holiday.

Growth Analyst and Strategist: Cadence and Calendar

PROMPT A — Claude Code: the analyst

Use the growth-analyst subagent. From data/channel-benchmarks.csv and data/team-capacity.md, recommend how often to publish on LinkedIn, Facebook, the newsletter and YouTube. Compare the extra clicks from one more post, the unfollows or unsubscribes it costs, and the hours. Show the numbers → strategy/cadence.md.

Growth Analyst and Strategist: Cadence and Calendar (cont.)

PROMPT B — Claude Code: the strategist

Use the growth-strategist subagent to build a four-week calendar from Monday 2 November 2026 in strategy/calendar.csv (columns as data/calendar-format.csv), from the cadence, strategy/storyboard.md and strategy/content-spec.md. review_by two business days before each date; no public holidays; no week over the team's hours. Show a week-by-week summary with hours.

NOTE Frequency is a capacity decision as much as an audience one. A calendar the team cannot deliver is fiction.

Growth Analyst and Strategist: Cadence and Calendar

Check your work before moving on.

CHECK YOUR WORK

- ✓ cadence.md gives a frequency per channel with the evidence.
- ✓ A third LinkedIn post a week costs far more unfollows than it earns.
- ✓ It explains the unsubscribe jump at weekly emails.
- ✓ calendar.csv covers four weeks in the required columns.
- ✓ Every review_by date is two business days before publishing.
- ✓ No week is over 12 hours.

Day 1: The Team Is Ready

Everything the campaigns need tomorrow is in place.



Multi-Channel Content Creation and Agent Workflow Coordination


- 1 Subagents or an agent team?
- 2 Human in the loop — enforced, not requested
- 3 Campaigns 2–4: social week, newsletter + landing page, video
- 4 From approval to LinkedIn, Facebook, email and YouTube

Subagents vs Agent Teams

Both run agents in parallel, each with its own context. The difference is who they talk to.

SUBAGENTS


Report back to one main agent



Each works in its own context and returns only its result. They do not talk to each other.

AGENT TEAM

A lead, a shared task list, teammates who talk



Teammates claim tasks and message each other directly — e.g. a reviewer sends fixes straight to the writer.

Use subagents when only the result matters. Use a team when the agents need to talk.

Which Pattern When?


From the Claude Code documentation, applied to the team.

	SUBAGENTS	AGENT TEAM
Talks to	Returns a result to the Lead	Teammates message each other directly
Coordination	The Lead manages everything	A shared task list; teammates claim tasks
Best for	Research, analysis, one-off reviews	Copy ↔ design ↔ review loops
Cost	Lower: results are summarised back	Higher: every teammate is its own Claude session
Campaigns	Labs 5, 6, 7, 12	Labs 8, 9, 10

TIP Agent teams are experimental: turn them on with `CLAUDE_CODE_EXPERIMENTAL_AGENT_TEAMS=1`.

Running an Agent Team

Claude Code in a terminal. Describe the team in plain words.



TIP Teammates load CLAUDE.md, skills and connectors — but not the Lead's conversation. One team per session.

Human in the Loop, Enforced

A rule in a prompt persuades. The publisher kit builds a gate that holds.

Only people approve

approve.mjs refuses to run without an interactive terminal — an agent's tool call cannot approve.

The approved words go out

The approval stores a hash. Edit one word and publish --live refuses until a person approves again.

A hook stops agents

A PreToolUse hook blocks agents from running approve.mjs or publishing unapproved items.

The Hook That Holds the Gate

Claude Code runs it before every Bash, Edit or Write. Exit code 2 blocks the call and tells the agent why.

`.claude/settings.json`

```
{
  "env": { "CLAUDE_CODE_EXPERIMENTAL_AGENT_TEAMS": "1" },
  "hooks": {
    "PreToolUse": [{
      "matcher": "Bash|Edit|Write|MultiEdit",
      "hooks": [{
        "type": "command",
        "command": "node \"$CLAUDE_PROJECT_DIR/scripts/gate-hook.mjs\""
      }]
    }]
  }
}
```

The Publishing APIs


publish.mjs is a dry run by default. Tokens live only in .env.

CHANNEL	ENDPOINT	PERMISSION
LinkedIn	POST api.linkedin.com/rest/posts (LinkedIn-Version: YYYYMM)	w_member_social — your own profile
Facebook	POST graph.facebook.com/<version>/<page-id>/feed	Page token: pages_manage_posts
YouTube	Resumable upload to youtube/v3/videos, privacy private	youtube.upload; unaudited projects stay private
Newsletter	No API — HTML + text files, then Gmail drafts	Consent = yes, unsubscribed = no

TIP No accounts? Every lab works in dry-run mode: you see the exact request that would be sent.

Campaign 2: A Social Media Week

Lab 8: two LinkedIn and three Facebook posts, each with a visual.



Campaign 2: A Social Media Week

LAB 08

DAY 2 • ~50 MIN

THE STORY SO FAR

Day 2. Week 1 of the campaign starts Monday: two LinkedIn posts and three Facebook posts, each with a visual — and Rachel's one rule: nothing goes out unless a person approved that exact version.

YOU'LL BUILD

content/social/week-01/: five posts with images, reviewed, approved by you and posted (or dry-run)

Surface: Claude Code (agent team) → LinkedIn, Facebook

Features: the publisher kit and its hook • an agent team: creator, designer, Lead • teammates message each other • human approval → post

Campaign 2: A Social Media Week

Claude Code (agent team) → LinkedIn, Facebook · full step-by-step with reasoning is in your Learner Guide.

- 1 Install the kit** Copy publisher-kit/scripts into scripts/ and merge its settings.json into .claude/settings.json (agent teams on, the approval hook). Restart Claude Code.
- 2 Test the gate** Ask Claude to run approve.mjs on anything. The hook must block it.
- 3 Start the team** Paste Prompt A. Teammates appear in the panel; ↑ ↓ and Enter open one; Ctrl+T shows the task list.
- 4 Watch them talk** The designer asks the creator for each post's hook; the Lead sends fixes back.
- 5 Approve as a person** Read each post and image. In your own terminal: `node scripts/approve.mjs <id> --by "Your Name"`.
- 6 Publish** Paste Prompt B. Dry run first; --live only if your accounts are connected.

Campaign 2: A Social Media Week

PROMPT A — Claude Code: the social team

Create an agent team for Week 1 social from data/social-brief.md and strategy/calendar.csv. Spawn two teammates using the agent types content-creator and creative-designer; you are the Lead.

- content-creator: 2 LinkedIn + 3 Facebook posts, one file each in content/social/week-01/, with the frontmatter in data/publishing-spec.md.
- creative-designer: one image per post (Canva or a rendered card) with alt text; message the creator for each post's hook before designing.
- You: fact-check and fin-compliance on every post; message fixes to the teammate who owns it. When a post passes, run node scripts/submit.mjs on it. Publish nothing.

Campaign 2: A Social Media Week (cont.)

PROMPT B — Claude Code: publish

For every Week 1 post approved in review/approvals.csv, run node scripts/publish.mjs <id> as a dry run and show me each request. Wait for me to say "go" before any --live run.

NOTE Subagents report to one boss; teammates also talk to each other. A designer who can ask the writer for the hook makes a better image.

Campaign 2: A Social Media Week

Check your work before moving on.

CHECK YOUR WORK

- ✓ The hook (or approve.mjs) stopped the agent from approving.
- ✓ Two teammates ran; the task list showed their tasks.
- ✓ The designer and the creator messaged each other.
- ✓ Five posts with images, alt text, frontmatter and UTM links.
- ✓ You approved in your own terminal; the rows show your name and a hash.
- ✓ Dry runs printed every request; nothing went live unless you said go.

Campaign 3: Newsletter + Landing Page

Lab 9: three specialists, one email and one page, consent first.



Consent Decides Who Gets Mail

The PDPA, applied by an agent that counts.

Opt-in only

consent = yes and unsubscribed = no.
Everyone else goes to excluded.csv with a reason.

Every email carries

An unsubscribe link, Horizon's address
and the short disclaimer.

The page collects little

Name, email, a required consent box, a
separate unticked newsletter opt-in.

Campaign 3: Newsletter and a Landing Page

LAB 09

DAY 2 • ~45 MIN

THE STORY SO FAR

The Sunny Sunday email books chats most cheaply. November's issue promotes a new landing page for the free checklist — and must reach only the 31 people who opted in.

YOU'LL BUILD

content/newsletter/2026-11/ (issue, HTML, text, recipients, excluded), web/checklist/ published, Gmail drafts

Surface: Claude Code (agent team) → Claude artifact, Gmail

Features: creator writes, website designer builds the lead-magnet page, analyst filters by consent, the Lead checks · drafts only

Campaign 3: Newsletter and a Landing Page

Claude Code (agent team) → Claude artifact, Gmail · full step-by-step with reasoning is in your Learner Guide.

- 1 Start the campaign** Paste Prompt A.
- 2 Approve the words** Read issue.md and the page copy; approve nl-2026-11 in your own terminal.
- 3 Publish the page** The website designer runs page-qa and publishes web/checklist/ as a Claude artifact.
- 4 Check the list** recipients.csv: only consent = yes and unsubscribed = no.
- 5 Create the drafts** Paste Prompt B.
- 6 Send one test** Send only the draft addressed to yourself, and read it on your phone.

Campaign 3: Newsletter and a Landing Page

PROMPT A — Claude Code: the newsletter team

Create an agent team for the November Sunny Sunday email (data/newsletter-brief.md, data/newsletter-spec.md) and its landing page (data/landing-page-brief.md). Spawn three teammates: content-creator, website-designer and growth-analyst.

- content-creator: issue.md (id nl-2026-11) and the landing-page copy.
- website-designer: web/checklist/index.html with landing-page from the approved copy; page-qa.
- growth-analyst: recipients.csv and excluded.csv from data/subscribers.csv. You check everything, then submit nl-2026-11. Build newsletter.html and .txt, and publish the page as an artifact, only after I approve.

Campaign 3: Newsletter and a Landing Page (cont.)

PROMPT B — Claude Code: Gmail drafts

Using the Gmail connector, create a DRAFT of newsletter.html for each address in recipients.csv, plus one draft addressed to me only. Send nothing. Report how many drafts and how many people were excluded, and why.

NOTE Consent decides who gets mail — not the agent and not the deadline. Marketing to people who did not opt in breaks the PDPA.

Campaign 3: Newsletter and a Landing Page

Check your work before moving on.

CHECK YOUR WORK

- ✓ Three specialists worked on the campaign; the Lead checked all.
- ✓ You approved nl-2026-11 in your own terminal before the build.
- ✓ The landing page passed page-qa and is published.
- ✓ recipients.csv holds 31 people; excluded.csv lists 9 with reasons.
- ✓ The email has unsubscribe, address and disclaimer.
- ✓ Drafts only — nothing was sent to the list.

Campaign 4: The YouTube Explainer

Lab 10: the team writes and designs; you approve twice.

Strategist and Creator

The angle, then 5–7 scenes with words on screen and a description with chapters and an AI disclosure.

Creative Designer

Storyboard, thumbnail, then the render with video-render — ffmpeg, 1920×1080, under 90 s.

Upload privately

Approve the words, then the finished video. It uploads as private; a person makes it public.

Campaign 4: The YouTube Explainer

LAB 10

DAY 2 • ~45 MIN

THE STORY SO FAR

Pre-retirees say they would rather watch than read. Rachel wants the first 'Money in Plain English' video, uploaded privately until she has watched it.

YOU'LL BUILD

content/video/ep01/: scenes.json, thumbnail, ep01.mp4, ep01.md, and a private YouTube upload (or dry run)

Surface: Claude Code (agent team, render, upload)

Features: strategist angle • creator script • designer storyboard, thumbnail and render • approve twice • private upload

Campaign 4: The YouTube Explainer

Claude Code (agent team, render, upload) • full step-by-step with reasoning is in your Learner Guide.

- 1 Script it as a team** Paste Prompt A.
- 2 Approve the script** Read scenes.json and ep01.md; approve yt-ep01-script in your own terminal.
- 3 Render it** Paste Prompt B. The designer runs video-render.
- 4 Approve the video** Submit ep01.md, watch ep01.mp4 to the end, then approve yt-ep01.
- 5 Upload** Paste Prompt C. It goes up as private.
- 6 Disclose** In YouTube Studio, answer the altered or synthetic content question before making it public.

Campaign 4: The YouTube Explainer

PROMPT A — Claude Code: the video team

Create an agent team for YouTube episode 1 from data/youtube-brief.md. Spawn three teammates: growth-strategist (the angle and hook), content-creator (content/video/ep01/scenes.json, 5-7 scenes, and ep01.md with title, tags, description, chapters and the AI disclosure) and creative-designer (a storyboard of the scenes and a thumbnail, using brand-visuals). You fact-check and fin-compliance the script, then submit it as yt-ep01-script. Publish nothing.

Campaign 4: The YouTube Explainer (cont.)

PROMPTS B and C — Claude Code: render and upload

PROMPT B

Use the creative-designer subagent with video-render on the approved scenes.json → content/video/ep01/ep01.mp4 (data/video-spec.md). Check it with ffprobe. Do not change approved words.

PROMPT C

Run node scripts/publish.mjs yt-ep01 as a dry run. If I say "go", run it with --live: it uploads as PRIVATE. Give me the YouTube Studio link.

NOTE Approve twice: once for the words, once for the finished video. Rendering can break what the script got right.

Campaign 4: The YouTube Explainer

Check your work before moving on.

CHECK YOUR WORK

- ✓ The strategist, creator and designer each delivered their part.
- ✓ You approved the script, then the finished video.
- ✓ ep01.mp4 is 1920x1080, under 90 seconds, every scene present.
- ✓ Every figure on screen matches facts-2026.md.
- ✓ The description has chapters, link, disclaimer and AI disclosure.
- ✓ The upload is private (or the dry run printed the request).

Where You Are Now

Three campaigns, every piece checked by the Lead and approved by a person. But the team stops when you close the laptop.

Gate · Lab 8

People approve, hashes hold, a hook stops agents.

Campaigns · Labs 8–10

Social week, newsletter + landing page, the video.

What is missing

A team that works every week on its own — and rules to govern it. That is Topic 4.

Content Distribution, Strategy Guidelines and Responsible AI Practices

- 1 Campaign 5: always on, with a scheduled task
- 2 Responsible AI for financial content
- 3 Measure the patterns; write the team playbook
- 4 Optional demo: the team on Hermes Agent

From Skill to Schedule

Lab 11: do the job once, keep it as a skill, let a schedule run it.



TIP Pause scheduled tasks you are not using.

Automate the Preparation, Not the Decision

What a scheduled agent should and should not do for a regulated firm.

Prepare

Read the data, run the skill, draft the report, flag what needs a decision.

Draft, never send

Write "draft only — never send, publish or approve" in the skill AND the schedule.

Keep a human gate

The schedule fills an inbox. A named person still decides.

Campaign 5: Always On — the Weekly Growth Report

LAB 11

DAY 2 • ~35 MIN

THE STORY SO FAR

The team must keep working after class. Rachel wants the Growth Analyst's report in her inbox every Monday — as a draft she decides on.

YOU'LL BUILD

a scheduled task that drafts the weekly growth report every Monday at 8am, and one draft in Gmail

Surface: Claude (skill, scheduled task, Drive, Gmail)

Features: the Growth Analyst's campaign-report skill • do • verify • schedule • draft-only delivery

Campaign 5: Always On — the Weekly Growth Report

Claude (skill, scheduled task, Drive, Gmail) • full step-by-step with reasoning is in your Learner Guide.

- 1 Put the data on Drive** Upload campaign-results.csv to the Horizon Studio folder.
- 2 Bring the skill** Upload the campaign-report skill folder under Customize → Skills in Claude.
- 3 Do it once** Paste Prompt A. Refine the report until it is right.
- 4 Schedule it** Create a scheduled task for Mondays at 8am and paste scheduled-task-instructions.md into its Instructions.
- 5 Run it now** Run the task once and open the Gmail draft.
- 6 Pause it** After class, pause the scheduled task.

Campaign 5: Always On — the Weekly Growth Report

PROMPT A — Claude: do it once

Act as Horizon's Growth Analyst. Using Google Drive, read campaign-results.csv in the Horizon Studio folder. Use the campaign-report skill and data/weekly-report-brief.md. Leave the report as a Gmail DRAFT to me, subject "Horizon weekly growth report" and today's date. Send nothing.

NOTE Automate the preparation, not the decision. The schedule fills an inbox; a person still decides.

Campaign 5: Always On — the Weekly Growth Report

Check your work before moving on.

CHECK YOUR WORK

- ✓ The campaign-report skill is available in Claude.
- ✓ The report names the cheapest channel per chat (the newsletter).
- ✓ It recommends a budget split with a reason per channel.
- ✓ A scheduled task runs Mondays at 8am with draft-only instructions.
- ✓ A manual run left a draft — nothing was sent.
- ✓ The task is paused after class.

Responsible AI in the Team

Each risk has a control you built in this course.

RISK	CONTROL	WHERE
Wrong figure	fact-check skill; the Lead's review; facts sheet	Labs 3, 5, 8–10
Read as advice or a promise	fin-compliance; disclaimers; no returns	Labs 3, 8–10
An agent with too much access	Least-privilege connectors per agent	Lab 4
Unapproved post	approve.mjs (people only), hash, hook	Lab 8
Marketing without consent	The Analyst filters by consent	Lab 9
Undisclosed AI media	Description line + YouTube disclosure	Lab 10

TIP Keep keys in .env, client data out of prompts, and the approval records for audit.

Measure What Matters

Lab 12: the numbers decide next month — and which patterns earn their tokens.

Cost per chat

Spend ÷ chats booked shows which channel works: email first, YouTube last.

Defects before a person

A Lead that catches mistakes saves the reviewer's time; one agent alone caught none.

The playbook

Roster, instructions, skills, connectors, approval matrix, incidents — so the team runs without you.

Measure, Govern and Write the Team Playbook

LAB 12

DAY 2 • ~40 MIN

THE STORY SO FAR

Eight weeks in. Rachel asks which ways of running agents earned their tokens, and how the team runs safely after you leave.

YOU'LL BUILD

reports/agent-comparison.md and strategy/team-playbook.md

Surface: Claude Code (subagents)

Features: compare agent patterns • the roster: instructions, skills, connectors • approval matrix • responsible-AI controls

Measure, Govern and Write the Team Playbook

Claude Code (subagents) • full step-by-step with reasoning is in your Learner Guide.

- 1 Add the data** Copy the assets into data/.
- 2 Compare the patterns** Paste Prompt A.
- 3 Write the playbook** Paste Prompt B.
- 4 Review it as Rachel** Add one rule of your own and say why.

Measure, Govern and Write the Team Playbook

PROMPT A — Claude Code: compare

Use the growth-analyst subagent. From data/run-log.csv compare the patterns we used — one agent, subagents, agent teams and the scheduled task — on minutes, tokens, defects caught before a person and human edits. Which pattern pays off for which job? Save reports/agent-comparison.md.

Measure, Govern and Write the Team Playbook (cont.)

PROMPT B — Claude Code: the playbook

Write strategy/team-playbook.md with data/playbook-outline.md. Include the roster: for each agent its instructions file in .claude/agents/, its skills, connectors, model and what it may never do; the approval matrix by risk; the rules in data/responsible-ai-checklist.md; and what to do when a wrong post goes out. Then run fin-compliance on it.

NOTE Teams cost more tokens. The run log shows what you buy with them: defects caught before a person, and fewer edits.

Measure, Govern and Write the Team Playbook


Check your work before moving on.

CHECK YOUR WORK

- ✓ agent-comparison.md shows the one-agent baseline caught 0 defects.
- ✓ It says which pattern suits which job, with numbers.
- ✓ The roster lists all six agents with skills and connectors.
- ✓ The approval matrix says who approves CPF and tax figures.
- ✓ It has an incident procedure, and your own rule.


What Is Hermes Agent?

Nous Research's open-source (MIT) agent — the same team on any model. A demo, not assessed.



The Hermes Kanban Board

Every task is a row anyone can read; every hand-off is visible.



TIP The dispatcher runs inside hermes gateway start.

Optional Demo: the Team on Hermes Agent

LAB 13

DAY 2 • ~25 MIN

THE STORY SO FAR

Optional. Rachel's IT adviser asks what happens if the AI provider changes its prices. The trainer runs the same pipeline on open-source Hermes Agent.

YOU'LL BUILD

three Hermes profiles and a pipeline on the board, with the review task blocked for a person (demo)

Surface: Hermes Agent (profiles, Kanban) — trainer demo

Features: open-source Hermes Agent • one profile per role • a Kanban pipeline • a task blocked for a person • the same skills

Optional Demo: the Team on Hermes Agent

Hermes Agent (profiles, Kanban) — trainer demo · full step-by-step with reasoning is in your Learner Guide.

- 1 Install Hermes** Follow hermes-setup.md with an API key (a Claude subscription does not cover Hermes).
- 2 Create the bots** Create the profiles in hermes-roles.md and copy skills/ into each.
- 3 Start the board** hermes kanban init, then hermes gateway start in a second terminal.
- 4 Build the pipeline** Paste the prompt into hermes, or run kanban-pipeline.sh.
- 5 Be the gate** The Lead blocks for a person: unblock it with hermes kanban unblock <id>.

Optional Demo: the Team on Hermes Agent

PROMPT — Hermes: the planner

You are Horizon's Marketing Team Lead. Create a Kanban pipeline with dependencies:

1. strategist: the angle and facts for a LinkedIn post on CPF cash top-ups, from facts-2026.md.
 2. creator (parent 1): write it with sunny-voice.
 3. lead (parent 2): fact-check and fin-compliance; then block the task with the reason "Needs approval from a person".
- Show me the task ids and the board.

NOTE Optional, not assessed. The pattern carries over: roles, skills and a human gate work on an open-source agent too.

Optional Demo: the Team on Hermes Agent


Check your work before moving on.

CHECK YOUR WORK

- ✓ hermes profile list shows the three profiles.
- ✓ The board shows three tasks linked by parents.
- ✓ The review task is blocked for a person.
- ✓ Unblocking finished the pipeline.

How Each Pattern Is Triggered


Five ways your team runs. The difference is what sets each one off.



Who decides? You: the approval · the model: delegation and skills · the system, every time: hooks and the clock

Horizon: One Team, Five Campaigns

Two days, one regulated business — the Claude Edition.



Courseware & Assessment on the LMS



Download the courseware

Get the slides and learner guide from the LMS for the open-book assessment.



Take the assessment

Complete the Written Assessment and Practical Performance on the LMS.



Log in

Use your enrolled account to access the course materials and assessment.

A screenshot of the Tertiary Infotech Academy LMS login interface. It features a dark blue background with a white logo at the top. Below the logo, the text 'Tertiary Infotech Academy' and 'LMS cum TMS for WSQ Courses' is displayed. There is an 'Email' input field, a 'Send OTP' button, and a link to 'Login with password instead'. At the bottom, there are links for 'Privacy Policy', 'Acceptable Use Policy', and 'Feedback', along with the text 'Powered by Tertiary Infotech Academy Pte Ltd'.

LMS: <https://lms-tms.tertiaryinfotech.com/>

Certification & TRAQOM Survey



TRAQOM Survey (Mandatory)

Complete the certification & TRAQOM survey to receive your certificate.

<https://lms-tms.tertiaryinfotech.com/>



Digital Attendance

Remember to take AM, PM and Assessment digital attendance — required for your funding.

Practice Exam: Claude Certified Associate

exams.tertiaryinfotech.com/practice-exams/anthropic/anthropic-ccao-foundations

The screenshot displays the Tertiary Exams website interface for the Claude Certified Associate — Foundations (CCAO-F) practice exam. The page includes a navigation bar with links to Practice Exams, Vendors, About Us, FAQ, and Support, along with a 'Get started' button. The main content area features a breadcrumb trail (All exams / Anthropic / Anthropic / CCAO-F / Foundational) and a detailed description of the exam bundle. Key details include 60 questions, a 120-minute duration, and a 72% pass score. The page also offers a free practice teaser and a paid exam option for SGD \$20.00. A table lists the domains covered and their respective percentages.

Domains covered	
Prompting and Task Execution	14%
Output Evaluation and Validation	21%
Product and Model Selection	12%
Workflow Integration and Solution Design	16%

PRACTICE EXAM exams.tertiaryinfotech.com/practice-exams/anthropic/anthropic-ccao-foundations

CCAO-F · Foundations

For professionals who use Claude as a productivity tool.

What it covers

Prompting and task execution, output evaluation, model selection, workflow integration, responsible use.

Try it free

Start with the free practice teaser; exam mode is 60 questions in 120 minutes.

Assessment & Funding



Final Assessment

- Written Assessment (WA) — 1 hour, starts 4:00 PM
- Practical Performance (PP) — 1 hour
- Format: Open book — class ends 6:00 PM
- Allowed: slides, learner guide & approved materials only
- Appeal process available



Criteria for Funding

Minimum 75% attendance


Based on SSG digital attendance records.

Assessed as 'Competent'

Complete and pass both assessment components.

Assessment Flow

At the assessment stage, follow these five steps in order.



TIP Courseware and the assessment are on the LMS: <https://lms-tms.tertiaryinfotech.com/>

Digital Attendance



Attendance is mandatory for SSG funding.

Take AM, PM and Assessment digital attendance for every session.

Use the SSG / TMS QR code provided by your trainer.

**Minimum
75%
attendance
required**

Thank You!

Questions? We're here to help — during and after class.

enquiry@tertiaryinfotech.com | +65 6100 0613