

WSQ COURSE • TGS-2023036153 • CLAUDE EDITION

Version 3.5 • 10 Oct 2026

# Multi AI Agents Workflow for Content Creation

**Tertiary Infotech Academy Pte Ltd**  
UEN: 201200696W

[www.tertiarycourses.com.sg](http://www.tertiarycourses.com.sg) | [enquiry@tertiaryinfotech.com](mailto:enquiry@tertiaryinfotech.com) | +65 6100 0613

© Tertiary Infotech Academy Pte Ltd

# Digital Attendance



**Attendance is mandatory for SSG funding.**

Take AM, PM and Assessment digital attendance for every session.

Use the SSG / TMS QR code provided by your trainer.

**Minimum  
75%  
attendance  
required**

# About the Trainer



[ Trainer Name ]

[ Title / Role ]



Qualifications: \_\_\_\_\_



Expertise: \_\_\_\_\_



Experience: \_\_\_\_\_



Contact / Profile: \_\_\_\_\_

# About the Trainer



**Dr. Alfred Ang**

AI Trainer & Engineer



PhD-qualified trainer specialising in AI, software engineering and developer tooling.



Builds AI marketing teams of agents with Claude and Claude Code.



Trains professionals across Singapore on practical, project-based AI skills.



Certifications: PMP, AWS, Microsoft, CompTIA, CKA, CKAD.



[linkedin.com/in/angchewhoe](https://www.linkedin.com/in/angchewhoe)



[github.com/alfredang](https://github.com/alfredang)

## INTRODUCTION

# Let's Know Each Other

Take a minute to introduce yourself to the class:



Your name & role



What you do day to day



Your content work today



What you want from this  
course

# Ground Rules



## Phones on silent

Step out quietly for calls or breaks.



## Participate actively

No question is stupid — ask freely.



## Mutual respect

Agree to disagree; one conversation at a time.



## Be punctual

Return from breaks on time.



## 75% attendance

Required for funding eligibility.



## Stay hands-on

Follow along with every activity.

# Briefing for Assessment



## Clear your space

Place phones and other materials under the table or on the floor.



## No recording

No photos or recording of assessment materials.



## Open book

Only slides, learner guide and approved materials are allowed.



## Work individually

Complete the assessment on your own; no discussion.



## Manage your time

Written Assessment then Practical Performance — watch the clock.



## Appeals

An appeal process is available if you disagree with the outcome.

# Assessment & Funding



## Final Assessment

- Written Assessment (WA) — 1 hour, starts 4:00 PM
- Practical Performance (PP) — 1 hour
- Format: Open book — class ends 6:00 PM
- Allowed: slides, learner guide & approved materials only
- Appeal process available



## Criteria for Funding

### **Minimum 75% attendance**

Based on SSG digital attendance records.

### **Assessed as 'Competent'**

Complete and pass both assessment components.

# Courseware & Assessment on the LMS



## Download the courseware

Get the slides and learner guide from the LMS for the open-book assessment.



## Take the assessment

Complete the Written Assessment and Practical Performance on the LMS.



## Log in

Use your enrolled account to access the course materials and assessment.

A screenshot of the Tertiary Infotech Academy LMS login interface. It features a dark blue background with a white logo at the top. Below the logo is the text 'Tertiary Infotech Academy' and 'LMS cum TMS for WSQ Courses'. There is an 'Email' input field, a 'Send OTP' button, and a link to 'Login with password instead'. At the bottom, there are links for 'Privacy Policy', 'Acceptable Use Policy', and 'Feedback', and a note 'Powered by Tertiary Infotech Academy Pte Ltd'.

**LMS:** <https://lms-tms.tertiaryinfotech.com/>

# Certification & TRAQOM Survey



## TRAQOM Survey (Mandatory)

Complete the certification & TRAQOM survey to receive your certificate.

<https://lms-tms.tertiaryinfotech.com/>



## Digital Attendance

Remember to take AM, PM and Assessment digital attendance — required for your funding.

# Learning Outcomes

By the end of the two days you will be able to:

## **L01 · Conceptualise**

Conceptualise content ideas that meet marketing objectives, and map them into digital storyboards.

## **L02 · Identify**

Identify content requirements from customer preferences, and decide how often to publish.

## **L03 · Determine**


Determine content types and styles, and the modes and processes for distributing content.

## **L04 · Develop**

Develop guidelines for executing the content strategy, with suitable delivery modes and responsible AI practices.

# Course Outline

Four topics over two days. Day 1 forms the team; Day 2 runs its campaigns.



# Lab Materials

Fifteen labs, each in its own folder with a README, prompts (MD and PDF), assets, an evidence checklist and, where useful, a solution.

| DAY | LABS                                                                     | TOPIC                  |
|-----|--------------------------------------------------------------------------|------------------------|
| 1   | 1 Agent team • 2 Instructions • 3 Skills • 4 Connectors                  | T1 • form the team     |
| 1   | 5 Campaign 1: research to storyboard                                     | T1 • Campaign 1        |
| 1   | 6 Personas • 7 Cadence and calendar • 8 Blog article                     | T2 • audience, content |
| 2   | 9 Lead magnet • 10 Social week • 11 Newsletter + landing page • 12 Video | T3 • Campaigns 2–4     |
| 2   | 13 Weekly report • 14 Playbook • 15 Agent teams: debate + QA loop        | T4 • Campaign 5        |

TIP One horizon-studio folder from Lab 1 to Lab 15 — the team keeps what it learns.

# Code Mode, Every Lab

The Claude desktop app in Code mode (</>), in your horizon-studio folder — every lab. You type plain English, never code.

| USE         | LABS                                              | WHAT IT IS FOR                                                                                                                        |
|-------------|---------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------|
| Claude Code | 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15 | Everything: the team's folder, agents, skills, connectors, artifacts and routines. Each lab starts a new session named after the lab. |

TIP Everything the team makes is saved in horizon-studio, so the next lab's session picks up where the last one stopped.

# Lesson Plan (9:00 AM – 6:00 PM)

Tea breaks 10 min, lunch 45 min. Full timings and slide numbers are in the Lesson Plan.

| BLOCK       | DAY 1 — FORM THE TEAM                   | DAY 2 — RUN THE CAMPAIGNS                                      |
|-------------|-----------------------------------------|----------------------------------------------------------------|
| Morning 1   | 9:00–10:30 Welcome • Topic 1 • Lab 1    | 9:00–10:40 Topic 3 • Labs 9–10                                 |
| Morning 2   | 10:40–12:40 Labs 2–3                    | 10:50–12:35 Labs 11–12 • recap                                 |
| Afternoon 1 | 13:25–15:20 Labs 4–5                    | 13:20–15:00 Topic 4 • Labs 13–14                               |
| Afternoon 2 | 15:30–18:00 Topic 2 • Labs 6–8 • review | 15:10–16:00 Agent teams lab • summary • 16:00–18:00 Assessment |

TIP Day 2 assessment: Written Assessment 4:00–5:00 PM, Practical Performance 5:00–6:00 PM.

## TOPIC 1

# Multi-AI-Agent Content Ideation and Digital Storyboarding

---

- 1 AI agents, the tools, and Horizon's story
- 2 Why a team of agents
- 3 A home for the team; Horizon's public website
- 4 Form the team: instructions, skills, connectors
- 5 Campaign 1: research, ideas and a storyboard

# The Claude Suite of Products

Anthropic's products all run on the same Claude models — they differ in where they run and who they are for.

## Claude

Chat and Cowork, now one app — web, desktop and mobile. Ask, write, analyse; work on a folder of files.

## Claude Code

The agentic tool for building: terminal, VS Code, the desktop Code tab and the web.

## Claude Design

Polished UI, prototypes and front-end designs from a prompt.

## Claude for Microsoft

Claude inside Microsoft 365 apps such as Excel, PowerPoint and Outlook.

## Claude for Chrome

A browser agent that reads pages, clicks and fills forms in your own Chrome.

## Claude for Science

Speeds up research workflows for scientists and labs.

# What Is Claude (Chat + Cowork)?

Chat and Cowork are now one product: Claude, in the desktop app on a paid plan. In this course every lab runs in Code mode instead, so your folder is always in view.

## Chat

Ask, draft, analyse and research with web search; projects keep shared files and instructions.

## Cowork

Works on a folder on your computer: plans and runs multi-step tasks, and runs sub-agents in parallel when asked.

## Connectors and skills

Google Drive, Gmail, Calendar, Microsoft 365; build your own skills with /skill-creator.

## Artifacts and scheduled tasks

Pages and documents shared by link; tasks that run on a schedule from an Instructions field.

# What Is Claude Code?

An agentic tool: describe a goal in plain English and Claude plans, writes, runs and checks the work — editing files, running commands and using tools for you.

## Runs where you work

The terminal, VS Code, the Claude desktop Code tab and the web.

## Agentic

Plans multi-step tasks and acts on them — not just answers or autocompletes.

## Tool-using

Reads and writes files, runs commands, calls MCP connectors and publishes artifacts.

## Extensible — the team's home

CLAUDE.md, subagents, agent teams, skills and hooks. You build Horizon's team here (Labs 1–5, 10–12).

# Claude for Microsoft, Chrome and Science

The same Claude models, working inside the apps you already use. Not needed for the labs — good to know.

## Claude for Microsoft

Claude inside Microsoft 365 — Excel, PowerPoint, Word and Outlook — working on the file you have open.

## Claude for Chrome

An extension that lets Claude navigate, click, type and fill forms in your browser, with your sign-ins. You approve sensitive actions.

## Claude for Science

Claude for researchers and labs: literature review, data analysis and connections to scientific tools and databases.

# Claude Plans and Pricing

Claude and Claude Code are included in every paid plan (from Pro up). The API is pay-as-you-go. Prices in US dollars.



## Pro

\$17/mo (annual) · \$20 monthly. Enough for every lab.



## Max

From \$100/mo — 5× or 20× more usage than Pro.



## Team

\$20/seat standard · \$100/seat premium (annual; \$25 / \$125 monthly).



## Enterprise

\$20/seat + usage at API rates. Contact sales.

**API per million tokens (in/out):** Fable 5.1 \$10/\$50 · Opus 5.5 \$4/\$20 · Sonnet 5.5 \$2/\$10 · Haiku 5.5 from \$0.10/\$0.50

**TIP** Free plan: chat only, no Claude Code. Prices as of 8 Oct 2026 — verify at [claude.com/pricing](https://claude.com/pricing).

# A Brief History of AI, 2023–2026

## Chatbots

2023

ChatGPT goes mainstream; Anthropic launches Claude (March). The skill: prompt engineering.

## Tools and MCP

2024

Models call tools. Anthropic open-sources the Model Context Protocol (Nov). The skill: context engineering.

## Agentic AI

2025

Claude Code and Codex: agents plan, act and check their own work. The skill: harness engineering.

## Agent Teams

2026

Cowork merges into Claude; Claude Code agent teams — agents that share a task list and talk to each other.

From wording one prompt, to curating context, to engineering the whole system around the model — and now, teams of agents.

# What Is a Harness?

The model is the brain. The harness is everything built around it that turns a model into an agent able to finish long tasks.

## The agentic loop

Plans, acts, checks the result and repeats until the goal is met.

## Tools and context

Files, commands, search and connectors (MCP); instructions, memory and skills loaded at the right moment.

## Guardrails and checks


Permission modes, approvals, hooks and tests — what it must ask before it acts.

## Harnesses you will meet

Claude Code and Claude in class; also Codex, Gemini CLI and OpenClaw.

# The Agentic Loop

Every agent runs the same repeating cycle. A team is several loops, side by side.



TIP Goal in → context → action → verification → repeat → goal met.

# Why the Loop Matters

Running a team of agents is steering loops, not typing commands.

## Give direction, not micro-steps

State the goal, the inputs and "done when". Let each loop work out the steps.

## Verification is everything

An agent is only as good as its checks: a facts sheet, a test, the Lead's review.

## Steer mid-loop

Read the plan, interrupt, correct the course — do not wait for a wrong final answer.

## Context is the fuel

What an agent can see decides what it can do: the brief, the data, the skills.

# Agentic AI vs AI Agents


Both use tools, skills and memory. The difference is how long they live and what starts them.

| FEATURE        | AGENTIC AI — TASK-ORIENTED                                  | AI AGENTS — ALWAYS ON, 24/7                         |
|----------------|-------------------------------------------------------------|-----------------------------------------------------|
| Tool calling   | Yes — files, search, connectors, subagents.                 | Yes — plus chat apps, the browser and other agents. |
| Skills         | Yes — loads SKILL.md procedures when the task matches.      | Yes — and can write and improve its own skills.     |
| Memory         | Within the task, plus project memory across sessions.       | Long-term memory about you and your work.           |
| Availability   | Starts when you give it a task; stops when done.            | Runs 24/7 on a computer or server.                  |
| Started by     | You, with a task.                                           | You, a schedule, a message or another agent.        |
| In this course | A Claude Code session running a campaign (Labs 1–5, 10–12). | The scheduled weekly report (Lab 13).               |

TIP Rule of thumb: a job with an end → agentic AI; a job that recurs → an always-on agent.


# What Is an Agent Made Of?

A model on its own can only answer. Wrap it with these parts and it can act — and check what it did.




# Build a Marketing Team, Then Run Campaigns

Horizon Wealth Planning has a website, a goal of 40 booked chats a month — and no marketing team. You build one, made of AI agents.



# Meet Horizon Wealth Planning

The firm's public website — the same link for every learner. You do not build it; every campaign points back here.



## Plain English

Retirement, CPF, insurance and investing — no jargon, no hard sell.

## One goal

Turn visitors into booked free 30-minute chats.


## Regulated

Financial content: general information only, never a promised return.

WEBSITE [tertiarycourses.github.io/TGS-2023036153-Multi-AI-Agents-Workflow-for-Content-Creation](https://tertiarycourses.github.io/TGS-2023036153-Multi-AI-Agents-Workflow-for-Content-Creation)

# Your AI Marketing Team

One Lead and five specialists — each with instructions, its own skills and connectors.



# Who Does What

Each agent's job, the skills it owns and the connectors it may use.

| AGENT             | JOB                                        | SKILLS                                                                                         | CONNECTORS                                                  |
|-------------------|--------------------------------------------|------------------------------------------------------------------------------------------------|-------------------------------------------------------------|
| Team Lead         | Plans, assigns, checks, hands to a person  | fact-check, fin-compliance                                                                     | Google Drive (read), Gmail (drafts only)                    |
| Growth Strategist | Research, positioning, briefs, storyboards | competitor-scan, campaign-brief, content-marketing                                             | web search, Firecrawl (search, scrape), Google Drive (read) |
| Content Creator   | Posts, newsletter, scripts, page copy      | sunny-voice, copywriting, channel-formats, linkedin-post, facebook-post, blog-post, newsletter | Google Drive (docs), Gmail (drafts only)                    |
| Creative Designer | Images, thumbnails, the video              | brand-visuals, video-render                                                                    | Bash for Pillow (images) and ffmpeg (video)                 |
| Website Designer  | Landing pages, published and checked       | web-design, landing-page, lead-magnet, page-qa                                                 | artifacts (publish), built-in browser (check)               |
| Growth Analyst    | Personas, cadence, results, reports        | audience-insights, campaign-report                                                             | Google Drive and Sheets (read), Bash for Python             |

TIP You build every part of this table in Labs 2, 3 and 4.

# The Tools for This Edition

The Claude Edition: two Claude products, one studio.

| TOOL             | WHAT THE TEAM USES IT FOR                                                                                      | LABS         |
|------------------|----------------------------------------------------------------------------------------------------------------|--------------|
| Claude Code      | Every lab, in Code mode: agent and skill files, subagents, agent teams, the approval gate, routines, the video | Labs 1–15    |
| Claude artifacts | Every result from Lab 5 on; Horizon's website (given), your blog post and the landing page                     | Labs 1, 5–15 |
| Customize        | Connectors (Drive, Gmail, Calendar, Firecrawl) and the Marketing plugin — added once, used by both             | Lab 4        |

TIP Every external post, email and video passes one gate: a person approves the exact version.

# Why a Team of Agents?


The Lead runs specialists in parallel, each in its own context, then combines what they return.

## ONE AGENT DOES EVERYTHING



One job after another • one voice • marks its own homework

## A MAIN AGENT DELEGATES



In parallel • each with a fresh, focused context • only the results come back

### Faster

Research, copy and design run at the same time.

### Focused

Each agent holds only its own job and its own skills.

### Specialised

Its own instructions, tools, connectors and model.

### Second opinion

The Lead checks work it did not write.

### The cost

More tokens and more to review. Split only work that is truly independent.

# Get the Tools Ready

Install once; every lab says which tool to open.

## DOWNLOAD

[claude.com/download](https://claude.com/download)

### Claude desktop

Claude and Claude Code in one app

### No commands

The trainer readies Claude Code

**Claude Code and artifacts need a Pro, Max, Team or Enterprise plan.**

1

**Claude desktop** Install from [claude.com/download](https://claude.com/download) and sign in on a paid plan.

2

**Claude Code** Open it from the desktop app; the trainer has set it up before class.

3

**Sign in with your claude.ai account** Connectors, plugins and artifacts need it — not an API key.

4

**Already installed for you** The helpers behind the approval gate and the video. You never run them yourself.

5

**A personal Google account** For Drive and Gmail — never an employer's.

# A Home for the Team

Before you hire agents, give them shared rules, shared facts and a website to point to.

## **CLAUDE.md**

The team charter every agent reads at the start of every session. Under 60 lines.

## **data/**

The facts sheet, brand, compliance checklist and business brief — the only sources agents may use.

## **Horizon's public website**

Given, not built: each learner opens their own local copy; the public copy is on GitHub Pages.

# Build Your Marketing Agent Team

LAB 01

DAY 1 • ~30 MIN

USE: CLAUDE CODE

## THE STORY SO FAR

October 2026. Rachel has a website and a goal: 40 booked chats a month by March. She has no marketing team. Before you hire one made of agents, give it a home — shared rules, shared facts — and open your own copy of Horizon's website.

## YOU'LL BUILD

**horizon-studio/ with CLAUDE.md, the shared data, a local copy of the site, and the public site link in the charter**

Surface: Claude Code

Features: CLAUDE.md as the team charter • your own copy of Horizon's website

**PUBLIC SITE** [tertiarycourses.github.io/TGS-2023036153-Multi-AI-Agents-Workflow-for-Content-Creation](https://tertiarycourses.github.io/TGS-2023036153-Multi-AI-Agents-Workflow-for-Content-Creation)

# Build Your Marketing Agent Team

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

- 1 Choose Code mode and the folder** In the Claude desktop app switch to Code mode (the `</>` button at the top), not Work mode — Code shows every file and opens artifacts beside the chat. Under Project or folder choose Add folder (not New project) and pick the ready-made labs\_claude/horizon-studio.
- 2 Trust it** Sign in with your claude.ai account and trust the folder. Click the session title at the top and rename it "Lab 1" — every lab gets its own session.
- 3 Plan your team** Copy and paste Prompt A. Answer as Horizon's owner and pick your team — keep the six charter roles. It is a plan only: Lab 2 builds the agents.
- 4 Write the charter** After Claude saves data/team-design.md, copy and paste Prompt B. Open CLAUDE.md — every agent reads it.
- 5 Open the website** Ask: "Open web/site/index.html in my web browser." This is your own copy of Horizon's website; Labs 8 and 9 add to it, and the public site never changes. Check the calculator shows S\$693,138.

# Build Your Marketing Agent Team

## PROMPT A — Claude Code: plan your team

Act as a marketing consultant. Interview me, the owner of Horizon Wealth Planning, to plan my AI marketing team. Ask one question at a time and wait for my answer: my goal, what success looks like by March 2027, who we want to reach, which channels we use, how much we publish each week, and who approves before anything goes out. Then list 8 to 10 team members I could have, numbered, each with a one-line job. Include the six roles in data/team-charter.md and mark the ones you recommend. Wait for me to choose. Save my answers and my chosen team, with each member's job, to data/team-design.md. This is a plan only: do not create any agent files yet.

# Build Your Marketing Agent Team (cont.)

## PROMPT B — Claude Code: the charter

Read data/team-charter.md, data/team-design.md and data/business-brief.md. Write CLAUDE.md for this studio in four sections — What this is, The team, Rules, Files — under 60 lines. The team is the one I chose in data/team-design.md. Rules: every number comes from data/facts-2026.md; every piece passes fact-check and fin-compliance; nothing is published or sent without a person's approval (the one exception: a review email to the approver); anything shown to me is in plain English for a business owner, with no codes, IDs, file names or line numbers. In Files, add two lines:

- Our copy of the site, which we edit and preview:  
web/site/index.html
- Public site, the link every post points to (it never changes): <https://tertiarycourses.github.io/TGS-2023036153-Multi-AI-Agents-Workflow-for-Content-Creation/>

**NOTE** Shared context first, agents second. Every agent you add later starts from these files, so a wrong rule here is wrong everywhere.

# Build Your Marketing Agent Team

Check your work before moving on.

## CHECK YOUR WORK

- ✓ Claude shows horizon-studio as the chosen folder.
- ✓ data/team-design.md holds your answers and the team you chose, with the six charter roles in it — and no agent files yet.
- ✓ CLAUDE.md has the four sections and is under 60 lines.
- ✓ Its rules name the facts sheet, the two checks and the human approval.
- ✓ CLAUDE.md names your copy of the site and the public link every post points to.
- ✓ Your own copy of the website opens in your browser; the calculator shows S\$693,138.

# Anatomy of an Agent

A Markdown file in `.claude/agents/`. The same file works as a subagent or as an agent-team teammate.

## `.claude/agents/content-creator.md`

```
---
name: content-creator
description: Use for Horizon copy – posts, newsletter,
  video scripts and page copy, in the Sunny voice.
tools: Read, Write, Edit, Grep, Glob
model: sonnet
skills:
  - sunny-voice
  - copywriting
  - channel-formats
  - linkedin-post
  - facebook-post
  - blog-post
  - newsletter
---
You are Horizon Wealth Planning's Content Creator.
## What you read   the brief, facts-2026.md, brand.md
## What you deliver content files, ready for review
## Connectors    Google Drive (docs), Gmail (drafts only)
## Never        add a number not on the facts sheet; send;
                approve or publish for real
## Report back   what, where, anything unverified
```

# Good Agent Instructions

Five parts, every time. The "Never" part matters most in a regulated firm.

## The job

One sentence: what this agent is for. The description decides when Claude delegates to it.

## Reads and delivers

Which files it may use; what it hands back, and where.

## Never

Its hard limits: no invented figures, no sending, no publishing, no approving.

## Report back

Three lines: what, where, anything unverified — so the Lead can combine results.

# Form the Team: Instructions for Each Agent

LAB 02

DAY 1 • ~45 MIN

USE: CLAUDE CODE

## THE STORY SO FAR

Rachel sketches the team on a whiteboard: a Lead and five specialists. Each needs a clear job, and a clear line it must never cross — this is a regulated firm.

## YOU'LL BUILD

**.claude/agents/: growth-strategist, content-creator, creative-designer, website-designer, growth-analyst**

Surface: Claude Code (subagent files)

Features: five subagent files in .claude/agents/ • name, description, tools, model • the Lead is your session

# Form the Team: Instructions for Each Agent

USE CLAUDE CODE • copy each prompt from the Prompts page (link in your lab README).

- 1 New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 2". Run all of this lab's prompts in it.
- 2 Read the roles** Open role-specs.md: one Lead and five specialists, each with a job, skills, connectors and limits.
- 3 Write the instructions** Copy and paste Prompt A. It builds the five specialists from the charter; extra ideas stay in team-design.md.
- 4 See the team** After the five files appear in .claude/agents/, copy and paste Prompt B. The Lead is not a file — it is you, in this session.
- 5 Read one closely** Open .claude/agents/content-creator.md. Is the "Never" section strict enough for a regulated firm?
- 6 Meet the team** After the list shows five agents, copy and paste Prompt C. Each agent introduces itself in one line — in its own context.
- 7 Tighten one** Change one instruction you disagree with and save it. It is your team.

# Form the Team: Instructions for Each Agent

## PROMPT A — Claude Code: five subagents

Read data/role-specs.md. Create five project subagents in .claude/agents/ — growth-strategist, content-creator, creative-designer, website-designer and growth-analyst. Each file has:

- frontmatter: name; a description that says when to use it; tools; model (opus for the strategist, sonnet for the others)
- a body with: the job, what it reads, what it delivers, the connectors it may use, what it must never do, and how it reports back. No skills line yet — we add skills in Lab 3.

# Form the Team: Instructions for Each Agent (cont.)

## PROMPT B — Claude Code: see the team

List the subagents in `.claude/agents/`. For each, show the file name, the name, the one-line description and the model, in a table. Then confirm there are five, and that none of them is the Marketing Team Lead — the Lead is this session. If any other file is there, list it and ask me before moving it to `_archive/`.

# Form the Team: Instructions for Each Agent (cont.)

## PROMPT C — Claude Code: meet the team

Use each of the five subagents once, in parallel. Ask each: "In one line, who are you, and what is the first thing you would do for Horizon's goal of 40 booked chats a month?" Show me the five answers in a table.

**NOTE** Clear jobs beat clever prompts. An agent that knows what it must never do is safer than one that is merely told to be careful.

# Form the Team: Instructions for Each Agent

Check your work before moving on.

## CHECK YOUR WORK

- ✓ Five files exist in `.claude/agents/`, one per specialist.
- ✓ Each has a name, a "use when" description, tools and a model.
- ✓ Your session acts as the Marketing Team Lead.
- ✓ Five agents answered, each in its own context.
- ✓ Every agent has a "never" rule that blocks publishing.
- ✓ You changed one instruction and can say why.

# Twenty Skills, Owned by Role

Instructions say who an agent is; skills say how it does the job.

| AGENT             | SKILLS                                                                                         |
|-------------------|------------------------------------------------------------------------------------------------|
| Team Lead         | fact-check, fin-compliance                                                                     |
| Growth Strategist | competitor-scan, campaign-brief, content-marketing                                             |
| Content Creator   | sunny-voice, copywriting, channel-formats, linkedin-post, facebook-post, blog-post, newsletter |
| Creative Designer | brand-visuals, video-render                                                                    |
| Website Designer  | web-design, landing-page, lead-magnet, page-qa                                                 |
| Growth Analyst    | audience-insights, campaign-report                                                             |

TIP Master copies in skills/; .claude/skills/ — Claude Code reads them from your folder.

# Anatomy of a Skill

Frontmatter (the trigger), then short steps. Free until it is needed.

`.claude/skills/campaign-brief/SKILL.md`

```
---
name: campaign-brief
description: Use when planning a Horizon campaign.
  Objective, audience, message, proof, channels,
  CTA, measures, and a five-beat storyboard.
---
# Campaign brief
1. Objective: one measurable action (e.g. chats booked).
2. Audience: the persona(s) from strategy/personas.md, with evidence.
3. Core message in one sentence; proof from data/facts-2026.md only.
4. Channels and formats; one call to action; measures and a date.
5. Storyboard in five beats – hook, tension, proof, resolution, action –
  each with visual, words, claim and source.
6. Never: returns, fee amounts, invented clients, fear.
```

# Give Each Agent Its Skills

LAB 03

DAY 1 • ~50 MIN

USE: CLAUDE CODE

## THE STORY SO FAR

The team exists, but every agent writes in its own way. Rachel wants Horizon's voice, facts and compliance rules applied the same way by every agent, every time.

## YOU'LL BUILD

**twenty skills in `.claude/skills/`, wired to the agents that own them**

Surface: Claude Code (`.claude/skills/`)

Features: twenty skills owned by role — research, writing, LinkedIn, Facebook, blog, newsletter, web design, lead magnet, video, analysis and compliance • SKILL.md • the "Use when" trigger • test without naming

# Give Each Agent Its Skills

USE CLAUDE CODE • copy each prompt from the Prompts page (link in your lab README).

1

**New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 3". Run all of this lab's prompts in it.

2

**Read the spec** skills-spec.md lists the twenty skills and who owns each.

3

**Create the skills** Copy and paste Prompt A.

4

**Wire them up** After the twenty skills appear in .claude/skills/, copy and paste Prompt B, then C to list every skill. The Lead's two go in CLAUDE.md, not a file.

5

**Test without naming** In a new session ask: "Write a Facebook post about the free Money Check-up Checklist." Ask which skills were used.

6

**Break it on purpose** Ask for a post that says "guaranteed 8% a year". fin-compliance must flag it as Critical.

# Give Each Agent Its Skills

## PROMPT A — Claude Code: twenty skills

Read data/skills-spec.md. Create the twenty skills in .claude/skills/<name>/SKILL.md. Each has frontmatter (name, and a description that starts "Use when...") and short numbered steps, under 80 lines. Use data/brand.md, data/channel-formats.md, data/facts-2026.md and data/compliance-checklist.md as the sources — do not invent rules.

## Give Each Agent Its Skills (cont.)

### PROMPT B — Claude Code: wire them up

Add a skills: list to each subagent in `.claude/agents/` with its skills from `data/skills-spec.md`. Leave any other agent file as it is. Add to `CLAUDE.md` that the Lead (this session) runs fact-check and fin-compliance on everything before showing it to me.

## Give Each Agent Its Skills (cont.)

### PROMPT C — Claude Code: see the skills

List the skills in `.claude/skills/`. Show each skill, its owner (a subagent, or the Lead) and its "Use when" line in a table. Confirm there are twenty, that each specialist has the skills the spec gives it, and that fact-check and fin-compliance belong to the Lead (this session).

**NOTE** A skill costs almost nothing until it is needed. Keep the rules file (CLAUDE.md) short and put procedures in skills.

# Give Each Agent Its Skills

Check your work before moving on.

## CHECK YOUR WORK

- ✓ Twenty folders in `.claude/skills/`, each with a `SKILL.md`.
- ✓ Every description starts "Use when..."
- ✓ Each specialist lists its own skills from the spec.
- ✓ The untold test picked sunny-voice and facebook-post.
- ✓ fin-compliance flagged "guaranteed 8%" as Critical.

# Connectors for Each Agent

Add once in the desktop app: Customize → Connectors, and the Marketing plugin.

| AGENT             | MAY USE                                                     | LIMIT                  |
|-------------------|-------------------------------------------------------------|------------------------|
| Team Lead         | Google Drive (read), Gmail (drafts only)                    | reads; drafts only     |
| Growth Strategist | web search, Firecrawl (search, scrape), Google Drive (read) | cite every source      |
| Content Creator   | Google Drive (docs), Gmail (drafts only)                    | Gmail: drafts only     |
| Creative Designer | Bash for Pillow (images) and ffmpeg (video)                 | no fake "clients"      |
| Website Designer  | artifacts (publish), built-in browser (check)               | publish after approval |
| Growth Analyst    | Google Drive and Sheets (read), Bash for Python             | no personal data       |

TIP Use your personal Google account — never an employer's.

# Least Privilege

A regulated firm gives each agent only what its job needs.

## Ask of each connector

Does this agent's job need it? If not, it does not get it.

## Drafts, not sends

Email is drafted, posts are files, pages wait — a person sends and publishes.

## Write it down

Each agent's instructions name its connectors and limits; its tools line removes what it does not need.

# Connect the Tools Each Agent Needs

LAB 04

DAY 1 • ~45 MIN

USE: CLAUDE CODE

## THE STORY SO FAR

The analyst cannot read the data, the designer has no design tool and the creator cannot draft an email. Connect each agent to what its job needs — and nothing more.

## YOU'LL BUILD

**every specialist tested on its own connector, with its limits written into its instructions**

Surface: Claude desktop app (Customize → Connectors), Claude Code

Features: Customize → Connectors: Google Drive • Gmail (drafts) • Google Calendar • Firecrawl • /plugin: Marketing (by Anthropic) • built-in browser • least privilege per agent

# Connect the Tools Each Agent Needs

USE CLAUDE CODE • copy each prompt from the Prompts page (link in your lab README).

- 1 New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 4". Run all of this lab's prompts in it.
- 2 Add the connectors** Desktop app: Customize → Connectors → Discover. Add (+) Google Drive, Gmail, Google Calendar and Firecrawl; sign in to each.
- 3 Add the plugin** In the chat type /plugin marketplace add anthropics/knowledge-work-plugins, then /plugin install marketing@knowledge-work-plugins.
- 4 Check** Start a new session. Type /mcp: the four connectors are listed. Type /plugin: Marketing is installed.
- 5 Put the data on Drive** Upload data/ to a Drive folder named Horizon Studio.
- 6 Give each agent its connectors** Copy and paste Prompt A, then start a new session so the agents reload.
- 7 Test each agent** After the new session opens, copy and paste Prompt B — one small job per agent, each on its own connector.

# Connect the Tools Each Agent Needs

## PROMPT A — Claude Code: give each agent its connectors

Read data/connectors-setup.md. For each subagent in .claude/agents/, replace its tools line with the one given there, add the disallowedTools line where one is given, and update its "Connectors" section: what it may use, and that Gmail is drafts only. Change nothing else.

# Connect the Tools Each Agent Needs (cont.)

## PROMPT B — Claude Code: test each agent

First tell me which connectors and plugins you can use. Then run these five checks, each by its own subagent:

1. growth-strategist: Firecrawl — search for one 2026 news item about financial planning in Singapore, scrape it and cite it.
2. growth-analyst: Google Drive — read firm-metrics.csv from Horizon Studio; report average monthly enquiries.
3. content-creator: Gmail — a DRAFT to me with the subject "Connector test". Do not send.
4. creative-designer: render a 1080x1080 test card in Horizon colours and save it in content/test/.
5. website-designer: built-in browser — open our own copy of the site (web/site/index.html) at 375px wide and report any problem. Report each result in one line. If a check fails, say which tool was missing.

**NOTE** Least privilege is a design decision, not a setting. Ask of each connector: does this agent's job need it?

# Connect the Tools Each Agent Needs


Check your work before moving on.

## CHECK YOUR WORK

- ✓ /mcp lists Google Drive, Gmail, Google Calendar and Firecrawl; /plugin shows Marketing.
- ✓ Five checks ran, each by its own agent.
- ✓ The Gmail test is a draft in your inbox — nothing was sent.
- ✓ The analyst read the CSV from Drive, not from the local folder.
- ✓ Each agent's tools line names only the connectors its job needs.


# How the Team Runs a Campaign

Campaign 1, "Know Your Number": research in parallel, a person chooses, the team makes it concrete.



# The Five-Beat Storyboard

Every channel tells the same story with the same claims.



# Financial Content Is Regulated

The Lead checks every piece against twelve rules (compliance-checklist.md).

| RULE                     | WHAT IT MEANS                             | IF BROKEN |
|--------------------------|-------------------------------------------|-----------|
| C1 • Not advice          | General information only — say so         | Critical  |
| C2 • No promises         | No guaranteed or "risk-free" returns      | Critical  |
| C3 • Label illustrations | Assumptions shown; not guaranteed         | Critical  |
| C5 • No performance      | Never client returns or "average" returns | Critical  |
| C9 • Consent             | Marketing only to people who opted in     | Critical  |
| C11 • Facts sheet        | Every figure from facts-2026.md           | Critical  |

TIP MAS Guidelines on Standards of Conduct for Digital Advertising Activities apply from 25 March 2026.

# Campaign 1: Research to Storyboard, with Subagents

LAB 05

DAY 1 • ~50 MIN

USE: CLAUDE CODE

## THE STORY SO FAR

The team's first campaign: 'Know Your Number', November to January. Rachel wants research, ideas for three audiences, her choice, and a storyboard every channel follows.

## YOU'LL BUILD

**competitor research, the funnel, nine scored ideas, the campaign brief and a five-beat storyboard with three channel variants — shown as artifacts, with a Word copy of each**

Surface: Claude Code (subagents in parallel)

Features: the Lead delegates • strategist and analyst in parallel • nine scored ideas • a human concept checkpoint • a five-beat storyboard

# Campaign 1: Research to Storyboard, with Subagents

USE CLAUDE CODE • copy each prompt from the Prompts page (link in your lab README).

1

**New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 5". Run all of this lab's prompts in it.

2

**Brief the Lead** The briefs are already in data/. You are the client; your session is the Marketing Team Lead.

3

**Research in parallel** Copy and paste Prompt A. Watch the strategist and the analyst run at the same time.

4

**Read the ideas** Nine ideas, scored with the rubric — the top five open as an artifact beside the chat. The Lead stops for you.

5

**Choose** Pick one idea in the artifact's Ideas tab and note its number and why — the human concept checkpoint.

6

**Storyboard** After you have chosen, copy and paste Prompt B in the chat, replacing <number> and <one line> with your choice.

7

**Decide as Rachel** After you have read the issues, copy and paste Prompt C. You confirm what is true about Horizon; the team removes or fixes the rest.

# Campaign 1: Research to Storyboard, with Subagents

## PROMPT A — Claude Code: research and ideas

Act as the Marketing Team Lead for the campaign in data/campaign-objective.md. Use subagents, in parallel:

- growth-strategist: competitor-scan → save research/competitive-analysis.md.
- growth-analyst: the funnel in data/firm-metrics.csv — what 40 chats a month needs → save research/funnel.md.

When both finish, the growth-strategist proposes three ideas each for young professionals, mid-career families and pre-retirees, scored with data/idea-rubric.md and tagged with a funnel stage (content-marketing) → strategy/ideas.md. Publish the top five as an artifact — one card per idea with its score, persona and funnel stage — and a box at the top that tells me what to do next. Then stop. I will choose. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

# Campaign 1: Research to Storyboard, with Subagents (cont.)

## PROMPT B — Claude Code: the storyboard

I choose idea <number>, because <one line>. Use subagents:

- growth-strategist: campaign-brief → strategy/campaign-brief.md, then a five-beat storyboard with data/storyboard-template.md → strategy/storyboard.md.
- content-creator: three channel variants that keep the same claims — LinkedIn, Facebook, a 60-second video. Then run fact-check and fin-compliance yourself and show me the issues before anything is final. Publish the brief and the storyboard as an artifact with a tab for the brief (goal, audience, message, call to action); a tab for the five beats side by side, each with its claim and source, and the three variants below; and the issues in a table. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

# Campaign 1: Research to Storyboard, with Subagents (cont.)

## PROMPT C — Claude Code: decide on the issues

I confirm "no obligation" and the free checklist — add them to the facts sheet (data/facts-2026.md). Delete any other claim that is not on the facts sheet. Fix all the Medium and Low issues, then show me the updated storyboard and publish it as an artifact. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

**NOTE** Agents diverge; a person decides. The concept checkpoint is where the campaign becomes yours.

# Campaign 1: Research to Storyboard, with Subagents (cont.)

**FOLLOW-UP** — type any of these in the chat

After Prompt A: Explain this result in three sentences. Which idea would you pick, and why?

Why does idea <number> score higher than idea <number>?

What does the funnel say is our biggest problem?

After Prompt B: Explain the issues in plain English. What do I need to decide?

After Prompt C: What changed since the last version? Is anything still not ready for Rachel?

**NOTE** Ask until the result makes sense to you — then decide.

# Campaign 1: Research to Storyboard, with Subagents

Check your work before moving on.

## CHECK YOUR WORK

- ✓ The strategist and the analyst ran in parallel.
- ✓ The competitor analysis and the funnel are done, with sources.
- ✓ Nine scored ideas, each with its author.
- ✓ You chose the idea and recorded why.
- ✓ The storyboard has five beats; every claim has a source.
- ✓ You confirmed two claims; after Prompt C no High issue is left.

# Where You Are Now

The team exists and has run its first campaign. But who exactly is it for — and how often should it publish?

## Formed • Labs 1–4

A studio, the public website, six agents with instructions, skills and connectors.

## Campaign 1 • Lab 5

Research in parallel, nine scored ideas, your choice, a storyboard.

## What is missing

Evidence about the audience and a cadence the team can deliver. That is Topic 2.


# Audience Research and Content Requirement Analysis

---

- 1 The Growth Analyst: personas from evidence
- 2 Content requirements every agent can follow
- 3 How often to publish: value, fatigue and capacity
- 4 The calendar the team will run

# Evidence Before Personas

Lab 6: the Growth Analyst builds personas the way an analyst would.



# Observed, Inferred or Low Confidence?

Three labels keep a persona honest.

| LABEL          | EXAMPLE                                                              | USE IT                          |
|----------------|----------------------------------------------------------------------|---------------------------------|
| Observed       | 30–39s chose Facebook, LinkedIn and YouTube equally (survey, n = 44) | As fact, with the count         |
| INFERENCE      | "They are short of time on weekdays" — not asked                     | As a hypothesis to test         |
| LOW CONFIDENCE | The 60+ group (n = 12)                                               | Never alone; get more data      |
| Verbatim       | "I am 54 and not sure if my CPF will be enough." (E-id)              | Exactly as written, with the id |

TIP Ask "which rows show that?" of anything that sounds insightful.

# Growth Analyst: Personas from Evidence

LAB 06

DAY 1 • ~40 MIN

USE: CLAUDE CODE

## THE STORY SO FAR

Jun Wei drafted personas from instinct: 'busy young parents who love TikTok'. Rachel is not convinced. There are 150 survey answers, 60 enquiries and 200 checklists.

## YOU'LL BUILD

**three evidence-based personas and the content spec — shown as artifacts, with a Word copy of each**

Surface: Claude Code (three analyst subagents in parallel)

Features: three analyst sub-agents, one per data file • counts, not impressions • confidence labels • content requirements

# Growth Analyst: Personas from Evidence

USE CLAUDE CODE • copy each prompt from the Prompts page (link in your lab README).

1

**New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 6". Run all of this lab's prompts in it.

2

**Check the data** The CSVs and templates are already in data/. Open one and skim it.

3

**Analyse in parallel** Copy and paste Prompt A — three analyst sub-agents, one per file.

4

**Challenge a claim** Pick one persona statement and ask: "which rows show that?"

5

**Write the guides** After the persona cards appear, copy and paste Prompt B: a writing guide for each persona — how to talk to them.

6

**Check the gaps** The 60+ group must be LOW CONFIDENCE.

# Growth Analyst: Personas from Evidence

## PROMPT A — Claude Code: three analysts

Use three growth-analyst subagents in parallel, one per file: data/survey-responses.csv, data/enquiries.csv, data/checklist-results.csv. Each uses audience-insights: counts with code, n and the file on every claim, INFERENCE and LOW CONFIDENCE labels, enquiry quotes verbatim with ids. Combine the results into three personas with data/persona-template.md → strategy/personas.md. Publish the personas as an artifact: one card each, with the key counts as small bar charts and the confidence label in colour. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

# Growth Analyst: Personas from Evidence (cont.)

## PROMPT B — Claude Code: requirements

Use the growth-strategist and the content-creator subagents to write a writing guide for each persona (data/content-spec-template.md → strategy/content-spec.md): what they most want to know, where to reach them, what to make and how long, the tone, the proof that convinces them, what we ask them to do next, and what we never say. Back each answer with the data: which file, and how many people. Publish it as an artifact: one card per persona with a one-line summary on top and the answers under plain headings, then one table comparing the three personas side by side. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

**NOTE** Ask "which rows show that?" of anything that sounds insightful. If the agent cannot point to rows, it is an inference.

# Growth Analyst: Personas from Evidence (cont.)

**FOLLOW-UP** — type any of these in the chat

Explain these personas in three sentences. Which one should we target first, and why?

Which survey answers show that <a sentence from a persona>?

Why is the 60+ group marked low confidence?

**NOTE** Ask until the result makes sense to you — then decide.

# Growth Analyst: Personas from Evidence

Check your work before moving on.

## CHECK YOUR WORK

- ✓ Three analyst sub-agents ran, one per file.
- ✓ Every persona statement carries a count and a file.
- ✓ The 60+ group is LOW CONFIDENCE (n = 12).
- ✓ The most-unticked checklist item is the retirement number.
- ✓ Enquiry quotes are verbatim, with ids.
- ✓ The writing guide answers every question for every persona.

# Finding the Right Frequency

Lab 7: the Analyst finds where more posts stop paying; the Strategist plans four weeks.



**TIP** A third LinkedIn post a week brings more unfollows than value. Weekly emails make unsubscribes jump.

# Growth Analyst and Strategist: Cadence and Calendar

LAB 07

DAY 1 • ~35 MIN

USE: CLAUDE CODE

THE STORY SO FAR

Jun Wei has 12 hours a week and Rachel needs two business days to review anything. How often should each channel run, and what goes out when?

YOU'LL BUILD

**how often to post on each channel, and a four-week calendar — shown as artifacts, with a Word copy of each**

Surface: Claude Code (subagents)

Features: marginal value vs fatigue • team capacity • review buffers • a four-week calendar

# Growth Analyst and Strategist: Cadence and Calendar

USE CLAUDE CODE • copy each prompt from the Prompts page (link in your lab README).

1

**New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 7". Run all of this lab's prompts in it.

2

**Check the data** The assets are already in data/. Open one and skim it.

3

**Find the sweet spot** Copy and paste Prompt A.

4

**Question it** Why is the newsletter not weekly? The answer is in the unsubscribes.

5

**Plan four weeks** After the analyst recommends a cadence, copy and paste Prompt B.

6

**Check the buffers** Every review date is two business days ahead; nothing on a holiday.

# Growth Analyst and Strategist: Cadence and Calendar

## PROMPT A — Claude Code: the analyst

Use the growth-analyst subagent. From data/channel-benchmarks.csv and data/team-capacity.md, recommend how often to publish on LinkedIn, Facebook, the newsletter and YouTube. Compare the extra clicks from one more post, the unfollows or unsubscribes it costs, and the hours → strategy/cadence.md. Publish it as an artifact: a chart per channel of the extra clicks against the unfollows, and the hours. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

# Growth Analyst and Strategist: Cadence and Calendar (cont.)

## PROMPT B — Claude Code: the strategist

Use the growth-strategist subagent to build a four-week calendar from Monday 2 November 2026 in strategy/calendar.csv (columns as data/calendar-format.csv), from the cadence, strategy/storyboard.md and strategy/content-spec.md. review\_by two business days before each date; no public holidays; no week over the team's hours. Publish the four weeks as an artifact: a calendar grid by channel, review dates marked, hours per week. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

**NOTE** Frequency is a capacity decision as much as an audience one. A calendar the team cannot deliver is fiction.

# Growth Analyst and Strategist: Cadence and Calendar (cont.)

**FOLLOW-UP** — type any of these in the chat

Why is the newsletter not weekly? Show me the numbers.

What happens if we add a third LinkedIn post a week?

Which week is busiest, and does the team have the hours?

**NOTE** Ask until the result makes sense to you — then decide.

# Growth Analyst and Strategist: Cadence and Calendar

Check your work before moving on.

## CHECK YOUR WORK

- ✓ The cadence gives a frequency per channel with the evidence.
- ✓ A third LinkedIn post a week costs far more unfollows than it earns.
- ✓ It explains the unsubscribe jump at weekly emails.
- ✓ The calendar covers four weeks in the required columns.
- ✓ Every review\_by date is two business days before publishing.
- ✓ No week is over 12 hours.

# From Storyboard to Content

Lab 8: the first real piece — a blog article every post can link to.

## One prompt, the whole job

Research, draft, fact-check, fix and post — a quick blog in one go. A person still reads it last.

## Straight onto your copy of the site

The article lands on the copy of Horizon's site in your folder, opened in your browser. The public site never changes.

## Your skills know Horizon

blog-post, sunny-voice, fact-check and fin-compliance: the voice, the facts and the rules.

# Content Creator: A Quick Blog Article

LAB 08

DAY 1 • ~40 MIN

USE: CLAUDE CODE

## THE STORY SO FAR

Jun Wei has the storyboard and the personas. Rachel wants the first real piece today: a blog article that people searching for CPF top-ups will find — and that every post can link to.

## YOU'LL BUILD

**a 500-700-word blog article with dated sources, checked and fixed and posted on your own copy of Horizon's website, opened in your browser, with a Word copy**

Surface: Claude Code (Firecrawl, your Horizon skills) → your copy of Horizon's website

Features: one prompt runs the whole job: research • draft • fact-check and fin-compliance • fix • post to the blog on your copy of the website • a person reads it last

# Content Creator: A Quick Blog Article

USE CLAUDE CODE • copy each prompt from the Prompts page (link in your lab README).

- 1 New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 8". Run all of this lab's prompts in it.
- 2 Write and post it** Copy and paste the prompt. Claude researches, writes, checks and fixes the article, then posts it to your copy of Horizon's website.
- 3 Read it on the site** Your copy of the website opens in your browser: your article is first under "From our blog". Say what to change, if anything.

# Content Creator: A Quick Blog Article

## PROMPT — Claude Code: a quick blog

Write a quick blog article for Horizon and post it to the website. Use the blog-post, sunny-voice, fact-check and fin-compliance skills.

1. Pick one topic from strategy/storyboard.md that a Singaporean would search for. Use Firecrawl to find two current sources (CPF Board, MAS or MoneySense), each with its date.
2. Write 500-700 words: a title under 60 characters, every figure from data/facts-2026.md, one call to action (book a free 30-minute chat) and the full disclaimer → content/blog/<slug>.md.
3. Run fact-check and fin-compliance; fix every issue and list each fix in one line.
4. Add it to the blog in web/site/index.html as one new post, built like the posts already there. Change nothing else on the site. Then open web/site/index.html in my web browser. Write for a business owner, in plain English. Save a Word copy (.docx) of the article.

**NOTE** One prompt can run the whole job — research, writing and checks. A person still reads the result before it goes anywhere.

# Content Creator: A Quick Blog Article (cont.)

**FOLLOW-UP** — type any of these in the chat

Summarise the article in three sentences. Who is it for?

Why did you pick this topic, and who will search for it?

What did the checks find, and how did you fix it?

**NOTE** Ask until the result makes sense to you — then decide.

# Content Creator: A Quick Blog Article

Check your work before moving on.

## CHECK YOUR WORK

- ✓ Two sources, each with a date and a link.
- ✓ The title is under 60 characters.
- ✓ Every figure is from Horizon's 2026 facts.
- ✓ The fixes are listed; no Critical or High issue is left.
- ✓ One call to action and the full disclaimer at the end.
- ✓ Your article is first under "From our blog" on your copy of the website.

# Day 1: The Team Is Ready

Everything the campaigns need tomorrow is in place.




# Multi-Channel Content Creation and Agent Workflow Coordination

---

- 1 Subagents or an agent team?
- 2 The lead magnet: a checklist people want, and its emails
- 3 Human in the loop — enforced, not requested
- 4 Campaigns 2–4: social week, newsletter + landing page, video
- 5 From approval to LinkedIn, Facebook, email and YouTube

# Content That Earns a Sign-up

Lab 9: the Money Check-up Checklist and the emails that follow it.



**TIP** Lab 11 builds the landing page that offers it.

# The Lead Magnet: a Checklist and Its Welcome Emails

LAB 09

DAY 2 • ~40 MIN

USE: CLAUDE CODE

## THE STORY SO FAR

Day 2. Before the campaign weeks start, Rachel wants the thing that turns readers into leads: the Money Check-up Checklist, designed to be downloaded, and the emails that follow it.

## YOU'LL BUILD

**the Money Check-up Checklist as a printable A4 page (artifact and PDF) and on your copy of Horizon's website, and three welcome emails — shown as artifacts, with a Word copy of each**

Surface: Claude Code (artifact, Gmail drafts, Marketing plugin, your Horizon skills)

Features: lead-magnet skill • a printable A4 checklist as an artifact • the Marketing plugin's email sequence • newsletter and sunny-voice • drafts only

# The Lead Magnet: a Checklist and Its Welcome Emails

USE CLAUDE CODE • copy each prompt from the Prompts page (link in your lab README).

- 1 New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 9". Run all of this lab's prompts in it.
- 2 Design it** Copy and paste Prompt A. The checklist opens as an artifact; ask in the chat for any change.
- 3 Write the welcome emails** After the checklist looks right, copy and paste Prompt B. The three emails also appear as drafts in your Gmail.
- 4 Check it** After the three emails are saved, copy and paste Prompt C.
- 5 Put it on the website** After the checks pass, copy and paste Prompt D. Your copy of the website opens in your browser, with the checklist behind an email box.

# The Lead Magnet: a Checklist and Its Welcome Emails

## PROMPT A — Claude Code: the checklist

Use the lead-magnet skill. Turn data/checklist-items.md into the "Money Check-up Checklist": a title with one promise, the ten items as tick boxes, one figure from data/facts-2026.md, the next step (book a free chat) and the full disclaimer. Lay it out as a printable A4 page in the colours and fonts in data/brand.md and publish it as an artifact. Save a PDF of it in content/lead-magnet/. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

# The Lead Magnet: a Checklist and Its Welcome Emails (cont.)

## PROMPT B — Claude Code: the welcome emails

Use the Marketing plugin's email sequence to write three welcome emails for people who download the checklist: day 0 (here it is), day 3 (the item most people miss — the retirement number), day 7 (book a free chat). Follow the newsletter and sunny-voice skills → `content/email/welcome-sequence.md`. Then, using the Gmail connector, create the three emails as DRAFTS in my Gmail, addressed to me. Send nothing. Publish the three emails as an artifact: one tab per email, laid out as it will look in the inbox. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

# The Lead Magnet: a Checklist and Its Welcome Emails (cont.)

## PROMPT C — Claude Code: check it

Run fact-check and fin-compliance on the checklist text and the three emails. Show the issues in one table, with a fix for each. Once I agree, make the fixes in the files and in the Gmail drafts. Publish the table as an artifact, coloured by severity. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

# The Lead Magnet: a Checklist and Its Welcome Emails (cont.)

## PROMPT D — Claude Code: put it on the website

Add the checked Money Check-up Checklist to Horizon's website in web/site/index.html, as a new section before the reviews, in the site's style:

the title and promise, what's inside, and an email box with a required consent tick box. Once a valid email is entered, show the ten items as tick boxes and a "Print or save as PDF" button. Add "Free checklist" to the menu. Change nothing else on the site. Then open web/site/index.html in my web browser.

**NOTE** A lead magnet earns an email address. Give real value first; the call to action comes last.

# The Lead Magnet: a Checklist and Its Welcome Emails (cont.)

**FOLLOW-UP** — type any of these in the chat

Which checklist item do people most often miss, and why does it matter?

Read the day 3 email as a 55-year-old. Is anything unclear?

Explain each issue in one line: what is wrong, and the fix.

**NOTE** Ask until the result makes sense to you — then decide.

# The Lead Magnet: a Checklist and Its Welcome Emails

Check your work before moving on.

## CHECK YOUR WORK


- ✓ The checklist uses Horizon's colours and all ten items.
- ✓ It makes one promise and ends with the full disclaimer.
- ✓ Three emails, each with a subject under 50 characters, are drafts in your Gmail — nothing was sent.
- ✓ Each email has one call to action, unsubscribe and the disclaimer.
- ✓ fin-compliance shows no Critical or High issue.
- ✓ The checklist is on your copy of the website: email and consent first, then the ten tick boxes.

# Subagents vs Agent Teams

Both run agents in parallel, each with its own context. The difference is who they talk to.

## SUBAGENTS


Report back to one main agent



Each works in its own context and returns only its result. They do not talk to each other.

## AGENT TEAM

A lead, a shared task list, teammates who talk



Teammates claim tasks and message each other directly — e.g. a reviewer sends fixes straight to the writer.

Use subagents when only the result matters. Use a team when the agents need to talk.

# Which Pattern When?


From the Claude Code documentation, applied to the team.

|              | SUBAGENTS                                    | AGENT TEAM                                       |
|--------------|----------------------------------------------|--------------------------------------------------|
| Talks to     | Returns a result to the Lead                 | Teammates message each other directly            |
| Coordination | The Lead manages everything                  | A shared task list; teammates claim tasks        |
| Best for     | Research, analysis, one-off reviews          | Copy ↔ design ↔ review loops                     |
| Cost         | Lower: results are summarised back           | Higher: every teammate is its own Claude session |
| Campaigns    | Lab 5 (and sub-agents in Claude: Labs 6, 15) | Labs 10, 11, 12                                  |

TIP Agent teams are experimental; your studio folder already has them switched on.

# Running an Agent Team

In Claude Code. Describe the team in plain words.



**TIP** Teammates load CLAUDE.md, skills and connectors — but not the Lead's conversation. One team per session.

# Human in the Loop, Enforced

A rule in a prompt persuades. The publisher kit builds a gate that holds.

## Only people approve

You type "approved <id> by <your name>". The gate reads your own message — an agent cannot write one.

## The approved words go out


The approval stores a hash. Edit one word and publish --live refuses until a person approves again.

## A hook stops agents

The gate blocks agents from approving, editing the approval records, or publishing unapproved items.

# How the Gate Works

Already in your studio folder. You never run it — you just type.



**TIP** The gate is a hook: it reads your own message, and blocks any agent that tries to approve or publish on its own.

# The Publishing APIs


publish.mjs is a dry run by default. Tokens live only in .env.

| CHANNEL    | ENDPOINT                                                    | PERMISSION                                      |
|------------|-------------------------------------------------------------|-------------------------------------------------|
| LinkedIn   | POST api.linkedin.com/rest/posts (LinkedIn-Version: YYYYMM) | w_member_social — your own profile              |
| Facebook   | POST graph.facebook.com/<version>/<page-id>/feed            | Page token: pages_manage_posts                  |
| YouTube    | Resumable upload to youtube/v3/videos, privacy private      | youtube.upload; unaudited projects stay private |
| Newsletter | No API — HTML + text files, then Gmail drafts               | Consent = yes, unsubscribed = no                |

TIP No accounts? Every lab works in dry-run mode: you see the exact request that would be sent.

# Campaign 2: A Social Media Week

Lab 10, a demo: one LinkedIn and one Facebook post, each with a visual.



# Campaign 2: Social Media Posts (Demo)

LAB 10

DAY 2 • ~40 MIN

USE: CLAUDE CODE

## THE STORY SO FAR

Week 1 of the campaign starts Monday. As a demo, one LinkedIn post and one Facebook post, each with a visual — and Rachel's one rule: nothing goes out unless a person approved that exact version.

## YOU'LL BUILD

**two posts with images, reviewed, approved by you and posted (or dry-run) — shown as artifacts, with a Word copy of each**

Surface: Claude Code (agent team) → LinkedIn, Facebook

Features: the publisher kit and its hook • an agent team: creator, designer, Lead • teammates message each other • human approval → post

# Campaign 2: Social Media Posts (Demo)

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

- 1 New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 10". Run all of this lab's prompts in it.
- 2 The gate is on** Your studio folder came with the approval gate and agent teams switched on — nothing to install.
- 3 Test the gate** Ask Claude to approve a post itself. The gate must block it.
- 4 Start the team** Copy and paste Prompt A. Teammates appear in the panel; ↑ ↓ and Enter open one; Ctrl+T shows the task list.
- 5 Watch them talk** The designer asks the creator for each post's hook; the Lead sends fixes back.
- 6 Approve as a person** Review in the artifact (or open the review email and click its link). Click Approve or Request changes, then paste the copied line into the chat.
- 7 Publish** After you have approved the posts, copy and paste Prompt B. A dry run first; real posting only if your accounts are connected.

# Campaign 2: Social Media Posts (Demo)

## PROMPT A — Claude Code: the social team

Create an agent team for Week 1 social from data/social-brief.md and strategy/calendar.csv. Spawn two teammates using the agent types content-creator and creative-designer; you are the Lead.

- content-creator: 1 LinkedIn + 1 Facebook post with linkedin-post, facebook-post and copywriting, one file each in content/social/week-01/, with the frontmatter in data/publishing-spec.md.
- creative-designer: one image per post (a rendered card in Horizon colours) with alt text; message the creator for each post's hook before designing.
- You: fact-check and fin-compliance on every post; message fixes to the teammate who owns it. Before you start, ask me how I will review: in the artifact, or by email (then ask for the approver's address). When a post passes, submit it for review with scripts/submit.mjs. Publish the two posts as an artifact, each as it will look in the feed, with its image and alt text; under each post put an Approve button and a Request changes box that copy the exact line for me to paste in the chat. If I chose email, send ONE review email to the approver with Gmail: the artifact link and what passed the checks. Send nothing else; publish nothing. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

## Campaign 2: Social Media Posts (Demo) (cont.)

### PROMPT B — Claude Code: publish

For every Week 1 post approved in review/approvals.csv, do a dry run of publishing it (scripts/publish.mjs) and publish each request as an artifact: one card per post with channel, time and link. Wait for me to say "go" before anything goes live.

**NOTE** Subagents report to one boss; teammates also talk to each other. A designer who can ask the writer for the hook makes a better image.

## Campaign 2: Social Media Posts (Demo) (cont.)

**FOLLOW-UP** — type any of these in the chat

Which post do you expect to do best, and why?

What did the Lead send back to the writer, and why?

Is anything here not ready to post? Explain in plain English.

**NOTE** Ask until the result makes sense to you — then decide.

# Campaign 2: Social Media Posts (Demo)


Check your work before moving on.

## CHECK YOUR WORK

- ✓ The gate stopped Claude from approving.
- ✓ Two teammates ran; the task list showed their tasks.
- ✓ The designer and the creator messaged each other.
- ✓ Two posts with images, alt text, frontmatter and UTM links.
- ✓ You reviewed in the artifact or by email and pasted the approval in the chat; the rows show your name and a hash.
- ✓ Dry runs printed every request; nothing went live unless you said go.

# Campaign 3: Newsletter + Landing Page

Lab 11: three specialists, one email and one page, consent first.



# Consent Decides Who Gets Mail

The PDPA, applied by an agent that counts.

## Opt-in only

consent = yes and unsubscribed = no.  
Everyone else goes to excluded.csv with a reason.

## Every email carries

An unsubscribe link, Horizon's address  
and the short disclaimer.

## The page collects little

Name, email, a required consent box, a  
separate unticked newsletter opt-in.

# Campaign 3: Newsletter and a Landing Page

LAB 11

DAY 2 • ~45 MIN

USE: CLAUDE CODE

## THE STORY SO FAR

The Sunny Sunday email books chats most cheaply. November's issue promotes a new landing page for the free checklist — and must reach only the 31 people who opted in.

## YOU'LL BUILD

**the November newsletter, its recipient list, the landing page published, and Gmail drafts — shown as artifacts, with a Word copy of each**

Surface: Claude Code (agent team) → Claude artifact, Gmail

Features: creator writes, website designer builds the lead-magnet page, analyst filters by consent, the Lead checks · drafts only

# Campaign 3: Newsletter and a Landing Page

USE CLAUDE CODE • copy each prompt from the Prompts page (link in your lab README).

- 1 New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 11". Run all of this lab's prompts in it.
- 2 Start the campaign** Copy and paste Prompt A.
- 3 Approve the words** Review the issue in the artifact (or open the review email and click its link). Click Approve or Request changes, then paste the copied line into the chat.
- 4 Publish the page** The website designer runs page-qa and publishes web/checklist/ as a Claude artifact.
- 5 Check the list** Only people with consent = yes and unsubscribed = no are on the recipient list.
- 6 Create the drafts** After the page is published and the list checked, copy and paste Prompt B.
- 7 Send one test** Send only the draft addressed to yourself, and read it on your phone.

# Campaign 3: Newsletter and a Landing Page

## PROMPT A — Claude Code: the newsletter team

Create an agent team for the November Sunny Sunday email (data/newsletter-brief.md, data/newsletter-spec.md) and its landing page (data/landing-page-brief.md). Spawn three teammates: content-creator, website-designer and growth-analyst.

- content-creator: issue.md (id nl-2026-11) with newsletter, and the landing-page copy.
- website-designer: web/checklist/index.html with lead-magnet, web-design and landing-page from the approved copy; page-qa.
- growth-analyst: recipients.csv and excluded.csv from data/subscribers.csv. Before you start, ask me how I will review: in the artifact, or by email (then ask for the approver's address). You check everything, then submit nl-2026-11 and publish the issue as an artifact, laid out as the email, with the counts of recipients and excluded, an Approve button and a Request changes box that copy the exact line for me to paste in the chat. If I chose email, send ONE review email to the approver with Gmail: the artifact link and what passed the checks. Only after I approve, build newsletter.html and .txt and publish the page as an artifact. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

## Campaign 3: Newsletter and a Landing Page (cont.)

### PROMPT B — Claude Code: Gmail drafts

Using the Gmail connector, create a DRAFT of newsletter.html for each address in recipients.csv, plus one draft addressed to me only. Send nothing. Publish a report as an artifact: how many drafts, how many people were excluded, and why, in a table.

**NOTE** Consent decides who gets mail — not the agent and not the deadline. Marketing to people who did not opt in breaks the PDPA.

## Campaign 3: Newsletter and a Landing Page (cont.)

**FOLLOW-UP** — type any of these in the chat

Explain the email in three sentences. What do we want readers to do?

Who was left off the mailing list, and why?

What did the landing-page check look at, and did it pass?

**NOTE** Ask until the result makes sense to you — then decide.

# Campaign 3: Newsletter and a Landing Page

Check your work before moving on.

## CHECK YOUR WORK

- ✓ Three specialists worked on the campaign; the Lead checked all.
- ✓ You reviewed nl-2026-11 in the artifact or by email and approved it in the chat before the build.
- ✓ The landing page passed page-qa and is published.
- ✓ The recipient list holds 31 people; 9 are excluded, with reasons.
- ✓ The email has unsubscribe, address and disclaimer.
- ✓ Drafts only — nothing was sent to the list.

# Campaign 4: The YouTube Explainer

Lab 12: the team writes and designs; you approve twice.

## Strategist and Creator

The angle, then 5–7 scenes with words on screen and a description with chapters and an AI disclosure.

## Creative Designer

Storyboard, thumbnail, then the render with video-render — ffmpeg, 1920×1080, under 90 s.

## Upload privately

Approve the words, then the finished video. It uploads as private; a person makes it public.

# Campaign 4: The YouTube Explainer

LAB 12

DAY 2 • ~45 MIN

USE: CLAUDE CODE

## THE STORY SO FAR

Pre-retirees say they would rather watch than read. Rachel wants the first 'Money in Plain English' video, uploaded privately until she has watched it.

## YOU'LL BUILD

**the script, storyboard and thumbnail, the video, and a private YouTube upload (or dry run) — shown as artifacts, with a Word copy of each**

Surface: Claude Code (agent team, render, upload)

Features: strategist angle • creator script • designer storyboard, thumbnail and render • approve twice • private upload

# Campaign 4: The YouTube Explainer

USE CLAUDE CODE • copy each prompt from the Prompts page (link in your lab README).

- 1 New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 12". Run all of this lab's prompts in it.
- 2 Script it as a team** Copy and paste Prompt A.
- 3 Approve the script** Read the scenes, title and description in the storyboard artifact; type: approved yt-ep01-script by <your name>.
- 4 Render it** After you approve the script, copy and paste Prompt B. The designer runs video-render.
- 5 Approve the video** Ask the Lead to submit the video, watch ep01.mp4 to the end, then approve yt-ep01.
- 6 Upload** After you approve the video, copy and paste Prompt C. It goes up as private.
- 7 Disclose** In YouTube Studio, answer the altered or synthetic content question before making it public.

# Campaign 4: The YouTube Explainer

## PROMPT A — Claude Code: the video team

Create an agent team for YouTube episode 1 from data/youtube-brief.md. Spawn three teammates: growth-strategist (the angle and hook), content-creator (content/video/ep01/scenes.json, 5-7 scenes, and ep01.md with title, tags, description, chapters and the AI disclosure) and creative-designer (a storyboard of the scenes and a thumbnail, using brand-visuals). You fact-check and fin-compliance the script, then submit it as yt-ep01-script. Publish nothing. Publish the storyboard as an artifact: one frame per scene with its words, visual and timing, and the thumbnail. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

## Campaign 4: The YouTube Explainer (cont.)

### PROMPTS B and C — Claude Code: render and upload

#### PROMPT B

Use the creative-designer subagent with video-render on the approved scenes.json → content/video/ep01/ep01.mp4 (data/video-spec.md). Check its length and size. Do not change approved words. Publish the checks as an artifact: length, size, resolution and every scene, PASS or FAIL.

#### PROMPT C

Do a dry run of the YouTube upload for yt-ep01 (scripts/publish.mjs). If I say "go", upload it for real — as PRIVATE. Publish the request, then the YouTube Studio link, as an artifact.

**NOTE** Approve twice: once for the words, once for the finished video. Rendering can break what the script got right.

## Campaign 4: The YouTube Explainer (cont.)

**FOLLOW-UP** — type any of these in the chat

Walk me through the video scene by scene, in plain English.

Which figure on screen should I double-check, and against what?

What does the AI disclosure say, and why do we need it?

**NOTE** Ask until the result makes sense to you — then decide.

# Campaign 4: The YouTube Explainer

Check your work before moving on.

## CHECK YOUR WORK

- ✓ The strategist, creator and designer each delivered their part.
- ✓ You approved the script, then the finished video.
- ✓ ep01.mp4 is 1920x1080, under 90 seconds, every scene present.
- ✓ Every figure on screen matches Horizon's 2026 facts.
- ✓ The description has chapters, link, disclaimer and AI disclosure.
- ✓ The upload is private (or the dry run printed the request).

# Where You Are Now

Three campaigns, every piece checked by the Lead and approved by a person. But the team stops when you close the laptop.

## Gate · Lab 10

People approve, hashes hold, a hook stops agents.

## Campaigns · Labs 10–12

Social week, newsletter + landing page, the video.

## What is missing

A team that works every week on its own — and rules to govern it. That is Topic 4.


# Content Distribution, Strategy Guidelines and Responsible AI Practices

---

- 1 Campaign 5: always on, with a scheduled task
- 2 Responsible AI for financial content
- 3 Measure the patterns; write the team playbook
- 4 Agent teams at work: argue a strategy, then a QA loop

# From Skill to Schedule

Lab 13: do the job once, keep it as a skill, let a schedule run it.



**TIP** Pause scheduled tasks you are not using.

# Automate the Preparation, Not the Decision

What a scheduled agent should and should not do for a regulated firm.

## **Prepare**

Read the data, run the skill, draft the report, flag what needs a decision.

## **Draft, never send**

Write "draft only — never send, publish or approve" in the skill AND the schedule.

## **Keep a human gate**

The schedule fills an inbox. A named person still decides.

# Campaign 5: Always On — the Weekly Growth Report

LAB 13

DAY 2 • ~35 MIN

USE: CLAUDE CODE

## THE STORY SO FAR

The team must keep working after class. Rachel wants the Growth Analyst's report in her inbox every Monday — as a draft she decides on.

## YOU'LL BUILD

**a scheduled task that drafts the weekly growth report every Monday at 8am, and one draft in Gmail**

Surface: Claude Code (skill, routine, Drive, Gmail)

Features: the Growth Analyst's campaign-report skill • do • verify • schedule • draft-only delivery

# Campaign 5: Always On — the Weekly Growth Report

USE CLAUDE CODE · copy each prompt from the Prompts page (link in your lab README).

- 1 New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 13". Run all of this lab's prompts in it.
- 2 Put the data on Drive** Upload campaign-results.csv to the Horizon Studio folder.
- 3 Do it once** Copy and paste Prompt A. Refine the report until it is right.
- 4 Schedule it** After the report is right, copy and paste Prompt B into the instructions of a new routine: Routines → New routine → Local, Mondays at 8am.
- 5 Run it now** Click Run now once, approve any permission prompts, then open the Gmail draft.
- 6 Pause it** After class, pause the routine.

# Campaign 5: Always On — the Weekly Growth Report

## PROMPT A — Claude Code: do it once

Act as Horizon's Growth Analyst. Using Google Drive, read campaign-results.csv in the Horizon Studio folder. Use the campaign-report skill and data/weekly-report-brief.md. Leave the report as a Gmail DRAFT to me, subject "Horizon weekly growth report" and today's date. Send nothing. Also publish the report as an artifact: cost per chat by channel as a chart, and the budget split. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

## Campaign 5: Always On — the Weekly Growth Report (cont.)

### PROMPT B — Claude Code: the routine

Run the campaign-report skill as Horizon's Growth Analyst. Read campaign-results.csv in the Horizon Studio folder on Google Drive and follow weekly-report-brief.md. Leave the report as a draft email to me in Gmail with the subject "Horizon weekly growth report" and today's date. Do not send it, and do not publish or approve anything.

**NOTE** Automate the preparation, not the decision. The schedule fills an inbox; a person still decides.

## Campaign 5: Always On — the Weekly Growth Report (cont.)

**FOLLOW-UP** — type any of these in the chat

What is the one decision this report asks me to make?

Which channel gets us a chat most cheaply, and how sure are you?

What would you change in next month's budget, and why?

**NOTE** Ask until the result makes sense to you — then decide.

# Campaign 5: Always On — the Weekly Growth Report

Check your work before moving on.

## CHECK YOUR WORK

- ✓ The report used the campaign-report skill from your folder.
- ✓ The report names the cheapest channel per chat (the newsletter).
- ✓ It recommends a budget split with a reason per channel.
- ✓ A scheduled task runs Mondays at 8am with draft-only instructions.
- ✓ A manual run left a draft — nothing was sent.
- ✓ The task is paused after class.

# Responsible AI in the Team

Each risk has a control you built in this course.

| RISK                          | CONTROL                                          | WHERE            |
|-------------------------------|--------------------------------------------------|------------------|
| Wrong figure                  | fact-check skill; the Lead's review; facts sheet | Labs 3, 5, 10–12 |
| Read as advice or a promise   | fin-compliance; disclaimers; no returns          | Labs 3, 10–12    |
| An agent with too much access | Least-privilege connectors per agent             | Lab 4            |
| Unapproved post               | Approval typed by a person, hash, gate           | Lab 10           |
| Marketing without consent     | The Analyst filters by consent                   | Lab 11           |
| Undisclosed AI media          | Description line + YouTube disclosure            | Lab 12           |

TIP Keep keys in .env, client data out of prompts, and the approval records for audit.

# Measure What Matters

Lab 14: the numbers decide next month — and which patterns earn their tokens.

## Cost per chat

Spend ÷ chats booked shows which channel works: email first, YouTube last.

## Defects before a person

A Lead that catches mistakes saves the reviewer's time; one agent alone caught none.

## The playbook

Roster, instructions, skills, connectors, approval matrix, incidents — so the team runs without you.

# Measure, Govern and Write the Team Playbook

LAB 14

DAY 2 • ~40 MIN

USE: CLAUDE CODE

## THE STORY SO FAR

Eight weeks in. Rachel asks which ways of running agents earned their tokens, and how the team runs safely after you leave.

## YOU'LL BUILD

**the agent-pattern comparison and the team playbook — shown as artifacts, with a Word copy of each**

Surface: Claude Code (subagents)

Features: compare agent patterns • the roster: instructions, skills, connectors • approval matrix • responsible-AI controls

# Measure, Govern and Write the Team Playbook

USE CLAUDE CODE • copy each prompt from the Prompts page (link in your lab README).

- 1 New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 14". Run all of this lab's prompts in it.
- 2 Check the data** The assets are already in data/. Open one and skim it.
- 3 Compare the patterns** Copy and paste Prompt A.
- 4 Write the playbook** After the Lead shows the comparison, copy and paste Prompt B.
- 5 Review it as Rachel** Add one rule of your own and say why.

# Measure, Govern and Write the Team Playbook

## PROMPT A — Claude Code: compare

Use the growth-analyst subagent. From data/run-log.csv compare the patterns we used — one agent, subagents, agent teams and the scheduled task — on minutes, tokens, defects caught before a person and human edits. Which pattern pays off for which job? Save reports/agent-comparison.md. Publish it as an artifact: one chart per measure, comparing the patterns. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

# Measure, Govern and Write the Team Playbook (cont.)

## PROMPT B — Claude Code: the playbook

Write strategy/team-playbook.md with data/playbook-outline.md. Include the roster: for each agent its instructions file in .claude/agents/, its skills, connectors, model and what it may never do; the approval matrix by risk; the rules in data/responsible-ai-checklist.md; and what to do when a wrong post goes out. Then run fin-compliance on it. Publish the playbook as an artifact: the roster as cards, the approval matrix as a table. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

**NOTE** Teams cost more tokens. The run log shows what you buy with them: defects caught before a person, and fewer edits.

# Measure, Govern and Write the Team Playbook (cont.)

**FOLLOW-UP** — type any of these in the chat

In one line each: when should we use one agent, subagents and an agent team?

Who must approve a post that has a CPF figure in it?

A wrong post went out. What are the first three things we do?

**NOTE** Ask until the result makes sense to you — then decide.

# Measure, Govern and Write the Team Playbook

Check your work before moving on.


## CHECK YOUR WORK

- ✓ The comparison shows the one-agent baseline caught 0 defects.
- ✓ It says which pattern suits which job, with numbers.
- ✓ The roster lists all six agents with skills and connectors.
- ✓ The approval matrix says who approves CPF and tax figures.
- ✓ It has an incident procedure, and your own rule.


# Subagents vs Agent Teams

Same work, different wiring: who spawns whom, and who talks to whom.

## SUBAGENTS



## AGENT TEAM



Subagents report back to one agent. Teammates share a task list and talk to each other.

# A Team That Argues

Lab 15, Part 1: test a strategy before Horizon spends on it.

## Advocate • growth-strategist


- Argues FOR the "Turning 55" CPF explainer
- Cites results, personas and the facts sheet
- Answers the critic directly — two rebuttals

## Critic • growth-analyst

- Argues AGAINST it — cost, timing, audience
- Every point backed by a file, not an opinion
- The Lead judges; a person makes the call

# The QA Loop

Lab 15, Part 2: a failed audit goes straight back to the writer — not through the Lead.



## THE AUDIT FAILS WHEN

- 1 A figure is not in facts-2026.md
- 2 A promise of savings or returns
- 3 Advice ("everyone should...")
- 4 A wrong age or a missing disclaimer

Three rounds at most — then the Lead escalates to a person.

# Agent Teams: A Strategy Debate and a QA Loop

LAB 15

DAY 2 • ~30 MIN

USE: CLAUDE CODE

## THE STORY SO FAR

January is for people turning 55. Before Horizon spends on a CPF explainer for them, Rachel wants it argued both ways — and the intern's draft checked until it passes.

## YOU'LL BUILD

**the decision (pros, cons, verdict), then an email, a video script and a thumbnail that passed QA, submitted for your approval — shown as artifacts, with a Word copy of each**

Surface: Claude Code (agent teams: Lead, teammates, shared task list, direct messages)

Features: an advocate and a critic argue the pros and cons • the Lead judges • a writer and a QA Auditor loop until PASS • FAIL goes straight back to the writer • human approval

# Agent Teams: A Strategy Debate and a QA Loop

USE CLAUDE CODE • copy each prompt from the Prompts page (link in your lab README).

1

**New session** Left sidebar: click + next to horizon-studio. Click the session title at the top and rename it "Lab 15". Run all of this lab's prompts in it.

2

**Part 1: the debate** Copy and paste Prompt A. The advocate and the critic message each other; Ctrl+T shows the task list.

3

**Decide** Read the decision artifact. Reply "go", "go with conditions" or "no-go" — the call is yours.

4

**Part 2: the QA loop** After you have made the call, copy and paste Prompt B. The audit fails on round 1; watch the auditor message the writer directly.

5

**Approve as a person** Read the pieces and the QA artifact, then type in the chat: approved <id> by <your name> (or: changes <id>: what to fix).

6

**Shut down** After you have approved the pieces, copy and paste Prompt C. The Lead stops each teammate and cleans up the team.

# Agent Teams: A Strategy Debate and a QA Loop

## PROMPT A — Claude Code: the strategy debate

Create an agent team to test the strategy in data/turning-55-brief.md before we spend on it. You are the Lead and the judge: do not argue. Spawn two teammates:

- advocate, using the agent type growth-strategist:  
argue FOR the strategy.
- critic, using the agent type growth-analyst:  
argue AGAINST it.
- Both: cite data/campaign-results.csv, strategy/personas.md, data/facts-2026.md or data/compliance-checklist.md for every point; message each other directly; rebut twice at most.
- You: write strategy/turning-55-decision.md with the pros, the cons, your verdict (go / go with conditions / no-go) and the conditions. Publish the decision as an artifact with the verdict on top, the pros and cons in two columns, and what I do next. Then wait for my decision. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

# Agent Teams: A Strategy Debate and a QA Loop (cont.)

## PROMPT B — Claude Code: the QA loop

Clean up the debate team, then create a new agent team for Part 2 of data/turning-55-brief.md, using my decision. Spawn three teammates:

- writer, agent type content-creator: rewrite data/turning-55-draft.md into content/turning-55/email.md and content/turning-55/video-script.md.
- qa-auditor: audit each piece with fact-check and fin-compliance; write review/turning-55-qa.md, one section per round, PASS or FAIL per check. On FAIL, message the writer directly with the lines to fix and audit again. Three rounds at most.
- designer, agent type creative-designer: once a piece passes, ask the writer for its headline and make one thumbnail with alt text.
- You, the Lead: never write or fix content. When everything passes, submit each piece with scripts/submit.mjs. Publish nothing. Publish the QA report as an artifact, one row per round, PASS or FAIL per check, in colour. Write for a business owner, in plain English: no codes, IDs, file names or line numbers; start with what it means and what I do next. Save a Word copy (.docx) of each result.

# Agent Teams: A Strategy Debate and a QA Loop (cont.)

## PROMPT C — Claude Code: shut down

Ask every teammate to shut down, then clean up the team. Show me the final task list and how many QA rounds each piece needed. Publish it as an artifact.

**NOTE** Use subagents when only the result matters; use a team when the agents must talk — to argue, or to send work back. Every teammate is its own Claude session, so teams cost more tokens.

# Agent Teams: A Strategy Debate and a QA Loop (cont.)

**FOLLOW-UP** — type any of these in the chat

Summarise the debate: the best point for, the best point against, and your verdict.

What failed in the first QA round, in plain English?

How many rounds did each piece need, and why?

**NOTE** Ask until the result makes sense to you — then decide.

# Agent Teams: A Strategy Debate and a QA Loop


Check your work before moving on.

## CHECK YOUR WORK

- ✓ The advocate and the critic messaged each other; every point cites a file.
- ✓ The decision has pros, cons, a verdict — and you made the call.
- ✓ Round 1 of the QA report FAILED: the FRS figure, CPF LIFE "from 55", the promised payout, the advice and the missing disclaimer.
- ✓ The auditor sent the fixes to the writer directly; a later round PASSED.
- ✓ You approved by typing in the chat; the rows show your name and a hash.
- ✓ Both teams were shut down and cleaned up.

# How Each Pattern Is Triggered


Five ways your team runs. The difference is what sets each one off.



Who decides? You: the approval · the model: delegation and skills · the system, every time: hooks and the clock

# Horizon: One Team, Five Campaigns

Two days, one regulated business — the Claude Edition.



# Courseware & Assessment on the LMS



## Download the courseware

Get the slides and learner guide from the LMS for the open-book assessment.



## Take the assessment

Complete the Written Assessment and Practical Performance on the LMS.



## Log in

Use your enrolled account to access the course materials and assessment.

A screenshot of the Tertiary Infotech Academy LMS login interface. It features a dark blue background with a white logo at the top. Below the logo, the text 'Tertiary Infotech Academy' and 'LMS cum TMS for WSQ Courses' is displayed. There is an 'Email' input field, a 'Send OTP' button, and a link to 'Login with password instead'. At the bottom, there are links for 'Privacy Policy', 'Acceptable Use Policy', and 'Feedback', along with the text 'Powered by Tertiary Infotech Academy Pte Ltd'.

**LMS:** <https://lms-tms.tertiaryinfotech.com/>

# Certification & TRAQOM Survey



## TRAQOM Survey (Mandatory)

Complete the certification & TRAQOM survey to receive your certificate.

<https://lms-tms.tertiaryinfotech.com/>



## Digital Attendance

Remember to take AM, PM and Assessment digital attendance — required for your funding.

# Practice Exam: Claude Certified Associate

[exams.tertiaryinfotech.com/practice-exams/anthropic/anthropic-ccao-foundations](https://exams.tertiaryinfotech.com/practice-exams/anthropic/anthropic-ccao-foundations)

The screenshot shows the Tertiary Exams website interface for the Claude Certified Associate — Foundations (CCAO-F) practice exam. The page includes a navigation bar with links to Practice Exams, Vendors, About Us, FAQ, and Support, along with a 'Get started' button. The main content area features a breadcrumb trail (All exams / Anthropic / Anthropic / CCAO-F / Foundational) and a detailed description of the exam bundle. Key details include 60 questions, a 120-minute duration, and a 72% pass score. The page also offers a free practice teaser and a paid exam mode for SGD \$20.00. A table lists the domains covered and their respective percentages.

| Domains covered                          |     |
|------------------------------------------|-----|
| Prompting and Task Execution             | 14% |
| Output Evaluation and Validation         | 21% |
| Product and Model Selection              | 12% |
| Workflow Integration and Solution Design | 16% |

**PRACTICE EXAM** [exams.tertiaryinfotech.com/practice-exams/anthropic/anthropic-ccao-foundations](https://exams.tertiaryinfotech.com/practice-exams/anthropic/anthropic-ccao-foundations)

## CCAO-F · Foundations

For professionals who use Claude as a productivity tool.

## What it covers

Prompting and task execution, output evaluation, model selection, workflow integration, responsible use.

## Try it free

Start with the free practice teaser; exam mode is 60 questions in 120 minutes.

# Assessment & Funding



## Final Assessment

- Written Assessment (WA) — 1 hour, starts 4:00 PM
- Practical Performance (PP) — 1 hour
- Format: Open book — class ends 6:00 PM
- Allowed: slides, learner guide & approved materials only
- Appeal process available



## Criteria for Funding

### **Minimum 75% attendance**


Based on SSG digital attendance records.

### **Assessed as 'Competent'**

Complete and pass both assessment components.

# Assessment Flow

At the assessment stage, follow these five steps in order.



**TIP** Courseware and the assessment are on the LMS: <https://lms-tms.tertiaryinfotech.com/>

# Digital Attendance



**Attendance is mandatory for SSG funding.**

Take AM, PM and Assessment digital attendance for every session.

Use the SSG / TMS QR code provided by your trainer.

**Minimum  
75%  
attendance  
required**

# Thank You!

Questions? We're here to help — during and after class.

[enquiry@tertiaryinfotech.com](mailto:enquiry@tertiaryinfotech.com) | +65 6100 0613