

Lab 01 — A Research Team of Sub-agents

Multi AI Agents Workflow for Content Creation · Day 1 · Claude Cowork

Course: Multi AI Agents Workflow for Content Creation (TGS-2023036153)

Day 1 · Topic 1 · about 45 minutes · slides 27–31

Surface: Claude Cowork

Features: sub-agents in parallel · market research · business analysis · competitive analysis · a recommendation

The story so far

October 2026. Rachel wants 40 booked chats a month by March. Before she spends a dollar, her partners want three answers: who to reach, what the funnel needs, and what competitors already publish.

Your goal

Before Horizon spends a dollar on content, find out who to reach, what the numbers need and what competitors do — with four specialist agents instead of one overloaded chat.

You'll build

research/: market research, business analysis, competitive analysis and a recommendation, every figure cited

What is in this folder

- `assets/business-brief.md`
- `assets/firm-metrics.csv`
- `assets/services.csv`
- `assets/research-agents-spec.md`
- `prompts.md` / `prompts.pdf` — every prompt, ready to paste
- `evidence/checklist.md` — what to capture as proof

Step by step

1. **Make the studio folder** — Create horizon-studio with a `data/` folder. Copy this lab's assets into `data/`.
2. **Open Cowork** — In the Claude desktop app choose Cowork, then point it at the horizon-studio folder.
3. **Brief the team** — Paste Prompt A. Cowork confirms the four roles and the plan — correct anything before it starts.
4. **Run them in parallel** — Paste Prompt B. Watch three sub-agents work at once in the progress panel.

5. **Audit the sources** — Open competitive-analysis.md. Every figure needs a source and a date; gaps say UNKNOWN.
6. **Decide** — Read recommendation.md. Change one thing you disagree with — it is your recommendation now.

The prompts

PROMPT A — Cowork: brief the team

I'm the marketing lead at Horizon Wealth Planning, an independent financial-planning firm in Singapore. Read data/business-brief.md and data/research-agents-spec.md.

We will run a research team of four sub-agents:

1. market-researcher — Singapore demand for financial planning: who, which channels, why now.
2. business-analyst — our funnel from data/firm-metrics.csv: what 40 chats a month needs.
3. competitive-analyst — six named competitors: offer, pricing model, channels, content themes.
4. strategy-advisor — reads the other three and writes the recommendation.

Rules for all: cite every figure with source and date, or write UNKNOWN. Never publish the INTERNAL columns. Confirm each role in one line and show me your plan. Do not start yet.

PROMPT B — Cowork: run the team

Run the research team now, using sub-agents.

- Start market-researcher, business-analyst and competitive-analyst in parallel, each as its own sub-agent, each saving its report to research/.
- When all three are done, the strategy-advisor writes research/recommendation.md exactly as research-agents-spec.md describes.

Then give me a five-line summary and a list of everything marked UNKNOWN.

Check your work

- ☐ Cowork confirmed all four roles before starting.
- ☐ Three research sub-agents ran at the same time.
- ☐ research/ holds market-research, business-analysis, competitive-analysis and recommendation.
- ☐ The competitive analysis names six real competitors, each with a source and a date.
- ☐ The business analysis shows conversion rates and the visits needed for 40 chats a month.
- ☐ No uncited figure; no INTERNAL fee anywhere in the reports.

If it goes wrong

- **Cowork did it all itself** — Say it directly: "use sub-agents — one per role, in parallel".
- **No time for research** — Use recommendation-sample.md from Lab 2's assets and carry on.

Stretch

- Add a fifth sub-agent, regulatory-scout, that lists the MAS and PDPA rules a financial firm's marketing must follow — with sources.

Why it matters: A sub-agent does its research in its own context and hands back only the result. Your main conversation stays clear for the decision.

Next

Lab 2 — Plan and Build the Horizon Website with Codex. Keep what you built — the next lab starts from it.

Safety

Use only this lab's synthetic data and your own accounts. Never paste a real API key, password, token or client data into a prompt or a file — tokens live only in .env. Every email address in the data must resolve to your own inbox. Nothing is published or sent without a person approving that exact version.