

Lab 12 — Hermes Agent Bots on a Kanban Board

Multi AI Agents Workflow for Content Creation · Day 2 · Hermes Agent (profiles, Kanban)

Course: Multi AI Agents Workflow for Content Creation (TGS-2023036153)

Day 2 · Topic 4 · about 45 minutes · slides 118–122

Surface: Hermes Agent (profiles, Kanban)

Features: open-source Hermes Agent · one profile per role · a Kanban role pipeline · review and human unblock · the same skills

The story so far

Rachel's IT adviser asks what happens if the studio's AI provider changes its prices or its rules. Try the same team on Hermes — open source, any model — with every hand-off on a Kanban board.

Your goal

The studio should not depend on one vendor. Hermes runs a team of named agent bots on any model, coordinated on a durable Kanban board you can watch, block and unblock.

You'll build

Three Hermes profiles, a researcher → writer → reviewer pipeline on the board, and Week 2 content approved by you

What is in this folder

- assets/hermes-setup.md
- assets/hermes-roles.md
- assets/kanban-pipeline.sh
- prompts.md / prompts.pdf — every prompt, ready to paste
- evidence/checklist.md — what to capture as proof

Step by step

1. **Install Hermes** — Follow hermes-setup.md; hermes setup with the API key your trainer gives you.
2. **Create the bots** — Run the three hermes profile create commands in hermes-roles.md, then copy skills/ into each profile.
3. **Start the board** — hermes kanban init, then hermes gateway start in a second terminal — it runs the dispatcher.
4. **Build the pipeline** — Paste the prompt into hermes (or run kanban-pipeline.sh).
5. **Watch it** — hermes kanban watch, or hermes dashboard for the board in your browser.
6. **Be the gate** — The reviewer blocks for you. Read the output, then hermes kanban unblock <id> — or request changes.

7. **Hand it to the gate** — Save the approved post as content/social/week-02/li-w02-01.md and submit it.

The prompts

PROMPT — Hermes: the planner

You are the planner for Horizon Wealth Planning.
Create a Kanban pipeline for Week 2 content, with dependencies:

1. researcher: facts for a LinkedIn post and a newsletter tip on CPF cash top-ups, from facts-2026.md only.
2. writer (parent: 1): write both in the sunny-voice skill.
3. reviewer (parent: 2): check with fact-check and fin-compliance; request changes until it passes, then block the task with the reason "Needs approval from a person".

Work in ~/horizon-studio/hermes. Show me the task ids and the board.

COMMANDS — terminal

```
hermes profile create researcher --description "Gathers and checks facts for Horizon content."  
hermes profile create writer --description "Writes Horizon posts in the sunny-voice skill."  
hermes profile create reviewer --description "Checks Horizon content; blocks for human approval."  
for p in researcher writer reviewer; do cp -R skills/* ~/.hermes/profiles/$p/skills/; done  
hermes kanban init  
hermes gateway start          # in a second terminal  
hermes kanban watch           # or: hermes dashboard  
hermes kanban unblock <task-id>
```

Check your work

- ☐ hermes profile list shows researcher, writer and reviewer.
- ☐ Each profile has the four studio skills.
- ☐ The board shows three tasks linked by parents.
- ☐ The writer started only after the researcher finished.
- ☐ The reviewer requested changes at least once, then blocked for you.
- ☐ You unblocked it; the post is submitted to approvals.csv.

If it goes wrong

- **Tasks stay in ready** — The gateway runs the dispatcher — start it with hermes gateway start.
- **Model errors** — A Claude subscription does not cover Hermes. Use the API key from your trainer.

Stretch

- Set a cheaper model for the researcher profile and a stronger one for the reviewer; compare the cost.

Why it matters: A Kanban board makes agent work visible and durable: every hand-off is a row anyone can read, and a blocked card is a person's turn.

Next

Lab 13 — Always On: a Telegram Bot That Asks Before It Acts. Keep what you built — the next lab starts from it.

Safety

Use only this lab's synthetic data and your own accounts. Never paste a real API key, password, token or client data into a prompt or a file — tokens live only in `.env`. Every email address in the data must resolve to your own inbox. Nothing is published or sent without a person approving that exact version.