

Lab 02 — Form the Team: Instructions for Each Agent

Multi AI Agents Workflow for Content Creation (ChatGPT Edition) · Day 1 · Codex (role cards, subagents)

Course: Multi AI Agents Workflow for Content Creation (ChatGPT Edition) (TGS-2023036153)

Day 1 · Topic 1 · about 45 minutes · slides 38–42

Surface: Codex (role cards, subagents)

Features: five role cards in agents/ · the roster in AGENTS.md · subagents by asking

The story so far

Rachel sketches the team on a whiteboard: a Lead and five specialists. Each needs a clear job, and a clear line it must never cross — this is a regulated firm.

Your goal

A team is a set of clear jobs. Write each agent's instructions once — its job, what it reads, what it delivers and what it must never do — and every campaign can call on it.

You'll build

agents/: five role cards, and the roster in AGENTS.md

What is in this folder

- assets/role-specs.md
- assets/team-org-chart.md
- solution/ — reference files from the verified build
- prompts.md / prompts.pdf — every prompt, ready to paste
- evidence/checklist.md — what to capture as proof

Step by step

1. **Read the roles** — Open role-specs.md: one Lead and five specialists, each with a job, skills, connectors and limits.
2. **Write the instructions** — Paste Prompt A.
3. **Read one closely** — Open agents/content-creator.md. Is the "Never" section strict enough for a regulated firm?
4. **Meet the team** — Paste Prompt B. Each agent introduces itself in one line — in its own context.
5. **Tighten one** — Change one instruction you disagree with and save it. It is your team.

The prompts

PROMPT A — Codex: five role cards

Read data/role-specs.md. Create five role cards in agents/ — growth-strategist.md, content-creator.md, creative-designer.md, website-designer.md and growth-analyst.md. Each has: the job, what it reads, what it delivers, the plugins it may use, what it must never do, and how it reports back. Then add "## The marketing team" to AGENTS.md: one line per agent pointing to its card, and the rule "when I ask for the team, use subagents — one per role — each following its card".

PROMPT B — Codex: meet the team

Use subagents — one per role, each following its card in agents/. Ask each: "In one line, who are you, and what is the first thing you would do for Horizon's goal of 40 booked chats a month?" Show me the five answers in a table.

Check your work

- ☐ Five files exist in agents/, one per specialist.
- ☐ Each card has job, reads, delivers, plugins, never and report sections.
- ☐ AGENTS.md lists the team and the subagent rule.
- ☐ Five agents answered, each in its own context.
- ☐ Every agent has a "never" rule that blocks publishing.
- ☐ You changed one instruction and can say why.

If it goes wrong

- **Codex answered alone** — Ask directly: "use subagents — one per role".

Stretch

- Add a sixth agent, community-manager, that drafts replies to comments — and decide what it must never do.

Why it matters: Clear jobs beat clever prompts. An agent that knows what it must never do is safer than one that is merely told to be careful.

Next

Lab 3 — Give Each Agent Its Skills. Keep what you built — the next lab starts from it.

Safety

Use only this lab's synthetic data and your own accounts. Never paste a real API key, password, token or client data into a prompt or a file — tokens live only in `.env`. Every email address in the data must resolve to your own inbox. Nothing is published or sent without a person approving that exact version.